

# Nina Stawski's (group 90) final project report

Illinois ID: ninas2

[GitHub repo link \(https://github.com/nstawski/dlh-final-project\)](https://github.com/nstawski/dlh-final-project)

## Introduction

### Background of the problem

#### Type of problem

This is a data preparation and processing problem. The authors of the article are testing a common belief that adding more data improves the resulting model performance. Their main hypothesis, which they subsequently prove, is that incorporating more data does not necessary improve the model performance. It can introduce spurious correlations, and hurt the resulting model performance rather than helping it.

#### What is the importance/meaning of solving the problem

The paper is challenging a common belief, meaning a lot of researchers are likely trying to incorporate as much data as they can expecting it would improve the performance of their models. The outcome of this research would provide guidance on the possible pitfalls and the cases where you wouldn't want to add external data - so it could set a new standard of processing and incorporating data for everyone in the field.

#### The difficulty of the problem

The problem is non-obvious and the paper is challenging the common belief held in the industry. The authors are putting a lot of state-of-the-art approaches to the test, and attempt to quantify the results as well as provide new standards and explanations. This is extremely hard to do so I believe the problem is difficult.

#### The state of the art methods and effectiveness

The "industry standard" way of improving model performance is adding more data from additional datasets, which the authors of this article prove to not be effective, and even being harmful in many cases.

One of the main issues causing the model performance decrease when adding more data from other sources is spurious correlations, which in case of x-rays could be coming even from the scanner artifacts, or other hospital-specific data. One of the state-of-the-art ways to mitigate this is balancing a dataset to reduce the influence of hospital-specific factors. While balancing definitely improved the situation, the resulting model performance was still in many cases worse than with a single-hospital dataset.

## **Paper explanation**

### **What did the paper propose**

The paper used four most-used chest x-ray datasets - MIMIC-CXR-JPG, CheXpert, PadChest, ChestXray8 - to disprove a popular belief that adding more data always would improve the performance of your model. They postulate that, for the specific x-ray data, even the scanners themselves, the way hospitals produce data, or send specific patients to specific places to do their scan, can introduce spurious correlations which, in many cases, significantly affect the worst group performance.

### **What is the innovations of the method**

Existing research (for example, John R Zech, Marcus A Badgeley, Manway Liu, Anthony B Costa, Joseph J Titano, and Eric Karl Oermann. Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: a cross-sectional study. PLoS medicine, 15(11): e1002683, 2018.) proves that adding a second dataset improves the average per-group accuracy. In contrast, the paper I am reproducing focuses on the worst per-group accuracy.

### **How well the proposed method work (in its own metrics)**

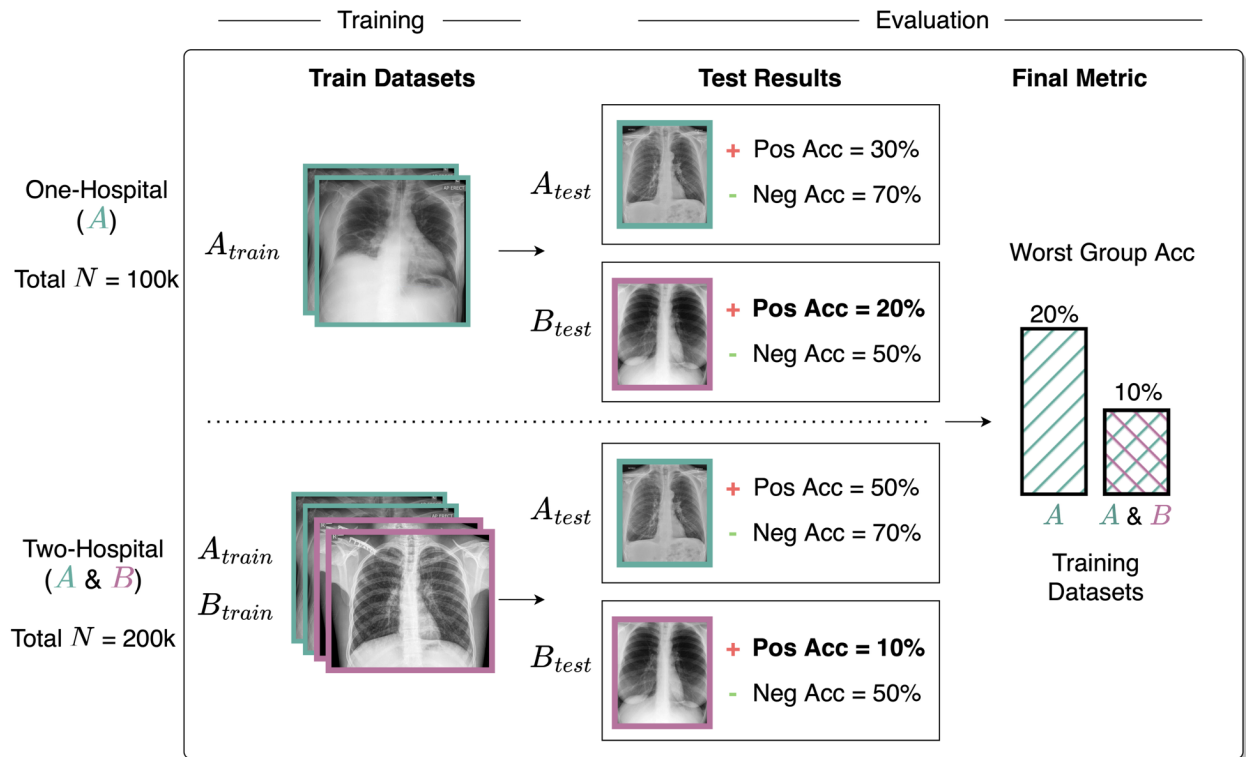
According to the article authors, their method works really well and proves that in nearly 50% of cases adding a second dataset, and even balancing it to reduce spurious correlations doesn't get the model to perform better than without that additional dataset. The models pick up on hospital-specific features even if those features weren't explicitly defined in the original data. They postulate that every CNN model, regardless of training disease or datasets, learns embeddings that can distinguish any of the hospital sources with near-perfect accuracy, even if the embeddings were trained via one or two hospitals' data.

**What is the contribution to the reasearch regime (referring the Background above, how important the paper is to the problem).**

The article cautions against blindly adding more datasets, and provides a number of approaches you can take if you still decide to do so. The conclusion is adding more data shouldn't be done blindly. The authors of the article definitely discourage the researchers from the most common approach of throwing data at the problem to improve model performance.

# Scope of Reproducibility:

List hypotheses from the paper you will test and the corresponding experiments you will run.



## Hypothesis 1

In 43% of training dataset/disease tasks, adding data from an external source hurts worst-group performance.

## Hypothesis 2

Balancing the dataset to reduce spurious correlations is often beneficial, but in the scenarios where adding an additional data source hurts generalization performance, it does not always improve generalization; in some cases, training on a balanced dataset achieves lower worst-group accuracy than training on datasets from one or two hospitals.

# Methodology

This methodology is the core of your project. It consists of run-able codes with necessary annotations to show the experiment you executed for testing the hypotheses.

The methodology at least contains two subsections **data** and **model** in your experiment.

## Python environment and package versions

I have originally set up the Jupyter Notebook and was running the code on my laptop. However, processing of the data took days and the training was promising to take months. My husband had a gaming computer with a powerful video card and Cuda available, so I ended up using his machine.

To keep the environment isolated, I installed `Anaconda` and created a separate environment for all my packages. I then ran `Jupyter` with access to local network and opened my notebook remotely from my laptop:

```
jupyter notebook --ip 192.168.x.xxx --port 8888
```

Original requirements listed a bit different versions, but due to package compatibility I had to update them. Most significant changes were: install a version of `torch` that supports Cuda, and upgrade `torchvision` to `0.9.1`

```
In [98]: env_df = pd.read_csv('env_packages.csv', header=None, names=['Package']
        styled_env_df = env_df.style.set_table_styles([
            {'selector': 'th, td',
             'props': [('min-width', '200px')]},
            {'selector': 'table',
             'props': [('width', '70%')]})
        ).set_properties(**{
            'background-color': 'white',
            'color': 'black',
            'border-color': 'black',
            'border-style': 'solid',
            'border-width': '1px'
        }).hide_index()

        styled_env_df
```

Out[98]:

Package	Version
python	3.6.13
pip	21.2.2
jupyter_core	4.8.1
imbalanced-learn	0.8.1
jupyter	1.0.0
matplotlib	2.2.2
numpy	1.19.5
pandas	1.1.0
pillow	8.4.0
scikit-learn	0.24.2
scipy	1.5.1
seaborn	0.11.2
torch	1.8.1+cu111
torchvision	0.9.1

In [44]:

```
# !pip install importlib  
# !pip install torch  
# !pip install torchvision  
# !pip install pandas  
# !pip install matplotlib  
# !pip install imblearn
```

```
In [99]: import numpy as np
import pandas as pd
from pathlib import Path
import os
from os.path import exists
import sys
import matplotlib.pyplot as plt
import seaborn as sns
from PIL import Image, ImageFile

import json
import random
from IPython.display import display
from datetime import datetime
from sklearn.metrics import f1_score, classification_report, accuracy_

import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset
from torchvision import datasets, models, transforms

from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler

import Data_Constants as Constants

#making sure all referenced files are reloaded
import importlib
importlib.reload(Constants)
```

```
Out[99]: <module 'Data_Constants' from 'C:\\Users\\Stan\\Documents\\GitHub\\dlh-final-project\\Data_Constants.py'>
```

It was a lot of debugging to make sure Cuda is available in the notebook, and it took me a few days to finally make it work. It made all processing and training code run a lot faster.

The next challenge I encountered was that the full source data did not fit into my laptop and I had to get an external storage to hold it - however, that external storage's speed wasn't keeping up. So instead, my husband got an upgrade to his internal storage with a very fast ssd.



```
In [100]: # os.environ['KMP_DUPLICATE_LIB_OK']='True'
# torch.set_default_device('cuda')

torch.set_default_tensor_type('torch.cuda.FloatTensor')
Tensor = torch.cuda.FloatTensor if torch.cuda.is_available() else torch.FloatTensor

print("Cuda is available:", torch.cuda.is_available())
```

Cuda is available: True

## Data

The study is using four datasets: MIMIC-CXR-JPG, CheXpert, PadChest, ChestXray8

The datasets are being filtered to include only frontal (PA/AP) images. Instances are labeled with one or more pathologies. Each dataset has a different set of diseases but they are preprocessed using code derived from ClinicalDG2 (Zhang et al., 2021) to extract the eight common labels and homogenize the datasets. Additionally, authors of the article created the Any label which indicates a positive label for any of the seven common disease labels, resulting in nine different binary labels. All experiments use the labels in a binary manner; a pathology is chosen as the target label, with an instance labeled 1 if the pathology of interest is present and 0 otherwise.

The authors apply an 80%/10%/10% subject-wise train/val/test split, with the same split used across seeds.

### MIMIC-CXR

1. [Obtain access \(https://mimic-cxr.mit.edu/about/access/\)](https://mimic-cxr.mit.edu/about/access/) to the MIMIC-CXR-JPG Database Database on PhysioNet and download the [dataset \(https://physionet.org/content/mimic-cxr-jpg/2.0.0/\)](https://physionet.org/content/mimic-cxr-jpg/2.0.0/). The best option is downloading from the GCP bucket:

```
gcloud auth login
mkdir MIMIC-CXR-JPG
gsutil -m rsync -d -r gs://mimic-cxr-jpg-2.0.0.physionet.org MIMIC-CXR-JPG
```

2. In order to obtain gender information for each patient, you will need to obtain access to [MIMIC-IV \(https://physionet.org/content/mimiciv/0.4/\)](https://physionet.org/content/mimiciv/0.4/). Download `core/patients.csv.gz` and place the file in the MIMIC-CXR-JPG directory.

### CheXpert

1. Sign up with your email address [here](#) (<https://stanfordmlgroup.github.io/competitions/chexpert/>).
2. Download either the original or the downsampled dataset (we recommend the downsampled version - `CheXpert-v1.0-small.zip`) and extract it.

## ChestX-ray8

1. Download the `images` folder and `Data_Entry_2017_v2020.csv` from the [NIH website](#) (<https://nihcc.app.box.com/v/ChestXray-NIHCC>).
2. Unzip all of the files in the `images` folder.

## PadChest

1. The paper uses a resized version of PadChest, which can be downloaded [here](#) (<https://academictorrents.com/details/96ebb4f92b85929eadfb16761f310a6d04105797>).
2. Unzip `images-224.tar`.
  - Statistics: include basic descriptive statistics of the dataset like size, cross validation split, label distribution, etc.
  - Data process: how do you manipulate the data, e.g., change the class labels, split the dataset to train/valid/test, refining the dataset.
  - Illustration: printing results, plotting figures for illustration.
  - You can upload your raw dataset to Google Drive and mount this Colab to the same directory. If your raw dataset is too large, you can upload the processed dataset and have a code to load the processed dataset.

## Data Processing

The original pre-processing for the article was done using the scripts outside of the Jupyter Notebook. Some of them didn't work for me, and the installation process didn't succeed despite multiple attempts either. Instead, I have adapted some of the original scripts to run in the notebook (with some modifications so they actually work with my data), using the external "Constants.py" file that points to the location of the datasets.

1. In `./Data_Constants.py`, update `image_paths` to point to each of the four directories that you downloaded.
2. Run the next two cells to pre-process the data

## Validating

I am using the validation and pre-processing code provided by the authors of the article, with some modifications to make it run as expected.

```
In [101]: #making sure constants are up to date if they were changed
importlib.reload(Constants)

def validate_mimic():
    img_dir = Path(Constants.image_paths['MIMIC'])
    meta_dir = Path(Constants.meta_paths['MIMIC'])

    print('meta_dir', meta_dir, os.getcwd())
    print('meta_dir', meta_dir/'mimic-cxr-2.0.0-metadata.csv')
    assert (meta_dir/'mimic-cxr-2.0.0-metadata.csv').is_file()
    assert (meta_dir/'mimic-cxr-2.0.0-negbio.csv').is_file()
    assert (meta_dir/'patients.csv').is_file()
    # modified the file that's being checked since I don't have the full
    # in the original script, the file in p19 was being checked.
    assert (img_dir/'p10/p10000032/s50414267/02aa804e-bde0afdd-112c0b3

def validate_cxp():
    img_dir = Path(Constants.image_paths['CXP'])
    if (img_dir/'CheXpert-v1.0').is_dir():
        cxp_subfolder = 'CheXpert-v1.0'
    else:
        cxp_subfolder = 'CheXpert-v1.0-small'
    assert (img_dir/cxp_subfolder/'train.csv').is_file()
    assert (img_dir/cxp_subfolder/'train/patient48822/study1/view1_fro
    assert (img_dir/cxp_subfolder/'valid/patient64636/study1/view1_fro

def validate_pad():
    img_dir = Path(Constants.image_paths['PAD'])
    meta_dir = Path(Constants.meta_paths['PAD'])
    assert (meta_dir/'PADCHEST_chest_x_ray_images_labels_160K_01.02.19
    assert (img_dir/'185566798805711692534207714722577525271_qb3lyn.pr

def validate_nih():
    img_dir = Path(Constants.image_paths['NIH'])
    meta_dir = Path(Constants.meta_paths['NIH'])
    assert (meta_dir/'Data_Entry_2017.csv').is_file()
    assert (img_dir/'images/00002072_003.png').is_file()

def validate_splits():
    for dataset in Constants.df_paths:
        for split in Constants.df_paths[dataset]:
            assert Path(Constants.df_paths[dataset][split]).is_file()
```

```
def validate_all():
    validate_mimic()
    validate_cxp()
    validate_nih()
    validate_pad()
```

## Data pre-processing setup

```
In [102]: # making sure constants are up to date if they were changed after running
importlib.reload(Constants)
```

```
def preprocess_mimic():
    img_dir = Path(Constants.image_paths['MIMIC'])
    meta_dir = Path(Constants.meta_paths['MIMIC'])
    out_folder = meta_dir/'clinicaldg'
    out_folder.mkdir(parents = True, exist_ok = True)

    patients = pd.read_csv(meta_dir/'patients.csv')
    labels = pd.read_csv(meta_dir/'mimic-cxr-2.0.0-negbio.csv')
    meta = pd.read_csv(meta_dir/'mimic-cxr-2.0.0-metadata.csv')

    df = meta.merge(patients, on = 'subject_id').merge(labels, on = ['subject_id', 'patient_id'])
    df['age_decile'] = pd.cut(df['anchor_age'], bins = list(range(0, 100)), labels = ['0-10', '10-20', '20-30', '30-40', '40-50', '50-60', '60-70', '70-80', '80-90', '90-100'])
    df['frontal'] = df['ViewPosition'].isin(['AP', 'PA'])

    df['path'] = df.apply(lambda x: os.path.join(f'p{str(x["subject_id"])}_{str(x["patient_id"])}_{str(x["study_id"])}_{str(x["view_position"])}_{str(x["projection"])}_{str(x["label"])}_{str(x["age_decile"])}_{str(x["frontal"])}'), axis=1)
    df.to_csv(out_folder/"preprocessed.csv", index=False)

def preprocess_pad():
    # I have modified this function from the original one, because I want to use the same folder for all datasets
    img_dir = Path(Constants.image_paths['PAD'])
    meta_dir = Path(Constants.meta_paths['PAD'])
    out_folder = meta_dir/'clinicaldg'
    out_folder.mkdir(parents=True, exist_ok=True)

    dtype_spec = {
        'ImageID': str,
        'StudyID': str,
        'PatientID': str,
        'PatientBirth': str, # converting this to the integer later to save space
        'PatientSex_DICOM': str,
        'ViewPosition_DICOM': str,
        'Projection': str,
        'Labels': str,
        'WindowCenter_DICOM': str,
        'WindowWidth_DICOM': str
    }
```

```

}

df = pd.read_csv(meta_dir/'PADCHEST_chest_x_ray_images_labels_160K')
df = df[['ImageID', 'StudyID', 'PatientID', 'PatientBirth', 'PatientAge']]
df = df[~df["Labels"].isnull()]
df = df[df["ImageID"].apply(lambda x: os.path.exists(os.path.join(meta_dir, x))) == True]
df = df[df.Projection.isin(['PA', 'L', 'AP_horizontal', 'AP'])]

df['frontal'] = ~(df['Projection'] == 'L')
df = df[~df['Labels'].apply(lambda x: 'exclude' in x or 'unchanged' in x)]

mapping = dict()
mapping['Effusion'] = ['hydropneumothorax', 'empyema', 'hemothorax']
mapping['Consolidation'] = ['air bronchogram']
mapping['No Finding'] = ['normal']

for pathology in Constants.take_labels:
    mask = df["Labels"].str.contains(pathology.lower())
    if pathology in mapping:
        for syn in mapping[pathology]:
            mask |= df["Labels"].str.contains(syn.lower())
    df[pathology] = mask.astype(int)

df['PatientBirth'] = df['PatientBirth'].dropna().astype(float).astype(int)
df['Age'] = 2017 - df['PatientBirth']
df.reset_index(drop=True).to_csv(out_folder/"preprocessed.csv", index=False)

def preprocess_cxp():
    img_dir = Path(Constants.image_paths['CXP'])
    out_folder = img_dir/'clinicaldg'
    if (img_dir/'CheXpert-v1.0/'/'train.csv').is_file():
        df = pd.concat([pd.read_csv(img_dir/'CheXpert-v1.0/'/'train.csv'),
                        pd.read_csv(img_dir/'CheXpert-v1.0/'/'valid.csv')],
                        ignore_index = True)
    elif (img_dir/'CheXpert-v1.0-small/'/'train.csv').is_file():
        df = pd.concat([pd.read_csv(img_dir/'CheXpert-v1.0-small/'/'train.csv'),
                        pd.read_csv(img_dir/'CheXpert-v1.0-small/'/'valid.csv')],
                        ignore_index = True)
    elif (img_dir/'train.csv').is_file():
        raise ValueError('Please set Constants.image_paths["CXP"] to be the directory and rerun this script.')
    else:
        raise ValueError("CheXpert files not found!")

    out_folder.mkdir(parents = True, exist_ok = True)

    df['subject_id'] = df['Path'].apply(lambda x: int(Path(x).parent.name))
    df['Path'] = df['Path'].apply(lambda x: str(x).replace("CheXpert-v1.0", "clinicaldg"))
    df.reset_index(drop = True).to_csv(out_folder/"preprocessed.csv", index=False)

```

```
def preprocess_nih():
    img_dir = Path(Constants.image_paths['NIH'])
    meta_dir = Path(Constants.meta_paths['NIH'])
    out_folder = meta_dir/'clinicaldg'
    out_folder.mkdir(parents = True, exist_ok = True)
    df = pd.read_csv(meta_dir/"Data_Entry_2017.csv")
    df['labels'] = df['Finding Labels'].apply(lambda x: x.split('|'))

    for label in Constants.take_labels:
        df[label] = df['labels'].apply(lambda x: label in x)
    df.reset_index(drop = True).to_csv(out_folder/"preprocessed.csv",
```

In [103]: %%script false --no-raise-error  
# skipping this cell since I already ran this.

```
if __name__ == '__main__':
    print("Validating paths...")
    validate_all()
    print("Preprocessing MIMIC-CXR...")
    preprocess_mimic()
    print("Preprocessing CheXpert...")
    preprocess_cxp()
    print("Preprocessing ChestX-ray8...")
    preprocess_nih()
    print("Preprocessing PadChest... This might take a few minutes...")
    preprocess_pad()
    print("Done.")
```

```
Validating paths...
meta_dir C:\Nina\e-root\data\mimic\physionet.org\files\mimic-cxr-jpg\
2.0.0 C:\Users\Stan\Documents\GitHub\dlh-final-project
meta_dir C:\Nina\e-root\data\mimic\physionet.org\files\mimic-cxr-jpg\
2.0.0\mimic-cxr-2.0.0-metadata.csv
Preprocessing MIMIC-CXR...
Preprocessing CheXpert...
Preprocessing ChestX-ray8...
Preprocessing PadChest... This might take a few minutes...
Done.
```

## Next, we need to resize and process the data.

I am using the code provided by the authors of the article to do this.

```
In [104]: def process_MIMIC(split, only_frontal):
    copy_subjectid = split['subject_id']
    split = split.drop(columns = ['subject_id']).replace(
        [[None], -1, "[False]", "[True]", "[ True]", 'UNABLE TO OE
        'DIVORCED'. 'SEPARATED'. '0-10'. '10-20'. '20-30'. '30-40
```

[http://192.168.4.249:8888/notebooks/Documents/GitHub/dlh-final-project/Nina\\_Stawski\\_final\\_project\\_stats.ipynb#Ablation-study](http://192.168.4.249:8888/notebooks/Documents/GitHub/dlh-final-project/Nina_Stawski_final_project_stats.ipynb#Ablation-study) Page 15 of 98

```

copy_subjectid = split['subject_id']
split = split.drop(columns = ['subject_id']).replace([[None], -1,
                                                    [0, 0, 0, 1, 1, "0-20", "20-40", "40-60",

split['subject_id'] = copy_subjectid.astype(str)
split['Sex'] = np.where(split['Sex']=='Female', 'F', split['Sex'])
split['Sex'] = np.where(split['Sex']=='Male', 'M', split['Sex'])
split = split.rename(
    columns = {
        'Pleural Effusion': 'Effusion',
        'Lung Opacity': 'Airspace Opacity'
    })
split['path'] = split['Path'].astype(str).apply(lambda x: os.path.
split['frontal'] = (split['Frontal/Lateral'] == 'Frontal')
if only_frontal:
    split = split[split['frontal']]
split['env'] = 'CXP'
split['study_id'] = split['path'].apply(lambda x: x[x.index('patie
return split[['subject_id', 'path', 'Sex', "Age", 'env', 'frontal', 's

def process_PAD(split, only_frontal):
    split['Age'] = np.where(split['Age'].between(0,19), 19, split['Age
    split['Age'] = np.where(split['Age'].between(20,39), 39, split['Ag
    split['Age'] = np.where(split['Age'].between(40,59), 59, split['Ag
    split['Age'] = np.where(split['Age'].between(60,79), 79, split['Ag
    split['Age'] = np.where(split['Age']>=80, 81, split['Age'])

    split = split.replace([[None], -1, "[False]", "[True]", "[ True]",
                            [0, 0, 0, 1, 1, "0-20", "20-40", "40-60",

    split.loc[split['Age'] == 0.0, 'Age'] = '0-20'
    split.loc[split['Age'].isnull(), 'Age'] = '0-20'
    split = split.rename(columns = {
        'PatientID': 'subject_id',
        'StudyID': 'study_id',
        'PatientSex_DICOM': 'Sex'
    })

    split.loc[~split['Sex'].isin(['M', 'F', '0']), 'Sex'] = '0'
    split['path'] = split['ImageID'].astype(str).apply(lambda x: os.p
    if only_frontal:
        split = split[split['frontal']]
    split['env'] = 'PAD'
    return split[['subject_id', 'path', 'Sex', "Age", 'env', 'frontal', 's

def split(df, split_portions = (0.8, 0.9), seed=0):
    # We don't want the data splits to be affected by seed
    # So lets temporarily set the seed to a static value

```



```

# ...lets temporarily set the seed to a static value...
rand_state = np.random.get_state()
np.random.seed(seed)

# Split our data (irrespective of the random seed provided in train)
subject_df = pd.DataFrame({'subject_id': np.sort(df['subject_id']).
subject_df['random_number'] = np.random.uniform(size=len(subject_d

train_id = subject_df[subject_df['random_number'] <= split_portion
valid_id = subject_df[(subject_df['random_number'] > split_portion
test_id = subject_df[subject_df['random_number'] > split_portions[

train_df = df[df.subject_id.isin(train_id.subject_id)]
valid_df = df[df.subject_id.isin(valid_id.subject_id)]
test_df = df[df.subject_id.isin(test_id.subject_id)]

# ...then return the random state back to what it was
np.random.set_state(rand_state)

return train_df, valid_df, test_df

def get_process_func(env):
    if env == 'MIMIC':
        return process_MIMIC
    elif env == 'NIH':
        return process_NIH
    elif env == 'CXP':
        return process_CXP
    elif env == 'PAD':
        return process_PAD
    else:
        raise NotImplementedError

```

In [105]: *# show data paths from constants*  
 Constants.df\_paths

```

def img_exists(path):
    return exists(path)

def is_diseased(row):
    # diseases = Constants.take_labels[1:]
    return int((row[Constants.take_labels[1:]]).sum() > 0)

```

**The following cell is pre-processing the data and will take a long time to run**

The cell below needs to run once, after that everything is saved into the CSV file and can be loaded from there. this block of code needs to re-run only if the data changed.

```

In [106]: %%script false --no-raise-error
# skipping this cell since I already ran this.

# loads data with random splits
print('This might take a while.')

for data_env in Constants.df_paths:
    print('Processing:', data_env)
    func = get_process_func(data_env)
    print('Got processing function, filtering by only frontal...')
    df_env = func(pd.read_csv(Constants.df_paths[data_env]), only_frontal=True)
    print('Filtering out the data without images...')
    df_env["img_exists"] = df_env["path"].apply(img_exists)
    print(df_env["img_exists"].value_counts())
    df_env = df_env[df_env["img_exists"]]

    df_env = df_env.fillna(0)

    print('Adding "All" column...')
    df_env["All"] = df_env.apply(is_diseased, axis=1)

    print('Saving results...')
    df_env.to_csv(f"{Constants.base_path}\\processed\\{data_env}.csv",
                  index=False)

    display(df_env)

print("Done.")

```

<b>144479</b>	112930952416074060371371014599496493673	C:\Nina\e-root\data\PadChest\images-224\128401...
<b>144480</b>	282743729971423358706056731890510600934	C:\Nina\e-root\data\PadChest\images-224\128401...
<b>144481</b>	52648743308541843883453242716226652771	C:\Nina\e-root\data\PadChest\images-224\128401...
<b>144482</b>	228646130593152933811948996634154201216	C:\Nina\e-root\data\PadChest\images-224\128401...
<b>144483</b>	137424047230303610602080410284588825286	C:\Nina\e-root\data\PadChest\images-224\128401...

99827 rows × 17 columns

Done.

## Resample data

```
In [107]: dfs = {}
print('Processing the data, splitting to all, train, val and test...')
for env in Constants.df_paths:
    func = get_process_func(env)
    df_env = pd.read_csv(f"{Constants.base_path}/processed/{env}.csv")

    print('Source:', env)
    print('Data length:', len(df_env))

    train_df, valid_df, test_df = split(df_env)
    dfs[env] = {
        'all': df_env,
        'train': train_df,
        'val': valid_df,
        'test': test_df
    }
    print(f'{env}: done.')

print('All done.')
```

```
Processing the data, splitting to all, train, val and test...
Source: MIMIC
Data length: 230693
MIMIC: done.
Source: CXP
Data length: 191229
CXP: done.
Source: NIH
Data length: 112120
NIH: done.
Source: PAD
Data length: 99827
PAD: done.
All done.
```

## Balancing the dataset

```
In [108]: def get_prop(df, column="Pneumonia"):
    num_instances = len(df)
    num_diseased = df[df[column] == 1][column].count()
    return num_diseased / (num_instances - num_diseased)

def get_resample_class(orig_prop, new_prop, resample_method):
    if new_prop > orig_prop:
```

```

    if new_prop > orig_prop:
        if resample_method == "over":
            return 1
        else:
            return 0
    if new_prop < orig_prop:
        if resample_method == "under":
            return 1
        else:
            return 0

def calculate_num_resample(df, orig_prop, new_prop, resample_method):
    pass

def balance_df_label(df, sampler, label_bal=0.05154780337262089, inverse_target = df["Pneumonia"] == 1):
    rus = sampler(random_state=0, sampling_strategy=label_bal if not inverse_target else 1 / label_bal)
    res_df, _ = rus.fit_resample(df, target)

    print(f"Previous pneumonia prop: {get_pneumonia_prop(df)} with {label_bal}")
    print(f"Resampled pneumonia prop: {get_pneumonia_prop(res_df)} with {label_bal}")

    return res_df

def balance_proportion(orig_df, new_df, resample_method="over", column):
    orig_df = orig_df.fillna(0.0)
    orig_prop = get_prop(orig_df, column)
    new_prop = get_prop(new_df, column)
    assert resample_method in ["over", "under"]
    resample_class = get_resample_class(orig_prop, new_prop, resample_method)
    print(f"Resampling '{column}' via '{resample_method}' on class {resample_class}")

    # Estimate the number of items we'll need to resample
    df_diseased = orig_df[orig_df[column] == 1.0]
    df_normal = orig_df[orig_df[column] == 0.0]
    num_diseased = len(df_diseased)
    num_normal = len(df_normal)
    assert num_diseased + num_normal == len(orig_df)

    if resample_method == "over":
        if resample_class == 0:
            new_num_normal = int(num_diseased / new_prop)
            print(f"Resampling normal samples from {num_normal} to {new_num_normal}")
            df_normal_rs = df_normal.sample(new_num_normal, replace=True)
            resampled_df = pd.concat([df_normal_rs, df_diseased])
        else:
            # Resample the pneumonia class
            # new_num_diseased = int(new_prop * num_normal)
            # print(f"Resampling diseased samples from {num_diseased} to {new_num_diseased}")
            # df_diseased_rs = df_diseased.sample(new_num_diseased, replace=True)
            # resampled_df = pd.concat([df_normal, df_diseased_rs])

```

```
target = df["Pneumonia"] == 1
rus = RandomOverSampler(random_state=0, sampling_strategy=
resampled_df, _ = rus.fit_resample(df, target)

resampled_df.sort_index(inplace=True)
print(f"New df proportion: {get_prop(resampled_df, column)}")
return resampled_df

# balance_proportion(dfs["MIMIC"]["train"], dfs["MIMIC"]["test"])
```

```
In [109]: dfs["CXP"]["train"]
```

```
Out[109]:
```

	subject_id	path
0	1	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
1	2	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
2	2	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
3	3	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
4	4	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
...	...	...
191222	64734	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
191223	64735	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
191225	64737	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
191227	64739	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
191228	64740	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...

153411 rows × 17 columns

```
In [119]: def balance_df_label(df, sampler, label_bal=0.05154780337262089, invert
target = df["Pneumonia"] == (1 if not invert else 0)
rus = sampler(random_state=42, sampling_strategy=label_bal if not
res_df, _ = rus.fit_resample(df, target)

print(f"Previous pneumonia prop: {get_prop(df)} with {len(df)} ins
print(f"Resampled pneumonia prop: {get_prop(res_df)} with {len(res

return res_df

# # uncomment this code if you want a balanced dataset
# print('Balancing...')
# mimic_balanced = balance_df_label(dfs["MIMIC"]["train"], RandomOverS
# cxp_balanced = balance_df_label(dfs["CXP"]["train"], RandomOverSampl
# print('Done.')

# # Balance the size of the two datasets
# n = len(cxp_balanced)
# mimic_balanced = mimic_balanced.sample(n)
```

```
In [117]: import warnings
warnings.filterwarnings('ignore')
```

## Calculating stats

```
In [120]: stat_rows = []
num_instances = []

disease_labels = ["Pneumonia", "Cardiomegaly", "Edema", "Effusion", "A
target_labels = disease_labels + ["Any", "No Finding"]
all_labels = target_labels + ["Num Instances"]

dfs2 = {}

for env in dfs:
    df = dfs[env]['all']
    df['Any'] = (df[disease_labels] > 0).any(axis=1).astype(int)

    # keep only every 30th sample for the dataset to reduce the size o
    # I am keeping the full dataset coe commented out to be able to ea
    df2 = df[df.index % 30 == 0]
    dfs2[env] = {}
    dfs2[env]['all'] = df2

    train_df, valid_df, test_df = split(df2)
    dfs2[env] = {
        'all': df env,
```



```

        'train': train_df,
        'val': valid_df,
        'test': test_df
    }

    totals = {}
    totals['Dataset'] = env
#     totals['Num Instances'] = len(df)
    totals['Num Instances'] = len(df2)
    num_instances.append(totals['Num Instances'])

    for label in target_labels:
#         if label in df.columns:
#             totals[label] = df[label].sum() / len(df)
        if label in df2.columns:
            totals[label] = df2[label].sum() / len(df2)
        else:
            totals[label] = 0.0

    stat_rows.append(totals)

stat_df = pd.DataFrame(stat_rows)
stat_df.set_index('Dataset', inplace=True)

ordered_cols = all_labels
stat_df = stat_df[ordered_cols]

transposed_stat_df = stat_df.T

styled_transposed_stat_df = transposed_stat_df.style.apply(
    lambda x: ["background-color: lightblue" if x.name != 'Num Instances' else ''],
    axis=1
).background_gradient(cmap='Blues', subset=pd.IndexSlice[target_labels, :])
styled_transposed_stat_df = styled_transposed_stat_df.format("{:.2%}")
styled_transposed_stat_df = styled_transposed_stat_df.format("{:,.0f}")

styled_transposed_stat_df

```

Out[120]:

Dataset	MIMIC	CXP
Pneumonia	6.87%	2.68%
Cardiomegaly	16.62%	12.27%
Edema	11.70%	25.77%
Effusion	22.94%	40.20%
Atelectasis	20.42%	16.02%
Pneumothorax	4.23%	9.33%

<b>Consolidation</b>	4.68%	6.59%
<b>Any</b>	50.51%	71.14%
<b>No Finding</b>	35.01%	8.69%
<b>Num Instances</b>	7,690	6,375

Here is the table from the article for comparison:

Table 1: Total number of instances and disease prevalence in each dataset.

Target Label	MIMIC	CXP	NIH	PAD
Pneumonia	6.82%	2.43%	1.31%	4.84%
Cardiomegaly	17.05%	12.38%	2.51%	9.15%
Edema	11.83%	26.01%	2.11%	1.23%
Effusion	23.18%	40.28%	11.94%	5.99%
Atelectasis	20.11%	15.47%	10.33%	5.50%
Pneumothorax	4.19%	9.25%	4.66%	0.31%
Consolidation	4.67%	6.81%	4.19%	1.56%
Any	50.73%	70.35%	28.04%	23.03%
No Finding	34.76%	8.98%	53.65%	36.12%
Num Instances	243k	192k	113k	100k

Looks like the distribution of the labels in the original dataset, while not the same, still is close enough.

## Citation to the original paper

- Rhys Compton; Lily Zhang; Aahlad Puli; Rajesh Ranganath, When More is Less: Incorporating Additional Datasets Can Hurt Performance By Introducing Spurious Correlations, arXiv preprint, 2023-08-09, Accepted at MLHC 2023, doi: [10.48550/arXiv.2308.04431](https://doi.org/10.48550/arXiv.2308.04431) (<https://doi.org/10.48550/arXiv.2308.04431>)

## Original paper repo

- [ood-generalization](https://github.com/basedrhys/ood-generalization/tree/master) (<https://github.com/basedrhys/ood-generalization/tree/master>)

## Model

The model includes the model definition which usually is a class, model training, and other necessary parts.

## Model architecture

In the article, the authors use the same model architecture as Zhang et al. (2021): a **DenseNet-121** network (Huang et al., 2017) initialized with pre-trained weights from ImageNet (Deng et al., 2009). The final layer is replaced with a **two-output linear layer** (for binary classification). For simplicity, the authors only consider binary disease classification.

## Model Training

For training the network, all images are resized to **224 × 224** and normalized to the ImageNet (Deng et al., 2009) mean and standard deviation.

During training, the following image augmentations are applied:

- random horizontal flip
- random rotation up to 10 degrees
- a crop of random size (75% - 100%) and aspect ratio (3/4 to 4/3)

All runs use **Adam** with **lr = 1e-5** and **batch size = 128**, which was found to be a performant configuration in early tuning ((Zhang et al., 2021) use lr = 5e-4 and batch size = 32).

All test results are obtained using the optimal model found during training as measured by the highest validation macro-F1 score (following (Fiorillo et al., 2021; Berenguer et al., 2022)) as it gives a robust ranking of model performance under imbalanced labels.

```
In [121]: # This is the model defined and provided by the authors of the article.
# While they are using densenet 121 for the article, the provided model
# is a resnet.

class EmbModel(nn.Module):
    # I had to add the num_labels parameter to reduce the resulting re
    def __init__(self, emb_type, feature_size_override, pretrain, conc
        super().__init__()
        self.emb_type = emb_type
        self.pretrain = pretrain
        self.concat_features = concat_features
        self.num_labels = num_labels

    assert emb_type in ["densenet121", "densenet201", "resnet"], f

    if emb_type == 'densenet121':
```

```

        model = models.densenet121()
        self.encoder = nn.Sequential(*list(model.children())[:-1])
        self.emb_dim = model.classifier.in_features
    elif emb_type == 'densenet201':
        model = models.densenet201()
        self.encoder = nn.Sequential(*list(model.children())[:-1])
        self.emb_dim = model.classifier.in_features
    elif emb_type == 'resnet':
        model = models.resnet50()
        self.encoder = nn.Sequential(*list(model.children())[:-1])
        self.emb_dim = list(model.children())[-1].in_features

    print("\nEmb Dim:")
    print(self.emb_dim)

    if feature_size_override:
        print(f"Manually setting output dim to {feature_size_override}")
        self.emb_dim = feature_size_override
        print(self.emb_dim)

    self.n_outputs = self.emb_dim + concat_features
    self.final_layer = nn.Linear(self.n_outputs, self.num_labels)

    nn.init.kaiming_normal_(self.final_layer.weight, mode='fan_out')

def forward(self, inp):
    if isinstance(inp, dict): # dict with image and additional features
        x = inp['img']
        concat = inp['concat']
        assert(concat.shape[-1] == self.concat_features)
    else: # tensor image
        assert(self.concat_features == 0)
        x = inp

    x = self.encoder(x).squeeze(-1).squeeze(-1)
    if "densenet" in self.emb_type:
        x = F.relu(x)
        x = F.avg_pool2d(x, kernel_size = 7).view(x.size(0), -1)

    if isinstance(inp, dict):
        x = torch.cat([x, concat], dim = -1)

    x = self.final_layer(x)
    return x

```

# Training

I wasn't able to run the training code provided by the authors of the article - the setup didn't work for me neither on my MacBook Pro laptop, nor on my husband's Windows 10 gaming computer.

To proceed, I instead wrote my own training code using the standard approach learned in class and homeworks.

## Hyperparameters used

- Model: densenet121
- Number of epochs for each model trainig: 10
- Hidden size: 1024 since I am setting the `feature_size_override` to 1024
- Batch size: 128
- Learning rate:  $1e-5$
- Optimizer: Adam

# Computational requirements

It is possible to run this code on a CPU with minor modifications. However, since I moved to another computer with GPU, some portions of this notebook send the computation to Cuda directly (todo: rewrite so it checks for Cuda and sends to the appropriate device).

## Hardware and software

- AMD Ryzen 7 7800X4D 8-Core Processor (4.20 GHz)
- 64 GB RAM
- NVMe Samsung SSD 970 EVO - 1TB
- Windows 10 64-bit

## Training requirements

I was not able to train on the full dataset since even on GPU one epoch of one model was running for 2-4 hours depending on the number of batches. This would require roughly 16 days to finish the whole training.

Initially, I attempted to run the training on the full dataset, but a number of circumstances (out of memory, kernel panic, random automatic Windows updates, power down, kids getting to the computer and switching the power supply off) proved that the expectation to run the training continuously for days to be completely unrealistic.

Instead, I modified my dataset to pick every 30th entry and discard the rest. As a result, I was able to run the training multiple times with different parameters when needed, both on balanced and unbalanced datasets.

- Average epoch running time: 6min
- Average time to complete all training: 12h
- Total number of attempts: 200+

## Creating a data loader

The authors of the article have a script to load the data in different configurations. I wasn't able to make it work because of the errors, so instead I am partially reusing it and creating my own Dataset class and a data loader.

In [122]:

```
ImageFile.LOAD_TRUNCATED_IMAGES = True # I was getting errors during t
```

In [123]:

```
class MultiEnvDataset(Dataset):
    def __init__(self, dataframes, subset='train', envs=None, transform=None):
        """
        Initializes the dataset with data from multiple environments and labels.
        :param dataframes: A dictionary with environment keys, each containing a DataFrame.
        :param subset: The subset to load ('train', 'val', or 'test').
        :param envs: A list of environment names to include. If None, all environments are used.
        :param transform: PyTorch transforms to apply to the images.
        """
        if envs is None:
            envs = list(dataframes.keys())

        self.data = pd.concat([dataframes[env][subset] for env in envs])
        self.label_columns = ["No Finding", "Atelectasis", "Cardiomegaly", "Pneumothorax", "Consolidation", "Edema"]
        self.transform = transform

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        img_path = self.data.iloc[idx]['path']
        image = Image.open(img_path).convert('RGB') # Converts to RGB

        if self.transform:
            image = self.transform(image)

        labels = Tensor(self.data.iloc[idx][self.label_columns].values)
        if torch.isnan(labels).any():
            raise ValueError("NaN values found in labels")

        return image, labels
```

In [124]:

```
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
```

```
In [125]: envs_list = ["CXP"], ["MIMIC"], ["NIH"], ["PAD"], ["CXP", "NIH"], ["CX
env_list_map = {
    "cxp": 0,
    "mimic": 1,
    "nih": 2,
    "pad": 3,
    "cxp_nih": 4,
    "cxp_pad": 5,
    "mimic_cxp": 6,
    "mimic_nih": 7,
    "mimic_pad": 8,
    "nih_pad": 9,
    "cxp_mimic_nih_pad": 10,
}

# a few functions to simplify getting the right names of the dataset c
def get_dataset_index(env_name):
    return env_list_map[env_name]

def get_env_from_list(env_name):
    return envs_list[env_list_map[env_name]]
```

```
In [126]: datasets = []
for env in envs_list:
    elem = {
        "env": env
    }
    for subset in ["train", "val"]:
        elem[subset] = {}
#         elem[subset]["dataset"] = MultiEnvDataset(dfs, subset=subset
        elem[subset]["dataset"] = MultiEnvDataset(dfs2, subset=subset,
        elem[subset]["loader"] = torch.utils.data.DataLoader(elem[subs
    datasets.append(elem)
print("Done.")
```

Done.

## Metrics

```
In [127]: loss_func = nn.BCEWithLogitsLoss()

max_batches = 10

def calculate_accuracies(outputs, labels):
    predictions = torch.sigmoid(outputs) > 0.5
    predictions = predictions.to(labels.device)
    correct_pred = (predictions == labels)
```



```

        accuracies = correct_pred.float().mean(axis=0)
        return accuracies

def calculate_f1(outputs, labels):
    predictions = torch.sigmoid(outputs) > 0.5
    predictions = predictions.to(labels.device)

    predictions = predictions.detach().cpu().numpy()
    labels = labels.detach().cpu().numpy()

    f1 = f1_score(labels, predictions, average=None)
    return f1

def train_model_one_epoch(model, train_loader, loss_func, optimizer):
    print("Starting training...")
    start = datetime.now()
    prev = start
    model.train()
    running_loss = 0
    total_accuracy = []
    total_f1_scores = []

    print('number of batches:', len(train_loader))
    for batch, (inputs, labels) in enumerate(train_loader):
        inputs = inputs.cuda()

        optimizer.zero_grad()
        outputs = model(inputs)

        if torch.isnan(outputs).any():
            raise ValueError("NaN detected in model outputs")

        loss = loss_func(outputs, labels)
        if torch.isnan(loss).any():
            raise ValueError("NaN detected in loss computation")

        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.)
        optimizer.step()
        running_loss += loss.item() * inputs.size(0)

        accuracies = calculate_accuracies(outputs, labels)
        f1_scores = calculate_f1(outputs, labels)
        total_accuracy.append(accuracies)
        total_f1_scores.append(f1_scores)

        if batch % 100 == 0:
            mid = datetime.now()
            print("time passed from the beginning", mid-start)
            print('batch', batch + 1, 'time passed:', mid-prev)
            prev = mid

```

prev - mid

```

epoch_loss = running_loss / len(train_loader.dataset)
end = datetime.now()
print("epoch done in", end-start, "number of batches:", batch)
epoch_accuracy = torch.stack(total_accuracy).mean(dim=0)
epoch_f1 = torch.tensor(total_f1_scores).mean(dim=0)
return epoch_loss, epoch_accuracy, epoch_f1

def validate_model(model, val_loader, loss_func):
    model.eval()
    running_loss = 0
    total_accuracy = []
    total_f1_scores = []
    with torch.no_grad():
        for inputs, labels in val_loader:
            inputs = inputs.cuda()
            outputs = model(inputs)
            loss = loss_func(outputs, labels)
            running_loss += loss.item() * inputs.size(0)

            accuracies = calculate_accuracies(outputs, labels)
            f1_scores = calculate_f1(outputs, labels)
            total_accuracy.append(accuracies)
            total_f1_scores.append(f1_scores)

    epoch_loss = running_loss / len(val_loader.dataset)
    epoch_accuracy = torch.stack(total_accuracy).mean(dim=0)
    epoch_f1 = torch.tensor(total_f1_scores).mean(dim=0)
    return epoch_loss, epoch_accuracy, epoch_f1

```

```

In [128]: def saveModel(model, env=None):
    now = datetime.now()
    dt_string = now.strftime("%d-%m-%Y-%H-%M-%S")

    model_file_name = "model/model-snapshot-"
    if (env):
        model_file_name += "env_" + "_".join(env) + "_"
    model_file_name += dt_string + ".pth"

    torch.save(model.state_dict(), model_file_name)

```

In [19]:

```

%%script false --no-raise-error
# skipping the training since I already ran it in different variations

num_epoch = 10

metrics_df = pd.DataFrame(columns=["env", "epoch", "train_loss", "valid_loss", "train_accuracy", "valid_accuracy", "train_f1", "valid_f1"])

for dataset in datasets:
    model = EmbModel(emb_type="densenet121", feature_size_override=1024)
    model.cuda()
    model.train()
    optimizer = torch.optim.Adam(model.parameters(), lr=1e-5)
    print("Processing dataset env:", dataset["env"])
    for i in range(num_epoch):
        train_loss, train_accuracy, train_f1 = train_model_one_epoch(model, dataset, optimizer)
        valid_loss, valid_accuracy, valid_f1 = validate_model(model, dataset)

        print("Epoch: %.2f, Train Loss: %.2f, Validation Loss: %.2f" % (i+1, train_loss, valid_loss))

        # Convert tensors to CPU for DataFrame update
        train_accuracy = train_accuracy.cpu().numpy()
        valid_accuracy = valid_accuracy.cpu().numpy()
        train_f1 = train_f1.cpu().numpy()
        valid_f1 = valid_f1.cpu().numpy()

        # Append metrics to DataFrame
        metrics_df = metrics_df.append({
            "env": dataset["env"],
            "epoch": i + 1,
            "train_loss": train_loss,
            "valid_loss": valid_loss,
            "train_accuracy": np.mean(train_accuracy),
            "valid_accuracy": np.mean(valid_accuracy),
            "worst_train_accuracy": np.min(train_accuracy),
            "worst_valid_accuracy": np.min(valid_accuracy),
            "train_f1": np.mean(train_f1),
            "valid_f1": np.mean(valid_f1)
        }, ignore_index=True)

    saveModel(model, env=dataset["env"])
print("All done.")

```

Emb Dim:

1024

Manually setting output dim to 1024

1024

Processing dataset env: ['CXP']

Starting training...

number of batches: 40

.....

```

time passed from the beginning 0:00:02.636734
batch 1 time passed: 0:00:02.636734
epoch done in 0:00:26.762730 number of batches: 39
Epoch: 1.00, Train Loss: 0.90, Validation Loss: 0.62
Starting training...
number of batches: 40
time passed from the beginning 0:00:00.602461
batch 1 time passed: 0:00:00.602461
epoch done in 0:00:24.425553 number of batches: 39

```

```

C:\Users\Stan\anaconda3\envs\gpu\lib\site-packages\sklearn\metrics\_c
lassification.py:1406: UndefinedMetricWarning: F-score is ill-defined

```

```

In [22]: %%script false --no-raise-error
# skipping this cell since I already ran this.

# Save DataFrame to CSV

df_now = datetime.now()
df_dt_string = df_now.strftime("%d-%m-%Y-%H-%M-%S")
metrics_df.to_csv(f"stats/{df_dt_string}_training_metrics.csv", index=

```

## Validating the saved models and visualizing results

```

In [129]: def predict(model, val_loader, device='cuda'):
            model.eval()
            model.to(device)
            all_preds = []
            all_preds_raw = []
            all_labels = []

            print('Started prediction validation')
            print('Number of batches:', len(val_loader))
            predict_start_time = datetime.now()
            with torch.no_grad():
                for batch, (images, labels) in enumerate(val_loader):
                    print("Batch number:", batch+1, "of", len(val_loader))
                    images = images.to(device)
                    labels = labels.to(device)

                    outputs = model(images)
                    probabilities = torch.sigmoid(outputs)

                    preds = (probabilities > 0.5)

                    any_disease = torch.any(preds[:, 1:], dim=1, keepdim=True)
                    any_probability = torch.max(probabilities[:, 1:], dim=1, keepdim=True)
                    any_label = torch.any(labels[:, 1:], dim=1, keepdim=True)

                    preds = torch.cat((preds, any_disease), dim=1)
                    probabilities = torch.cat((probabilities, any_probability), dim=1)
                    labels = torch.cat((labels, any_label), dim=1)

                    all_preds_raw.append(probabilities.cpu().numpy())
                    all_preds.append(preds.cpu().numpy())
                    all_labels.append(labels.cpu().numpy())

            all_preds = np.vstack(all_preds)
            all_preds_raw = np.vstack(all_preds_raw)
            all_labels = np.vstack(all_labels)

            predict_end_time = datetime.now()

            print('Done.')
            print('Prediction took:', predict_end_time-predict_start_time)

            return all_preds, all_preds_raw, all_labels

```

```
In [130]: def calculate_per_label_accuracy(predictions, labels):
    accuracies = {}
    num_labels = labels.shape[1]

    for i in range(num_labels):
        label_preds = predictions[:, i]
        label_true = labels[:, i]
        accuracies[target_labels[i]] = accuracy_score(label_true, label_preds)
    return accuracies

def calculate_stats(predictions, probabilities, labels, source):

    # Calculate overall accuracy
    accuracy = accuracy_score(labels, predictions)
    print(f"Overall Accuracy: {accuracy:.2%}")

    # Detailed classification report for each disease label
    report = classification_report(labels, predictions, target_names=target_labels)
    report_df = pd.DataFrame(report).transpose()
    report_df['source'] = source
    print("Classification Report:")
    print(report_df)

    label_accuracies = calculate_per_label_accuracy(predictions, labels)
    worst_label = min(label_accuracies, key=lambda x: label_accuracies[x])
    worst_accuracy = label_accuracies[worst_label]

    accuracy_df = pd.DataFrame(list(label_accuracies.items()), columns=['label', 'accuracy'])
    accuracy_df['Accuracy'] = accuracy_df['Accuracy'].apply(lambda x: f"{x:.2%}")
    accuracy_df['source'] = source

    print(accuracy_df)
    print(f"Worst Performing Label: {worst_label} with an accuracy of {worst_accuracy:.2%}")
    return report_df, accuracy_df
```

```
In [131]: # repeating the labels code so I don't have to re-run the cell way above
target_labels = ["No Finding", "Atelectasis", "Cardiomegaly", "Effusion", "Lung Lesion", "Lung Opacity", "Pneumothorax", "Vascular Calcification"]

device = 'cuda' if torch.cuda.is_available() else 'cpu'
```

## Loading trained models

I separated the loading code into different cells rather than had a cycle in one, to be able to pick and choose which parts I run. The model names are hardcoded with the ones currently in the repo - if you are running the training code above, the new snapshots will be created and the file names should be updated below.

## MIMIC only

```
In [133]: print("Model trained on MIMIC")
model_MIMIC = EmbModel(emb_type="densenet121", feature_size_override=1024)
model_MIMIC.load_state_dict(torch.load("model/balanced/model-snapshot-1024.pth"))
model_MIMIC.eval()
```

Model trained on MIMIC

Emb Dim:

1024

Manually setting output dim to 1024

1024

```
Out[133]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True)
```

```
In [134]: data_loader_MIMIC = datasets[get_dataset_index("mimic")]["val"]["loader"]
mimic_predictions, mimic_probabilities, mimic_labels = predict(model_M)
print('labels', mimic_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 6
Batch number: 1 of 6
Batch number: 2 of 6
Batch number: 3 of 6
Batch number: 4 of 6
Batch number: 5 of 6
Batch number: 6 of 6
Done.
Prediction took: 0:00:35.871596
labels [[0. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 ...
 [1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 1.]
 [1. 0. 0. ... 0. 0. 0.]]
Done
```



In [135]: `mimic_report, mimic_accuracy = calculate_stats(mimic_predictions, mimic_labels)`

Overall Accuracy: 25.81%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.715190	0.436293	0.541966	259.0	mimic
Atelectasis	0.269231	0.088050	0.132701	159.0	mimic
Cardiomegaly	0.153846	0.016393	0.029630	122.0	mimic
Effusion	0.471910	0.280000	0.351464	150.0	mimic
Pneumonia	0.000000	0.000000	0.000000	42.0	mimic
Pneumothorax	0.000000	0.000000	0.000000	29.0	mimic
Consolidation	0.000000	0.000000	0.000000	36.0	mimic
Edema	0.156250	0.070423	0.097087	71.0	mimic
Any	0.701389	0.274457	0.394531	368.0	mimic
micro avg	0.561866	0.224110	0.320416	1236.0	mimic
macro avg	0.274202	0.129513	0.171931	1236.0	mimic
weighted avg	0.474759	0.224110	0.299259	1236.0	mimic
samples avg	0.248288	0.231509	0.235901	1236.0	mimic

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic
5	Pneumothorax	96.08%	mimic
6	Consolidation	95.14%	mimic
7	Edema	87.43%	mimic
8	Any	58.11%	mimic

Worst Performing Label: Any with an accuracy of 58.11%

```
In [136]: def combine_with_existing(combined_accuracy, combined_report, add_accuracy, add_report):
    combined_accuracy = pd.concat([combined_accuracy, add_accuracy], ignore_index=True)
    combined_report = pd.concat([combined_report, add_report], ignore_index=True)

    print(combined_accuracy)
    print(combined_report)

    return combined_accuracy, combined_report
```

## PAD only

```
In [137]: print("Model trained on PAD")
model_PAD = EmbModel(emb_type="densenet121", feature_size_override=1024)
model_PAD.load_state_dict(torch.load("model/balanced/model-snapshot-emb.pkl"))
model_PAD.eval()
```

Model trained on PAD

Emb Dim:  
1024  
Manually setting output dim to 1024  
1024

```
Out[137]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True)
```

```
In [138]: data_loader_PAD = datasets[get_dataset_index("pad")]["val"]["loader"]
pad_predictions, pad_probabilities, pad_labels = predict(model_PAD, data_loader_PAD)
print('labels', pad_labels)
print("Done")
```

Started prediction validation  
Number of batches: 3  
Batch number: 1 of 3  
Batch number: 2 of 3  
Batch number: 3 of 3  
Done.  
Prediction took: 0:00:02.932136  
labels [[1. 0. 0. ... 0. 0. 0.]  
[0. 0. 0. ... 0. 0. 0.]  
[0. 0. 0. ... 0. 0. 1.]  
...  
[0. 0. 1. ... 0. 0. 1.]  
[1. 0. 0. ... 0. 0. 0.]  
[0. 1. 0. ... 0. 0. 1.]]  
Done

In [139]: pad\_report, pad\_accuracy = calculate\_stats(pad\_predictions, pad\_probab

Overall Accuracy: 46.04%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.656250	0.338710	0.446809	124.0	pad
Atelectasis	0.000000	0.000000	0.000000	16.0	pad
Cardiomegaly	0.000000	0.000000	0.000000	30.0	pad
Effusion	0.200000	0.043478	0.071429	23.0	pad
Pneumonia	0.000000	0.000000	0.000000	19.0	pad
Pneumothorax	0.000000	0.000000	0.000000	0.0	pad
Consolidation	0.000000	0.000000	0.000000	3.0	pad
Edema	0.000000	0.000000	0.000000	3.0	pad
Any	0.416667	0.065789	0.113636	76.0	pad
micro avg	0.533333	0.163265	0.250000	294.0	pad
macro avg	0.141435	0.049775	0.070208	294.0	pad
weighted avg	0.400142	0.163265	0.223413	294.0	pad
samples avg	0.135671	0.137195	0.136179	294.0	pad

	Label	Accuracy	source
0	No Finding	68.29%	pad
1	Atelectasis	94.21%	pad
2	Cardiomegaly	90.55%	pad
3	Effusion	92.07%	pad
4	Pneumonia	92.99%	pad
5	Pneumothorax	100.00%	pad
6	Consolidation	99.09%	pad
7	Edema	98.78%	pad
8	Any	76.22%	pad

Worst Performing Label: No Finding with an accuracy of 68.29%

In [140]: combined\_accuracy, combined\_report = combine\_with\_existing(mimic\_report

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic
5	0.000000	0.000000	0.000000	29.0	mimic
6	0.000000	0.000000	0.000000	36.0	mimic
7	0.156250	0.070423	0.097087	71.0	mimic
8	0.701389	0.274457	0.394531	368.0	mimic
9	0.561866	0.224110	0.320416	1236.0	mimic
10	0.274202	0.129513	0.171931	1236.0	mimic
11	0.474759	0.224110	0.299259	1236.0	mimic
12	0.248288	0.231509	0.235901	1236.0	mimic
13	0.656250	0.338710	0.446809	124.0	pad
14	0.000000	0.000000	0.000000	16.0	pad
15	0.000000	0.000000	0.000000	30.0	pad
16	0.200000	0.043478	0.071429	23.0	pad
17	0.000000	0.000000	0.000000	19.0	pad
18	0.000000	0.000000	0.000000	0.0	pad
19	0.000000	0.000000	0.000000	3.0	pad
20	0.000000	0.000000	0.000000	3.0	pad
21	0.416667	0.065789	0.113636	76.0	pad
22	0.533333	0.163265	0.250000	294.0	pad
23	0.141435	0.049775	0.070208	294.0	pad
24	0.400142	0.163265	0.223413	294.0	pad
25	0.135671	0.137195	0.136179	294.0	pad

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic
5	Pneumothorax	96.08%	mimic
6	Consolidation	95.14%	mimic
7	Edema	87.43%	mimic
8	Any	58.11%	mimic
9	No Finding	68.29%	pad
10	Atelectasis	94.21%	pad
11	Cardiomegaly	90.55%	pad
12	Effusion	92.07%	pad
13	Pneumonia	92.99%	pad
14	Pneumothorax	100.00%	pad
15	Consolidation	99.09%	pad
16	Edema	98.78%	pad
17	Any	76.22%	pad

## CXP only

```
In [141]: print("Model trained on CXP")
model_CXP = EmbModel(emb_type="densenet121", feature_size_override=1024)
model_CXP.load_state_dict(torch.load("model/balanced/model-snapshot-emb-1024.pth"))
model_CXP.eval()
```

Model trained on CXP

Emb Dim:

1024

Manually setting output dim to 1024

1024

```
Out[141]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True)
```

```
In [142]: data_loader_CXP = datasets[get_dataset_index("cxp")]["val"]["loader"]
cxp_predictions, cxp_probabilities, cxp_labels = predict(model_CXP, da

print('predictions', cxp_predictions)
print('labels', cxp_labels)
print("Done")
```

Started prediction validation

Number of batches: 5

Batch number: 1 of 5

Batch number: 2 of 5

Batch number: 3 of 5

Batch number: 4 of 5

Batch number: 5 of 5

Done.

Prediction took: 0:00:07.939821

predictions [[False False False ... False False False]

[False False False ... False False False]

[False False False ... False False False]

...

[False False False ... False False False]

[False False False ... False False False]

[False False False ... False False False]]

labels [[0. 0. 0. ... 0. 0. 1.]

[0. 1. 0. ... 0. 1. 1.]

[0. 0. 0. ... 0. 1. 1.]

...

[0. 0. 0. ... 0. 0. 1.]

[0. 0. 0. ... 0. 0. 0.]

[0. 1. 0. ... 0. 0. 1.]]

Done

In [143]: `cxp_report, cxp_accuracy = calculate_stats(cxp_predictions, cxp_probab`

Overall Accuracy: 22.86%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.000000	0.000000	0.000000	60.0	cxp
Atelectasis	0.000000	0.000000	0.000000	97.0	cxp
Cardiomegaly	0.000000	0.000000	0.000000	74.0	cxp
Effusion	0.631016	0.460938	0.532731	256.0	cxp
Pneumonia	0.000000	0.000000	0.000000	12.0	cxp
Pneumothorax	0.000000	0.000000	0.000000	55.0	cxp
Consolidation	0.000000	0.000000	0.000000	43.0	cxp
Edema	0.432203	0.356643	0.390805	143.0	cxp
Any	0.836207	0.437923	0.574815	443.0	cxp
micro avg	0.663620	0.306847	0.419653	1183.0	cxp
macro avg	0.211047	0.139500	0.166483	1183.0	cxp
weighted avg	0.501931	0.306847	0.377775	1183.0	cxp
samples avg	0.248677	0.221534	0.228012	1183.0	cxp

	Label	Accuracy	source
0	No Finding	89.84%	cxp
1	Atelectasis	83.81%	cxp
2	Cardiomegaly	88.10%	cxp
3	Effusion	67.14%	cxp
4	Pneumonia	98.10%	cxp
5	Pneumothorax	91.27%	cxp
6	Consolidation	93.17%	cxp
7	Edema	74.76%	cxp
8	Any	54.44%	cxp

Worst Performing Label: Any with an accuracy of 54.44%

In [144]: `combined_accuracy, combined_report = combine_with_existing(combined_ac`

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic
5	0.000000	0.000000	0.000000	29.0	mimic
6	0.000000	0.000000	0.000000	36.0	mimic
7	0.156250	0.070423	0.097087	71.0	mimic
8	0.701389	0.274457	0.394531	368.0	mimic
9	0.561866	0.224110	0.320416	1236.0	mimic
10	0.274202	0.129513	0.171931	1236.0	mimic
11	0.474759	0.224110	0.299259	1236.0	mimic
12	0.248288	0.231509	0.235901	1236.0	mimic
13	0.656250	0.338710	0.446809	124.0	pad
14	0.000000	0.000000	0.000000	16.0	pad
15	0.000000	0.000000	0.000000	30.0	pad

16	0.200000	0.043478	0.071429	23.0	pad
17	0.000000	0.000000	0.000000	19.0	pad
18	0.000000	0.000000	0.000000	0.0	pad
19	0.000000	0.000000	0.000000	3.0	pad
20	0.000000	0.000000	0.000000	3.0	pad
21	0.416667	0.065789	0.113636	76.0	pad
22	0.533333	0.163265	0.250000	294.0	pad
23	0.141435	0.049775	0.070208	294.0	pad
24	0.400142	0.163265	0.223413	294.0	pad
25	0.135671	0.137195	0.136179	294.0	pad
26	0.000000	0.000000	0.000000	60.0	cxp
27	0.000000	0.000000	0.000000	97.0	cxp
28	0.000000	0.000000	0.000000	74.0	cxp
29	0.631016	0.460938	0.532731	256.0	cxp
30	0.000000	0.000000	0.000000	12.0	cxp
31	0.000000	0.000000	0.000000	55.0	cxp
32	0.000000	0.000000	0.000000	43.0	cxp
33	0.432203	0.356643	0.390805	143.0	cxp
34	0.836207	0.437923	0.574815	443.0	cxp
35	0.663620	0.306847	0.419653	1183.0	cxp
36	0.211047	0.139500	0.166483	1183.0	cxp
37	0.501931	0.306847	0.377775	1183.0	cxp
38	0.248677	0.221534	0.228012	1183.0	cxp

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic
5	Pneumothorax	96.08%	mimic
6	Consolidation	95.14%	mimic
7	Edema	87.43%	mimic
8	Any	58.11%	mimic
9	No Finding	68.29%	pad
10	Atelectasis	94.21%	pad
11	Cardiomegaly	90.55%	pad
12	Effusion	92.07%	pad
13	Pneumonia	92.99%	pad
14	Pneumothorax	100.00%	pad
15	Consolidation	99.09%	pad
16	Edema	98.78%	pad
17	Any	76.22%	pad
18	No Finding	89.84%	cxp
19	Atelectasis	83.81%	cxp
20	Cardiomegaly	88.10%	cxp
21	Effusion	67.14%	cxp
22	Pneumonia	98.10%	cxp
23	Pneumothorax	91.27%	cxp
24	Consolidation	93.17%	cxp
25	Edema	74.76%	cxp



## NIH only

```
Out[145]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
          (relu1): ReLU(inplace=True)
          (conv1): Conv2d(64, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (norm2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=T
```

```
In [146]: data_loader_NIH = datasets[get_dataset_index("nih")]["val"]["loader"]
nih_predictions, nih_probabilities, nih_labels = predict(model_NIH, da

print('predictions', nih_predictions)
print('labels', nih_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 3
Batch number: 1 of 3
Batch number: 2 of 3
Batch number: 3 of 3
Done.
Prediction took: 0:00:07.337274
predictions [[False False False ... False False False]
 [ True False False ... False False False]
 [ True False False ... False False False]
 ...
 [ True False False ... False False False]
 [False False False ... False False False]
 [ True False False ... False False False]]
labels [[0. 0. 0. ... 0. 0. 1.]
 [0. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 1.]
 [0. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]]
Done
```

In [147]: nih\_report, nih\_accuracy = calculate\_stats(nih\_predictions, nih\_probab

Overall Accuracy: 47.95%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.594697	0.813472	0.687090	193.0	nih
Atelectasis	0.000000	0.000000	0.000000	34.0	nih
Cardiomegaly	0.000000	0.000000	0.000000	9.0	nih
Effusion	0.600000	0.061224	0.111111	49.0	nih
Pneumonia	0.000000	0.000000	0.000000	2.0	nih
Pneumothorax	0.000000	0.000000	0.000000	6.0	nih
Consolidation	0.000000	0.000000	0.000000	23.0	nih
Edema	0.000000	0.000000	0.000000	9.0	nih
Any	0.500000	0.037736	0.070175	106.0	nih
micro avg	0.585714	0.380510	0.461322	431.0	nih
macro avg	0.188300	0.101381	0.096486	431.0	nih
weighted avg	0.457486	0.380510	0.337567	431.0	nih
samples avg	0.435160	0.439726	0.436438	431.0	nih

	Label	Accuracy	source
0	No Finding	60.82%	nih
1	Atelectasis	90.14%	nih
2	Cardiomegaly	97.53%	nih
3	Effusion	86.85%	nih
4	Pneumonia	99.45%	nih
5	Pneumothorax	98.36%	nih
6	Consolidation	93.42%	nih
7	Edema	97.53%	nih
8	Any	70.96%	nih

Worst Performing Label: No Finding with an accuracy of 60.82%

In [148]: combined\_accuracy, combined\_report = combine\_with\_existing(combined\_ac

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic
5	0.000000	0.000000	0.000000	29.0	mimic
6	0.000000	0.000000	0.000000	36.0	mimic
7	0.156250	0.070423	0.097087	71.0	mimic
8	0.701389	0.274457	0.394531	368.0	mimic
9	0.561866	0.224110	0.320416	1236.0	mimic
10	0.274202	0.129513	0.171931	1236.0	mimic
11	0.474759	0.224110	0.299259	1236.0	mimic
12	0.248288	0.231509	0.235901	1236.0	mimic
13	0.656250	0.338710	0.446809	124.0	pad
14	0.000000	0.000000	0.000000	16.0	pad
15	0.000000	0.000000	0.000000	30.0	pad

16	0.200000	0.043478	0.071429	23.0	pad
17	0.000000	0.000000	0.000000	19.0	pad
18	0.000000	0.000000	0.000000	0.0	pad
19	0.000000	0.000000	0.000000	3.0	pad
20	0.000000	0.000000	0.000000	3.0	pad
21	0.416667	0.065789	0.113636	76.0	pad
22	0.533333	0.163265	0.250000	294.0	pad
23	0.141435	0.049775	0.070208	294.0	pad
24	0.400142	0.163265	0.223413	294.0	pad
25	0.135671	0.137195	0.136179	294.0	pad
26	0.000000	0.000000	0.000000	60.0	cxp
27	0.000000	0.000000	0.000000	97.0	cxp
28	0.000000	0.000000	0.000000	74.0	cxp
29	0.631016	0.460938	0.532731	256.0	cxp
30	0.000000	0.000000	0.000000	12.0	cxp
31	0.000000	0.000000	0.000000	55.0	cxp
32	0.000000	0.000000	0.000000	43.0	cxp
33	0.432203	0.356643	0.390805	143.0	cxp
34	0.836207	0.437923	0.574815	443.0	cxp
35	0.663620	0.306847	0.419653	1183.0	cxp
36	0.211047	0.139500	0.166483	1183.0	cxp
37	0.501931	0.306847	0.377775	1183.0	cxp
38	0.248677	0.221534	0.228012	1183.0	cxp
39	0.594697	0.813472	0.687090	193.0	nih
40	0.000000	0.000000	0.000000	34.0	nih
41	0.000000	0.000000	0.000000	9.0	nih
42	0.600000	0.061224	0.111111	49.0	nih
43	0.000000	0.000000	0.000000	2.0	nih
44	0.000000	0.000000	0.000000	6.0	nih
45	0.000000	0.000000	0.000000	23.0	nih
46	0.000000	0.000000	0.000000	9.0	nih
47	0.500000	0.037736	0.070175	106.0	nih
48	0.585714	0.380510	0.461322	431.0	nih
49	0.188300	0.101381	0.096486	431.0	nih
50	0.457486	0.380510	0.337567	431.0	nih
51	0.435160	0.439726	0.436438	431.0	nih

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic
5	Pneumothorax	96.08%	mimic
6	Consolidation	95.14%	mimic
7	Edema	87.43%	mimic
8	Any	58.11%	mimic
9	No Finding	68.29%	pad
10	Atelectasis	94.21%	pad
11	Cardiomegaly	90.55%	pad
12	Effusion	92.07%	pad

13	Pneumonia	92.99%	pad
14	Pneumothorax	100.00%	pad
15	Consolidation	99.09%	pad
16	Edema	98.78%	pad
17	Any	76.22%	pad
18	No Finding	89.84%	cxp
19	Atelectasis	83.81%	cxp
20	Cardiomegaly	88.10%	cxp
21	Effusion	67.14%	cxp
22	Pneumonia	98.10%	cxp
23	Pneumothorax	91.27%	cxp
24	Consolidation	93.17%	cxp
25	Edema	74.76%	cxp
26	Any	54.44%	cxp
27	No Finding	60.82%	nih
28	Atelectasis	90.14%	nih
29	Cardiomegaly	97.53%	nih
30	Effusion	86.85%	nih
31	Pneumonia	99.45%	nih
32	Pneumothorax	98.36%	nih
33	Consolidation	93.42%	nih
34	Edema	97.53%	nih
35	Any	70.96%	nih

## CXP and NIH

```
In [153]: print("Loading a model trained on both CXP and NIH")

model_CXP_NIH = EmbModel(emb_type="densenet121", feature_size_override=
model_CXP_NIH.load_state_dict(torch.load("model/balanced/model-snapsho
model_CXP_NIH.eval()
```

Loading a model trained on both CXP and NIH

Emb Dim:

1024

Manually setting output dim to 1024

1024

```
Out[153]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=Tr
```

```
In [154]: data_loader_CXP_NIH = datasets[get_dataset_index("cyp_nih")]["val"]
cyp_nih_predictions, cyp_nih_probabilities, cyp_nih_labels = predict(m

print('predictions', cyp_nih_predictions)
print('labels', cyp_nih_labels)
print("Done")
```

Started prediction validation

Number of batches: 5

Batch number: 1 of 5

Batch number: 2 of 5

Batch number: 3 of 5

Batch number: 4 of 5

Batch number: 5 of 5

Done.

Prediction took: 0:00:05.094706

predictions [[False False False ... False False False]

[False False False ... False True True]

[False False False ... False False False]

...

[False False False ... False False True]

[False False False ... False False False]

[False False False ... False False False]]

labels [[0. 1. 0. ... 0. 1. 1.]

[0. 0. 1. ... 0. 0. 1.]

[0. 1. 1. ... 0. 0. 1.]

...

[1. 0. 0. ... 0. 0. 0.]

[1. 0. 0. ... 0. 0. 0.]

[0. 0. 0. ... 0. 0. 1.]]

Done

In [155]: `cxp_nih_report, cxp_nih_accuracy = calculate_stats(cxp_nih_predictions`

Overall Accuracy: 24.29%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.000000	0.000000	0.000000	60.0	cxp+nih
Atelectasis	0.000000	0.000000	0.000000	97.0	cxp+nih
Cardiomegaly	0.000000	0.000000	0.000000	74.0	cxp+nih
Effusion	0.600000	0.539062	0.567901	256.0	cxp+nih
Pneumonia	0.000000	0.000000	0.000000	12.0	cxp+nih
Pneumothorax	0.000000	0.000000	0.000000	55.0	cxp+nih
Consolidation	0.000000	0.000000	0.000000	43.0	cxp+nih
Edema	0.433962	0.321678	0.369478	143.0	cxp+nih
Any	0.843750	0.487585	0.618026	443.0	cxp+nih
micro avg	0.662252	0.338123	0.447678	1183.0	cxp+nih
macro avg	0.208635	0.149814	0.172823	1183.0	cxp+nih
weighted avg	0.498257	0.338123	0.398989	1183.0	cxp+nih
samples avg	0.270503	0.247460	0.251977	1183.0	cxp+nih

	Label Accuracy	source
0 No Finding	90.32%	cxp+nih
1 Atelectasis	84.44%	cxp+nih
2 Cardiomegaly	87.14%	cxp+nih

In [156]: `combined_accuracy, combined_report = combine_with_existing(combined_ac`

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic
..	...	...	...	...	...
60	0.843750	0.487585	0.618026	443.0	cxp+nih
61	0.662252	0.338123	0.447678	1183.0	cxp+nih
62	0.208635	0.149814	0.172823	1183.0	cxp+nih
63	0.498257	0.338123	0.398989	1183.0	cxp+nih
64	0.270503	0.247460	0.251977	1183.0	cxp+nih

[65 rows x 5 columns]

	Label Accuracy	source
0 No Finding	74.19%	mimic
1 Atelectasis	75.27%	mimic
2 Cardiomegaly	82.30%	mimic
3 Effusion	79.05%	mimic
4 Pneumonia	93.65%	mimic
5 Pneumothorax	96.08%	mimic
6 Consolidation	95.14%	mimic
7 Edema	87.43%	mimic
8 Any	58.11%	mimic
9 No Finding	68.29%	pad



10	Atelectasis	94.21%	pad
11	Cardiomegaly	90.55%	pad
12	Effusion	92.07%	pad
13	Pneumonia	92.99%	pad
14	Pneumothorax	100.00%	pad
15	Consolidation	99.09%	pad
16	Edema	98.78%	pad
17	Any	76.22%	pad
18	No Finding	89.84%	cxp
19	Atelectasis	83.81%	cxp
20	Cardiomegaly	88.10%	cxp
21	Effusion	67.14%	cxp
22	Pneumonia	98.10%	cxp
23	Pneumothorax	91.27%	cxp
24	Consolidation	93.17%	cxp
25	Edema	74.76%	cxp
26	Any	54.44%	cxp
27	No Finding	60.82%	nih
28	Atelectasis	90.14%	nih
29	Cardiomegaly	97.53%	nih
30	Effusion	86.85%	nih
31	Pneumonia	99.45%	nih
32	Pneumothorax	98.36%	nih
33	Consolidation	93.42%	nih
34	Edema	97.53%	nih
35	Any	70.96%	nih
36	No Finding	90.32%	cxp+nih
37	Atelectasis	84.44%	cxp+nih
38	Cardiomegaly	87.14%	cxp+nih
39	Effusion	66.67%	cxp+nih
40	Pneumonia	98.10%	cxp+nih
41	Pneumothorax	90.79%	cxp+nih
42	Consolidation	93.17%	cxp+nih
43	Edema	75.08%	cxp+nih
44	Any	57.62%	cxp+nih

```
In [178]: # nih_cxp_report, nih_cxp_accuracy = calculate_stats(cxp_nih_prediction)
# combined_accuracy, combined_report = combine_with_existing(combined_report,
```

## CXP and PAD

In [157]: `print("Loading a model trained on both CXP and PAD")`

```
model_CXP_PAD = EmbModel(emb_type="densenet121", feature_size_override
model_CXP_PAD.load_state_dict(torch.load("model/balanced/model-snapsho
model_CXP_PAD.eval()
```

Loading a model trained on both CXP and PAD

Emb Dim:

1024

Manually setting output dim to 1024

1024

Out[157]:

```
EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=Tr
```

```
In [158]: data_loader_CXP_PAD = datasets[get_dataset_index("cxp_pad")]["val"]
cxp_pad_predictions, cxp_pad_probabilities, cxp_pad_labels = predict(m

print('predictions', cxp_pad_predictions)
print('labels', cxp_pad_labels)
print("Done")
```

Started prediction validation

Number of batches: 5

Batch number: 1 of 5

Batch number: 2 of 5

Batch number: 3 of 5

Batch number: 4 of 5

Batch number: 5 of 5

Done.

Prediction took: 0:00:05.910238

predictions [[False False False ... False False True]

[False False False ... False True True]

[False False False ... False False False]

...

[False False False ... False False False]

[False False False ... False False False]

[False False False ... False False True]]

labels [[0. 0. 0. ... 0. 1. 1.]

[0. 0. 0. ... 0. 1. 1.]

[0. 0. 0. ... 0. 0. 1.]

...

[1. 0. 0. ... 0. 0. 0.]

[0. 0. 0. ... 1. 0. 1.]

[0. 0. 0. ... 0. 0. 0.]]

Done

In [159]: `cxp_pad_report, cxp_pad_accuracy = calculate_stats(cxp_pad_predictions`

Overall Accuracy: 24.44%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	1.000000	0.016667	0.032787	60.0	cxp+pad
Atelectasis	0.000000	0.000000	0.000000	97.0	cxp+pad
Cardiomegaly	0.000000	0.000000	0.000000	74.0	cxp+pad
Effusion	0.630682	0.433594	0.513889	256.0	cxp+pad
Pneumonia	0.000000	0.000000	0.000000	12.0	cxp+pad
Pneumothorax	0.000000	0.000000	0.000000	55.0	cxp+pad
Consolidation	0.000000	0.000000	0.000000	43.0	cxp+pad
Edema	0.465347	0.328671	0.385246	143.0	cxp+pad
Any	0.842593	0.410835	0.552352	443.0	cxp+pad
micro avg	0.687500	0.288250	0.406194	1183.0	cxp+pad
macro avg	0.326513	0.132196	0.164919	1183.0	cxp+pad
weighted avg	0.558975	0.288250	0.366276	1183.0	cxp+pad
samples avg	0.237831	0.208757	0.217861	1183.0	cxp+pad

	Label	Accuracy	source
0	No Finding	90.63%	cxp+pad
1	Atelectasis	84.44%	cxp+pad
2	Cardiomegaly	88.10%	cxp+pad
3	Effusion	66.67%	cxp+pad
4	Pneumonia	98.10%	cxp+pad
5	Pneumothorax	91.27%	cxp+pad
6	Consolidation	93.17%	cxp+pad
7	Edema	76.19%	cxp+pad
8	Any	53.17%	cxp+pad

Worst Performing Label: Any with an accuracy of 53.17%

In [52]: `combined_accuracy, combined_report = combine_with_existing(combined_ac`

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic
..	...	...	...	...	...
73	0.841410	0.368015	0.512064	519.0	cxp+pad
74	0.679702	0.247123	0.362463	1477.0	cxp+pad
75	0.342983	0.131250	0.171620	1477.0	cxp+pad
76	0.581334	0.247123	0.328035	1477.0	cxp+pad
77	0.174322	0.156072	0.161537	1477.0	cxp+pad

[78 rows x 5 columns]

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic

2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic
5	Pneumothorax	96.08%	mimic
6	Consolidation	95.14%	mimic
7	Edema	87.43%	mimic
8	Any	58.11%	mimic
9	No Finding	68.29%	pad
10	Atelectasis	94.21%	pad
11	Cardiomegaly	90.55%	pad
12	Effusion	92.07%	pad
13	Pneumonia	92.99%	pad
14	Pneumothorax	100.00%	pad
15	Consolidation	99.09%	pad
16	Edema	98.78%	pad
17	Any	76.22%	pad
18	No Finding	89.84%	cxp
19	Atelectasis	83.81%	cxp
20	Cardiomegaly	88.10%	cxp
21	Effusion	67.14%	cxp
22	Pneumonia	98.10%	cxp
23	Pneumothorax	91.27%	cxp
24	Consolidation	93.17%	cxp
25	Edema	74.76%	cxp
26	Any	54.44%	cxp
27	No Finding	61.10%	nih
28	Atelectasis	90.14%	nih
29	Cardiomegaly	97.53%	nih
30	Effusion	86.85%	nih
31	Pneumonia	99.45%	nih
32	Pneumothorax	98.36%	nih
33	Consolidation	93.42%	nih
34	Edema	97.53%	nih
35	Any	70.96%	nih
36	No Finding	79.70%	cxp+nih
37	Atelectasis	86.73%	cxp+nih
38	Cardiomegaly	90.95%	cxp+nih
39	Effusion	73.57%	cxp+nih
40	Pneumonia	98.59%	cxp+nih
41	Pneumothorax	93.57%	cxp+nih
42	Consolidation	93.27%	cxp+nih
43	Edema	83.32%	cxp+nih
44	Any	62.51%	cxp+nih
45	No Finding	81.73%	cxp+pad
46	Atelectasis	88.00%	cxp+pad
47	Cardiomegaly	89.04%	cxp+pad
48	Effusion	75.26%	cxp+pad
49	Pneumonia	96.45%	cxp+pad
50	Pneumothorax	94.26%	cxp+pad
51	Consolidation	95.20%	cxp+pad

52	Edema	84.03%	cxp+pad
53	Any	62.00%	cxp+pad

```
In [181]: # pad_cxp_report, pad_cxp_accuracy = calculate_stats(cxp_pad_prediction)
# combined_accuracy, combined_report = combine_with_existing(combined_report,
```

## NIH and PAD

```
In [160]: print("Loading a model trained on both NIH and PAD")

model_NIH_PAD = EmbModel(emb_type="densenet121", feature_size_override=1024)
model_NIH_PAD.load_state_dict(torch.load("model/balanced/model-snapshots/model_NIH_PAD.pth"))
model_NIH_PAD.eval()
```

Loading a model trained on both NIH and PAD

Emb Dim:  
1024  
Manually setting output dim to 1024  
1024

```
Out[160]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
```

```
In [161]: data_loader_NIH_PAD = datasets[get_dataset_index("nih_pad")]["val"]
nih_pad_predictions, nih_pad_probabilities, nih_pad_labels = predict(m

print('predictions', nih_pad_predictions)
print('labels', nih_pad_labels)
print("Done")
```

Started prediction validation

Number of batches: 3

Batch number: 1 of 3

Batch number: 2 of 3

Batch number: 3 of 3

Done.

Prediction took: 0:00:07.671303

predictions [[False False False ... False False False]

[False False False ... False False False]

[ True False False ... False False False]

...

[False False False ... False False True]

[ True False False ... False False False]

[ True False False ... False False False]]

labels [[0. 0. 0. ... 0. 0. 0.]

[0. 1. 0. ... 0. 0. 1.]

[1. 0. 0. ... 0. 0. 0.]

...

[1. 0. 0. ... 0. 0. 0.]

[1. 0. 0. ... 0. 0. 0.]

[1. 0. 0. ... 0. 0. 0.]]

Done

In [162]: nih\_pad\_report, nih\_pad\_accuracy = calculate\_stats(nih\_pad\_predictions)

Overall Accuracy: 45.48%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.605809	0.756477	0.672811	193.0	nih+pad
Atelectasis	0.000000	0.000000	0.000000	34.0	nih+pad
Cardiomegaly	0.000000	0.000000	0.000000	9.0	nih+pad
Effusion	0.500000	0.020408	0.039216	49.0	nih+pad
Pneumonia	0.000000	0.000000	0.000000	2.0	nih+pad
Pneumothorax	0.000000	0.000000	0.000000	6.0	nih+pad
Consolidation	0.000000	0.000000	0.000000	23.0	nih+pad
Edema	0.000000	0.000000	0.000000	9.0	nih+pad
Any	0.666667	0.018868	0.036697	106.0	nih+pad
micro avg	0.603239	0.345708	0.439528	431.0	nih+pad
macro avg	0.196942	0.088417	0.083192	431.0	nih+pad
weighted avg	0.492083	0.345708	0.314766	431.0	nih+pad
samples avg	0.403653	0.404110	0.403836	431.0	nih+pad

	Label	Accuracy	source
0	No Finding	61.10%	nih+pad
1	Atelectasis	90.68%	nih+pad
2	Cardiomegaly	97.53%	nih+pad
3	Effusion	86.58%	nih+pad
4	Pneumonia	99.45%	nih+pad
5	Pneumothorax	98.36%	nih+pad
6	Consolidation	93.42%	nih+pad
7	Edema	97.53%	nih+pad
8	Any	71.23%	nih+pad

Worst Performing Label: No Finding with an accuracy of 61.10%

In [163]: combined\_accuracy, combined\_report = combine\_with\_existing(combined\_ac

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic
..	...	...	...	...	...
73	0.666667	0.018868	0.036697	106.0	nih+pad
74	0.603239	0.345708	0.439528	431.0	nih+pad
75	0.196942	0.088417	0.083192	431.0	nih+pad
76	0.492083	0.345708	0.314766	431.0	nih+pad
77	0.403653	0.404110	0.403836	431.0	nih+pad

[78 rows x 5 columns]

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic



2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic
5	Pneumothorax	96.08%	mimic
6	Consolidation	95.14%	mimic
7	Edema	87.43%	mimic
8	Any	58.11%	mimic
9	No Finding	68.29%	pad
10	Atelectasis	94.21%	pad
11	Cardiomegaly	90.55%	pad
12	Effusion	92.07%	pad
13	Pneumonia	92.99%	pad
14	Pneumothorax	100.00%	pad
15	Consolidation	99.09%	pad
16	Edema	98.78%	pad
17	Any	76.22%	pad
18	No Finding	89.84%	cxp
19	Atelectasis	83.81%	cxp
20	Cardiomegaly	88.10%	cxp
21	Effusion	67.14%	cxp
22	Pneumonia	98.10%	cxp
23	Pneumothorax	91.27%	cxp
24	Consolidation	93.17%	cxp
25	Edema	74.76%	cxp
26	Any	54.44%	cxp
27	No Finding	60.82%	nih
28	Atelectasis	90.14%	nih
29	Cardiomegaly	97.53%	nih
30	Effusion	86.85%	nih
31	Pneumonia	99.45%	nih
32	Pneumothorax	98.36%	nih
33	Consolidation	93.42%	nih
34	Edema	97.53%	nih
35	Any	70.96%	nih
36	No Finding	90.32%	cxp+nih
37	Atelectasis	84.44%	cxp+nih
38	Cardiomegaly	87.14%	cxp+nih
39	Effusion	66.67%	cxp+nih
40	Pneumonia	98.10%	cxp+nih
41	Pneumothorax	90.79%	cxp+nih
42	Consolidation	93.17%	cxp+nih
43	Edema	75.08%	cxp+nih
44	Any	57.62%	cxp+nih
45	No Finding	61.10%	nih+pad
46	Atelectasis	90.68%	nih+pad
47	Cardiomegaly	97.53%	nih+pad
48	Effusion	86.58%	nih+pad
49	Pneumonia	99.45%	nih+pad
50	Pneumothorax	98.36%	nih+pad
51	Consolidation	93.42%	nih+pad

52	Edema	97.53%	nih+pad
53	Any	71.23%	nih+pad

```
In [184]: # pad_nih_report, pad_nih_accuracy = calculate_stats(nih_pad_prediction)
# combined_accuracy, combined_report = combine_with_existing(combined_report,
```

## MIMIC and PAD

```
In [164]: print("Loading a model trained on both MIMIC and PAD")

model_MIMIC_PAD = EmbModel(emb_type="densenet121", feature_size_override=1024)
model_MIMIC_PAD.load_state_dict(torch.load("model/balanced/model-snaps"))
model_MIMIC_PAD.eval()
```

Loading a model trained on both MIMIC and PAD

Emb Dim:  
1024  
Manually setting output dim to 1024  
1024

```
Out[164]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
```

```
In [165]: data_loader_MIMIC_PAD = datasets[get_dataset_index("mimic_pad")]["val"]
mimic_pad_predictions, mimic_pad_probabilities, mimic_pad_labels = pre

print('predictions', mimic_pad_predictions)
print('labels', mimic_pad_labels)
print("Done")
```

Started prediction validation

Number of batches: 6

Batch number: 1 of 6

Batch number: 2 of 6

Batch number: 3 of 6

Batch number: 4 of 6

Batch number: 5 of 6

Batch number: 6 of 6

Done.

Prediction took: 0:00:37.308962

predictions [[ True False False ... False False False]

[ True False False ... False False False]

[ True False False ... False False False]

...

[False False True ... False False True]

[ True False False ... False False False]

[False False False ... False False True]]

labels [[1. 0. 0. ... 0. 0. 0.]

[1. 0. 0. ... 0. 0. 0.]

[0. 0. 0. ... 0. 0. 0.]

...

[1. 0. 0. ... 0. 0. 0.]

[0. 0. 0. ... 0. 0. 1.]

[0. 0. 1. ... 0. 1. 1.]]

Done

In [166]: `mimic_pad_report, mimic_pad_accuracy = calculate_stats(mimic_pad_predi`

Overall Accuracy: 31.35%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.681223	0.602317	0.639344	259.0	mimic+pad
Atelectasis	0.000000	0.000000	0.000000	159.0	mimic+pad
Cardiomegaly	0.250000	0.049180	0.082192	122.0	mimic+pad
Effusion	0.513274	0.386667	0.441065	150.0	mimic+pad
Pneumonia	0.000000	0.000000	0.000000	42.0	mimic+pad
Pneumothorax	0.000000	0.000000	0.000000	29.0	mimic+pad
Consolidation	0.000000	0.000000	0.000000	36.0	mimic+pad
Edema	0.294118	0.070423	0.113636	71.0	mimic+pad
Any	0.785185	0.288043	0.421471	368.0	mimic+pad
micro avg	0.618692	0.267799	0.373800	1236.0	mimic+pad
macro avg	0.280422	0.155181	0.188634	1236.0	mimic+pad
weighted avg	0.480387	0.267799	0.327627	1236.0	mimic+pad
samples avg	0.316892	0.294234	0.300727	1236.0	mimic+pad

	Label	Accuracy	source
0	No Finding	76.22%	mimic+pad
1	Atelectasis	76.22%	mimic+pad
2	Cardiomegaly	81.89%	mimic+pad
3	Effusion	80.14%	mimic+pad
4	Pneumonia	94.32%	mimic+pad
5	Pneumothorax	96.08%	mimic+pad
6	Consolidation	95.14%	mimic+pad
7	Edema	89.46%	mimic+pad
8	Any	60.68%	mimic+pad

Worst Performing Label: Any with an accuracy of 60.68%

In [167]: `combined_accuracy, combined_report = combine_with_existing(combined_ac`

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic
..	...	...	...	...	...
86	0.785185	0.288043	0.421471	368.0	mimic+pad
87	0.618692	0.267799	0.373800	1236.0	mimic+pad
88	0.280422	0.155181	0.188634	1236.0	mimic+pad
89	0.480387	0.267799	0.327627	1236.0	mimic+pad
90	0.316892	0.294234	0.300727	1236.0	mimic+pad

[91 rows x 5 columns]

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic
..	...	...	...
58	Pneumonia	94.32%	mimic+pad
59	Pneumothorax	96.08%	mimic+pad
60	Consolidation	95.14%	mimic+pad
61	Edema	89.46%	mimic+pad
62	Any	60.68%	mimic+pad

[63 rows x 3 columns]

In [187]: `# pad_mimic_report, pad_mimic_accuracy = calculate_stats(mimic_pad_pre`  
`# combined_accuracy, combined_report = combine_with_existing(combined_`

## MIMIC and NIH

In [168]: `print("Loading a model trained on both MIMIC and NIH")`

```
model_MIMIC_NIH = EmbModel(emb_type="densenet121", feature_size_overri
model_MIMIC_NIH.load_state_dict(torch.load("model/balanced/model-snaps
model_MIMIC_NIH.eval())
```

Loading a model trained on both MIMIC and NIH

Emb Dim:

1024

Manually setting output dim to 1024

1024

Out[168]:

```
EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=Tr
```

```
In [169]: data_loader_MIMIC_NIH = datasets[get_dataset_index("mimic_nih")]["val"]
mimic_nih_predictions, mimic_nih_probabilities, mimic_nih_labels = pre

print('predictions', mimic_nih_predictions)
print('labels', mimic_nih_labels)
print("Done")
```

Started prediction validation

Number of batches: 6

Batch number: 1 of 6

Batch number: 2 of 6

Batch number: 3 of 6

Batch number: 4 of 6

Batch number: 5 of 6

Batch number: 6 of 6

Done.

Prediction took: 0:00:37.731515

predictions [[False True False ... False False True]

[False False False ... False False False]

[ True False False ... False False True]

...

[False False False ... False False False]

[ True False False ... False False False]

[False True False ... True False True]]

labels [[0. 1. 0. ... 0. 0. 1.]

[0. 0. 0. ... 1. 0. 1.]

[1. 0. 0. ... 0. 0. 0.]

...

[0. 1. 0. ... 0. 0. 1.]

[1. 0. 0. ... 0. 0. 0.]

[0. 0. 1. ... 0. 1. 1.]]

Done

```
In [170]: mimc_nih_report, mimc_nih_accuracy = calculate_stats(mimc_nih_predic
```

Overall Accuracy: 26.49%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.743243	0.424710	0.540541	259.0	mimic+nih
Atelectasis	0.314815	0.106918	0.159624	159.0	mimic+nih
Cardiomegaly	0.275862	0.065574	0.105960	122.0	mimic+nih
Effusion	0.384615	0.533333	0.446927	150.0	mimic+nih
Pneumonia	0.000000	0.000000	0.000000	42.0	mimic+nih
Pneumothorax	0.000000	0.000000	0.000000	29.0	mimic+nih
Consolidation	0.000000	0.000000	0.000000	36.0	mimic+nih
Edema	0.200000	0.056338	0.087912	71.0	mimic+nih
Any	0.727273	0.456522	0.560935	368.0	mimic+nih
micro avg	0.558442	0.313107	0.401244	1236.0	mimic+nih
macro avg	0.293979	0.182600	0.211322	1236.0	mimic+nih
weighted avg	0.498171	0.313107	0.370560	1236.0	mimic+nih
samples avg	0.307275	0.286059	0.290798	1236.0	mimic+nih

	Label	Accuracy	source
0	No Finding	74.73%	mimic+nih
1	Atelectasis	75.81%	mimic+nih
2	Cardiomegaly	81.76%	mimic+nih
3	Effusion	73.24%	mimic+nih
4	Pneumonia	94.32%	mimic+nih
5	Pneumothorax	95.95%	mimic+nih
6	Consolidation	94.86%	mimic+nih
7	Edema	88.78%	mimic+nih
8	Any	64.46%	mimic+nih

Worst Performing Label: Any with an accuracy of 64.46%



```
In [171]: combined_accuracy, combined_report = combine_with_existing(combined_acc
```

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic
..	...	...	...	...	...
99	0.727273	0.456522	0.560935	368.0	mimic+nih
100	0.558442	0.313107	0.401244	1236.0	mimic+nih
101	0.293979	0.182600	0.211322	1236.0	mimic+nih
102	0.498171	0.313107	0.370560	1236.0	mimic+nih
103	0.307275	0.286059	0.290798	1236.0	mimic+nih

[104 rows x 5 columns]

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic
..	...	...	...
67	Pneumonia	94.32%	mimic+nih
68	Pneumothorax	95.95%	mimic+nih
69	Consolidation	94.86%	mimic+nih
70	Edema	88.78%	mimic+nih
71	Any	64.46%	mimic+nih

[72 rows x 3 columns]

```
In [172]: # nih_mimic_report, nih_mimic_accuracy = calculate_stats(mimic_nih_pre
# combined_accuracy, combined_report = combine_with_existing(combined_a
```

## MIMIC and CXP

In [173]: `print("Loading a model trained on both MIMIC and CXP")`

```
model_MIMIC_CXP = EmbModel(emb_type="densenet121", feature_size_overri
model_MIMIC_CXP.load_state_dict(torch.load("model/balanced/model-snaps
model_MIMIC_CXP.eval())
```

Loading a model trained on both MIMIC and CXP

Emb Dim:

1024

Manually setting output dim to 1024

1024

Out[173]:

```
EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=Tr
```

```
In [174]: data_loader_MIMIC_CXP = datasets[get_dataset_index("mimic_cxp")]["val"]
mimic_cxp_predictions, mimic_cxp_probabilities, mimic_cxp_labels = pre

print('predictions', mimic_cxp_predictions)
print('labels', mimic_cxp_labels)
print("Done")
```

Started prediction validation

Number of batches: 6

Batch number: 1 of 6

Batch number: 2 of 6

Batch number: 3 of 6

Batch number: 4 of 6

Batch number: 5 of 6

Batch number: 6 of 6

Done.

Prediction took: 0:00:35.213949

predictions [[False False False ... False False True]

[False False False ... False False True]

[False False False ... False False True]

...

[ True False False ... False False False]

[False False False ... False False False]

[ True False False ... False False False]]

labels [[1. 0. 0. ... 0. 0. 0.]

[0. 1. 0. ... 1. 0. 1.]

[1. 0. 0. ... 0. 0. 0.]

...

[1. 0. 0. ... 0. 0. 0.]

[0. 0. 0. ... 0. 0. 1.]

[0. 0. 0. ... 0. 0. 1.]]

Done

In [176]: `mimic_cxp_report, mimic_cxp_accuracy = calculate_stats(mimic_cxp_predi`

Overall Accuracy: 25.14%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.741259	0.409266	0.527363	259.0	mimic+cxp
Atelectasis	0.200000	0.062893	0.095694	159.0	mimic+cxp
Cardiomegaly	0.214286	0.024590	0.044118	122.0	mimic+cxp
Effusion	0.450382	0.393333	0.419929	150.0	mimic+cxp
Pneumonia	0.333333	0.023810	0.044444	42.0	mimic+cxp
Pneumothorax	0.000000	0.000000	0.000000	29.0	mimic+cxp
Consolidation	0.000000	0.000000	0.000000	36.0	mimic+cxp
Edema	0.176471	0.042254	0.068182	71.0	mimic+cxp
Any	0.767857	0.350543	0.481343	368.0	mimic+cxp
micro avg	0.590133	0.251618	0.352808	1236.0	mimic+cxp
macro avg	0.320399	0.145188	0.186786	1236.0	mimic+cxp
weighted avg	0.506947	0.251618	0.326874	1236.0	mimic+cxp
samples avg	0.265518	0.244730	0.250898	1236.0	mimic+cxp

	Label	Accuracy	source
0	No Finding	74.32%	mimic+cxp
1	Atelectasis	74.46%	mimic+cxp
2	Cardiomegaly	82.43%	mimic+cxp
3	Effusion	77.97%	mimic+cxp
4	Pneumonia	94.19%	mimic+cxp
5	Pneumothorax	96.08%	mimic+cxp
6	Consolidation	95.00%	mimic+cxp
7	Edema	88.92%	mimic+cxp
8	Any	62.43%	mimic+cxp

Worst Performing Label: Any with an accuracy of 62.43%

In [177]: `combined_accuracy, combined_report = combine_with_existing(combined_ac`

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic
..	...	...	...	...	...
112	0.767857	0.350543	0.481343	368.0	mimic+cxp
113	0.590133	0.251618	0.352808	1236.0	mimic+cxp
114	0.320399	0.145188	0.186786	1236.0	mimic+cxp
115	0.506947	0.251618	0.326874	1236.0	mimic+cxp
116	0.265518	0.244730	0.250898	1236.0	mimic+cxp

[117 rows x 5 columns]

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic
..	...	...	...
76	Pneumonia	94.19%	mimic+cxp
77	Pneumothorax	96.08%	mimic+cxp
78	Consolidation	95.00%	mimic+cxp
79	Edema	88.92%	mimic+cxp
80	Any	62.43%	mimic+cxp

[81 rows x 3 columns]

In [192]: `# cxp_mimic_report, cxp_mimic_accuracy = calculate_stats(mimic_cxp_pre`  
`# combined_accuracy, combined_report = combine_with_existing(combined_`

**All four datasets together: MIMIC, CXP, NIH and PAD**

```
In [178]: print("Loading a model trained on all four datasets: MIMIC, CXP, NIH a
model_CXP_MIMIC_NIH_PAD = EmbModel(emb_type="densenet121", feature_size=1024)
model_CXP_MIMIC_NIH_PAD.load_state_dict(torch.load("model/balanced/model_CXP_MIMIC_NIH_PAD.pth"))
model_CXP_MIMIC_NIH_PAD.eval()

(encoder): Sequential(
  (0): Sequential(
    (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
    (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    (denseblock1): _DenseBlock(
      (denselayer1): _DenseLayer(
        (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu1): ReLU(inplace=True)
        (conv1): Conv2d(64, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (norm2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu2): ReLU(inplace=True)
        (conv2): Conv2d(128, 32, kernel_size=(3, 3), stride=(1, 1),
```

```
In [179]: data_loader_CXP_MIMIC_NIH_PAD = datasets[get_dataset_index("cxp_mimic_
cxp_mimic_nih_pad_predictions, cxp_mimic_nih_pad_probabilities, cxp_mi

print('predictions', cxp_mimic_nih_pad_predictions)
print('labels', cxp_mimic_nih_pad_labels)
print("Done")
```

Started prediction validation

Number of batches: 17

Batch number: 1 of 17

Batch number: 2 of 17

Batch number: 3 of 17

Batch number: 4 of 17

Batch number: 5 of 17

Batch number: 6 of 17

Batch number: 7 of 17

Batch number: 8 of 17

Batch number: 9 of 17

Batch number: 10 of 17

Batch number: 11 of 17

Batch number: 12 of 17

Batch number: 13 of 17

Batch number: 14 of 17

Batch number: 15 of 17

Batch number: 16 of 17

Batch number: 17 of 17

Done.

Prediction took: 0:00:51.106543

predictions [[False False False ... False False False]

[False False False ... False False True]

[ True False False ... False False False]

...

[False False False ... False False False]

[ True False False ... False False False]

[False False False ... False False False]]

labels [[0. 0. 0. ... 0. 0. 0.]

[0. 0. 0. ... 0. 0. 0.]

[1. 0. 0. ... 0. 0. 0.]

...

[0. 0. 0. ... 0. 0. 1.]

[0. 0. 0. ... 0. 0. 1.]

[0. 0. 1. ... 0. 0. 1.]]

Done

In [180]: `cxp_mimic_nih_pad_report, cxp_mimic_nih_pad_accuracy = calculate_stats`

Overall Accuracy: 33.11%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.639155	0.523585	0.575627	636.0	cxp+mimic+nih+pad
Atelectasis	0.297297	0.035948	0.064140	306.0	cxp+mimic+nih+pad
Cardiomegaly	0.285714	0.008511	0.016529	235.0	cxp+mimic+nih+pad
Effusion	0.577558	0.366109	0.448143	478.0	cxp+mimic+nih+pad
Pneumonia	0.142857	0.040000	0.062500	75.0	cxp+mimic+nih+pad
Pneumothorax	0.000000	0.000000	0.000000	90.0	cxp+mimic+nih+pad
Consolidation	0.000000	0.000000	0.000000	105.0	cxp+mimic+nih+pad
Edema	0.398551	0.243363	0.302198	226.0	cxp+mimic+nih+pad
Any	0.798969	0.312185	0.448950	993.0	cxp+mimic+nih+pad
micro avg	0.628269	0.282761	0.389998	3144.0	cxp+mimic+nih+pad
macro avg	0.348900	0.169967	0.213121	3144.0	cxp+mimic+nih+pad
weighted avg	0.551798	0.282761	0.357065	3144.0	cxp+mimic+nih+pad
samples avg	0.276903	0.262579	0.266328	3144.0	cxp+mimic+nih+pad

	Label	Accuracy	source
0	No Finding	76.20%	cxp+mimic+nih+pad
1	Atelectasis	84.44%	cxp+mimic+nih+pad
2	Cardiomegaly	88.46%	cxp+mimic+nih+pad
3	Effusion	79.11%	cxp+mimic+nih+pad
4	Pneumonia	95.64%	cxp+mimic+nih+pad
5	Pneumothorax	95.64%	cxp+mimic+nih+pad
6	Consolidation	94.91%	cxp+mimic+nih+pad
7	Edema	87.69%	cxp+mimic+nih+pad
8	Any	63.11%	cxp+mimic+nih+pad

Worst Performing Label: Any with an accuracy of 63.11%



In [181]: combined\_accuracy, combined\_report = combine\_with\_existing(combined\_ac

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic
..	...	...	...	...	...
125	0.798969	0.312185	0.448950	993.0	cxp+mimic+nih+pad
126	0.628269	0.282761	0.389998	3144.0	cxp+mimic+nih+pad
127	0.348900	0.169967	0.213121	3144.0	cxp+mimic+nih+pad
128	0.551798	0.282761	0.357065	3144.0	cxp+mimic+nih+pad
129	0.276903	0.262579	0.266328	3144.0	cxp+mimic+nih+pad

[130 rows x 5 columns]

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic
..	...	...	...
85	Pneumonia	95.64%	cxp+mimic+nih+pad
86	Pneumothorax	95.64%	cxp+mimic+nih+pad
87	Consolidation	94.91%	cxp+mimic+nih+pad
88	Edema	87.69%	cxp+mimic+nih+pad
89	Any	63.11%	cxp+mimic+nih+pad

[90 rows x 3 columns]

In [ ]:

## Saving all data to data frame and to file

```
In [182]: print("Saving the evaluation stats to CSV")
combined_df_now = datetime.now()
combined_df_dt_string = combined_df_now.strftime("%d-%m-%Y-%H-%M-%S")
combined_accuracy.to_csv(f"stats/{combined_df_dt_string}_combined_accu
combined_report.to_csv(f"stats/{combined_df_dt_string}_combined_accura
print("Done.")

print("File names:")
print(f"stats/{combined_df_dt_string}_combined_accuracy_breakdown.csv")
print(f"stats/{combined_df_dt_string}_combined_accuracy_totals.csv")
```

Saving the evaluation stats to CSV

Done.

File names:

stats/05-05-2024-23-11-28\_combined\_accuracy\_breakdown.csv

stats/05-05-2024-23-11-28\_combined\_accuracy\_totals.csv

## Visualizing data

In [183]:

```

timestamp_to_load = "04-05-2024-23-34-18"
two_dataset_eval_timestamp = "04-05-2024-23-34-18"
one_dataset_eval_timestamp = "05-05-2024-16-23-37"
balanced_dataset_timestamp = "05-05-2024-22-03-40"
balanced_single_dataset_timestamp = "05-05-2024-23-11-28"

two_dataset_accuracy_breakdown_path = f"stats/{two_dataset_eval_timestamp}"
two_dataset_accuracy_totals_path = f"stats/{two_dataset_eval_timestamp}"

one_dataset_accuracy_breakdown_path = f"stats/{one_dataset_eval_timestamp}"
one_dataset_accuracy_totals_path = f"stats/{one_dataset_eval_timestamp}"

balanced_dataset_accuracy_breakdown_path = f"stats/{balanced_dataset_timestamp}"
balanced_dataset_accuracy_totals_path = f"stats/{balanced_dataset_timestamp}"

balanced_single_dataset_accuracy_breakdown_path = f"stats/{balanced_single_dataset_timestamp}"
balanced_single_dataset_accuracy_totals_path = f"stats/{balanced_single_dataset_timestamp}"

# Loading the DataFrames
two_dataset_combined_accuracy = pd.read_csv(two_dataset_accuracy_breakdown_path)
two_dataset_combined_report = pd.read_csv(two_dataset_accuracy_totals_path)

one_dataset_combined_accuracy = pd.read_csv(one_dataset_accuracy_breakdown_path)
one_dataset_combined_report = pd.read_csv(one_dataset_accuracy_totals_path)

balanced_dataset_combined_accuracy = pd.read_csv(balanced_dataset_accuracy_breakdown_path)
balanced_dataset_combined_report = pd.read_csv(balanced_dataset_accuracy_totals_path)

balanced_single_dataset_combined_accuracy = pd.read_csv(balanced_single_dataset_accuracy_breakdown_path)
balanced_single_dataset_combined_report = pd.read_csv(balanced_single_dataset_accuracy_totals_path)

print(two_dataset_combined_accuracy.head())
print(two_dataset_combined_report.head())

print(one_dataset_combined_accuracy.head())
print(one_dataset_combined_report.head())

print(balanced_dataset_combined_accuracy.head())
print(balanced_dataset_combined_report.head())

print(balanced_single_dataset_combined_accuracy.head())
print(balanced_single_dataset_combined_report.head())

```

	precision	recall	f1-score	support	source
0	0.654206	0.540541	0.591966	259.0	mimic
1	0.142857	0.006289	0.012048	159.0	mimic
2	0.200000	0.024590	0.043796	122.0	mimic
3	0.432692	0.300000	0.354331	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic

	Label	Accuracy	source
0	No Finding	73.92%	mimic
1	Atelectasis	77.84%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	77.84%	mimic
4	Pneumonia	93.92%	mimic

	precision	recall	f1-score	support	source
0	0.654206	0.540541	0.591966	259.0	mimic
1	0.142857	0.006289	0.012048	159.0	mimic
2	0.200000	0.024590	0.043796	122.0	mimic
3	0.432692	0.300000	0.354331	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic

	Label	Accuracy	source
0	No Finding	73.92%	mimic
1	Atelectasis	77.84%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	77.84%	mimic
4	Pneumonia	93.92%	mimic

	precision	recall	f1-score	support	source
0	0.654206	0.540541	0.591966	259.0	mimic
1	0.142857	0.006289	0.012048	159.0	mimic
2	0.200000	0.024590	0.043796	122.0	mimic
3	0.432692	0.300000	0.354331	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic

	Label	Accuracy	source
0	No Finding	73.92%	mimic
1	Atelectasis	77.84%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	77.84%	mimic
4	Pneumonia	93.92%	mimic

	precision	recall	f1-score	support	source
0	0.654206	0.540541	0.591966	259.0	mimic
1	0.142857	0.006289	0.012048	159.0	mimic
2	0.200000	0.024590	0.043796	122.0	mimic
3	0.432692	0.300000	0.354331	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic

	Label	Accuracy	source
0	No Finding	73.92%	mimic
1	Atelectasis	77.84%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	77.84%	mimic
4	Pneumonia	93.92%	mimic

```

In [184]: def get_accuracy_df(report):
    report['Accuracy'] = report['Accuracy'].replace('%', '', regex=True)

    baseline_accuracies = {}
    for source in report['source'].unique():
        if '+' not in source: # Only single sources
            source_data = report[report['source'] == source]
            for label in source_data['Label'].unique():
                key = (source, label)
                baseline_accuracies[key] = source_data[source_data['La

    # Calculate the changes for combinations involving these sources
    accuracy_changes = []
    for source in report['source'].unique():
        if '+' in source:
            parts = source.split('+')
            source_data = report[report['source'] == source]
            for label in source_data['Label'].unique():
                current_acc = source_data[source_data['Label'] == label]
                for part in parts:
                    base_key = (part, label)
                    base_acc = baseline_accuracies.get(base_key, 0) #
                    change = current_acc - base_acc
                    names_without_base = [x for x in parts if x != part]
                    name_base_first = [part] + names_without_base
                    if len(parts) > 2 and not(source.startswith(part)):
                        pass
                    else:
                        accuracy_changes.append({
                            'source_combination': "+".join(name_base_f
                            'part_source': part,
                            'label': label,
                            'change_in_accuracy': change
                        })

    # Create a DataFrame from the changes
    accuracy_changes_df = pd.DataFrame(accuracy_changes)
    print(accuracy_changes_df)

    return accuracy_changes_df

```

In [190]:

```

two_datasets_accuracy_changes_df = get_accuracy_df(two_dataset_combine
one_datasets_accuracy_changes_df = get_accuracy_df(one_dataset_combine
balanced_dataset_accuracy_changes_df = get_accuracy_df(balanced_dataset_combine
balanced_single_dataset_accuracy_changes_df = get_accuracy_df(balanced_single_dataset_combine

```

	source_combination	part_source	label	change_in_accuracy
0	cxp+nih	cxp	No Finding	-0.001169
1	nih+cxp	nih	No Finding	0.001989
2	cxp+nih	cxp	Atelectasis	0.007013
3	nih+cxp	nih	Atelectasis	-0.000461
4	cxp+nih	cxp	Cardiomegaly	0.002418
..	...	...	...	...
112	cxp+mimic+nih+pad	cxp	Pneumonia	-0.000071
113	cxp+mimic+nih+pad	cxp	Pneumothorax	0.007548
114	cxp+mimic+nih+pad	cxp	Consolidation	0.008808
115	cxp+mimic+nih+pad	cxp	Edema	0.006382
116	cxp+mimic+nih+pad	cxp	Any	-0.000871

[117 rows x 4 columns]

	source_combination	part_source	label	change_in_accuracy
0	cxp+nih	cxp	No Finding	-0.000032
1	nih+cxp	nih	No Finding	0.003126
2	cxp+nih	cxp	Atelectasis	0.006777
3	nih+cxp	nih	Atelectasis	-0.000697
4	cxp+nih	cxp	Cardiomegaly	0.002048
..	...	...	...	...
112	cxp+mimic+nih+pad	cxp	Pneumonia	-0.000071
113	cxp+mimic+nih+pad	cxp	Pneumothorax	0.007548
114	cxp+mimic+nih+pad	cxp	Consolidation	0.008808
115	cxp+mimic+nih+pad	cxp	Edema	0.006382
116	cxp+mimic+nih+pad	cxp	Any	-0.000871

[117 rows x 4 columns]

	source_combination	part_source	label	change_in_accuracy
0	cxp+nih	cxp	No Finding	-3.200000e-07
1	nih+cxp	nih	No Finding	3.126000e-05
2	cxp+nih	cxp	Atelectasis	6.777000e-05
3	nih+cxp	nih	Atelectasis	-6.970000e-06
4	cxp+nih	cxp	Cardiomegaly	2.048000e-05
..	...	...	...	...
112	cxp+mimic+nih+pad	cxp	Pneumonia	-7.100000e-07
113	cxp+mimic+nih+pad	cxp	Pneumothorax	7.548000e-05
114	cxp+mimic+nih+pad	cxp	Consolidation	8.808000e-05
115	cxp+mimic+nih+pad	cxp	Edema	6.382000e-05
116	cxp+mimic+nih+pad	cxp	Any	-8.710000e-06

[117 rows x 4 columns]

	source_combination	part_source	label	change_in_accuracy
0	cxp+nih	cxp	No Finding	-0.0032
1	nih+cxp	nih	No Finding	0.3126
2	cxp+nih	cxp	Atelectasis	0.6777
3	nih+cxp	nih	Atelectasis	-0.0697
4	cxp+nih	cxp	Cardiomegaly	0.2048
...	...	...	...	...
112	cxp+mimic+nih+pad	cxp	Pneumonia	-0.0071
113	cxp+mimic+nih+pad	cxp	Pneumothorax	0.7548
114	cxp+mimic+nih+pad	cxp	Consolidation	0.8808
115	cxp+mimic+nih+pad	cxp	Edema	0.6382
116	cxp+mimic+nih+pad	cxp	Any	-0.0871

[117 rows x 4 columns]

```
In [191]: def plot_breakdown_accuracy_changes(accuracy_changes_df, chart_name="")
    accuracy_changes_df['log_change'] = np.sign(accuracy_changes_df['c

    plt.figure(figsize=(14, 8))
    # Using seaborn's barplot to visualize the data, now using 'label'
    label_order = ['Pneumonia', 'Cardiomegaly', 'Edema', 'Effusion', '
    source_order = [
        'mimic+cxp', 'mimic+nih', 'mimic+pad',
        'cxp+mimic', 'cxp+nih', 'cxp+pad',
        'nih+mimic', 'nih+cxp', 'nih+pad',
        'pad+mimic', 'pad+cxp', 'pad+nih'
    ]


    ax = sns.barplot(data=accuracy_changes_df, y='change_in_accuracy',
                     linewidth=1.5)

    for y in ax.get_yticks():
        plt.axhline(y, color='gray', linewidth=0.5, linestyle='--')

    # Add a horizontal line at zero to indicate no change
    plt.axhline(0, color='gray', linewidth=0.8)

    plt.title(f'Change in Accuracy Relative to Baseline for Each Label')
    plt.ylabel('Relative difference in worst-group performance')
    plt.xlabel('Disease Label')
    plt.xticks()
    plt.legend(title='Source Combination', loc='upper right')
    plt.tight_layout()
    plt.show()
```

```
In [192]: t_breakdown_accuracy_changes(two_datasets_accuracy_changes_df, "evaluation")
t_breakdown_accuracy_changes(one_datasets_accuracy_changes_df, "evaluation")
t_breakdown_accuracy_changes(balanced_dataset_accuracy_changes_df, "evaluation")
t_breakdown_accuracy_changes(balanced_single_dataset_accuracy_changes_df, "evaluation")
```





```
In [193]: def plot_accuracy_per_dataset(report, chart_name=""):
dataset_names = ['nih', 'mimic', 'pad', 'cxp']

for dataset_name in dataset_names:
    dataset_specific = report[report['source'].str.contains(dataset_name)]

    sorted_sources = sorted(dataset_specific['source'].unique(), key=lambda x: len(x))
    dataset_specific = dataset_specific[dataset_specific['source'].isin(sorted_sources)]

    plt.figure(figsize=(12, 6))

    ax = sns.barplot(data=dataset_specific, x='Label', y='Accuracy')

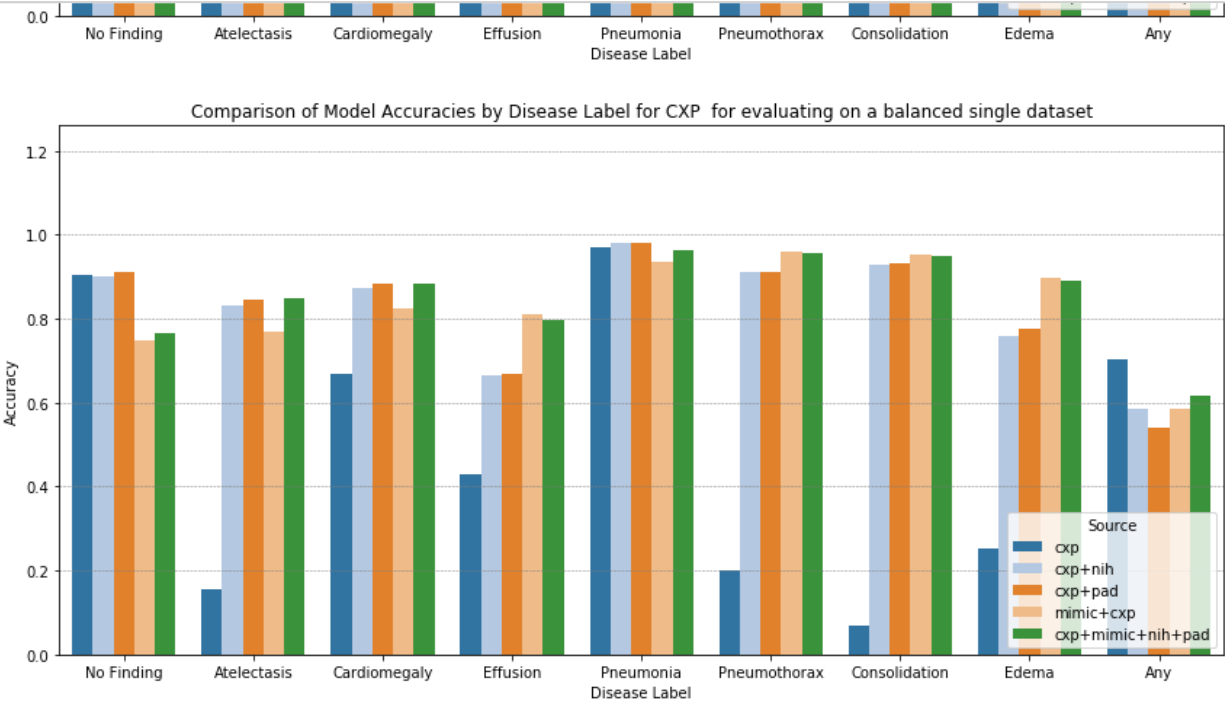
    for y in ax.get_yticks():
        plt.axhline(y, color='gray', linewidth=0.5, linestyle='--')

    plt.title(f'Comparison of Model Accuracies by Disease Label for {dataset_name}')
    plt.ylabel('Accuracy')
    plt.xlabel('Disease Label')
    plt.xticks()
    plt.legend(title='Source', loc='lower right')

    plt.tight_layout()

    # Show the plot
    plt.show()
```


```
In [194]: plot_accuracy_per_dataset(two_dataset_combined_report, "evaluating on  
plot_accuracy_per_dataset(one_dataset_combined_report, "evaluating on  
plot_accuracy_per_dataset(balanced_dataset_combined_report, "evaluating on  
plot_accuracy_per_dataset(balanced_single_dataset_combined_report, "evaluating on
```




Results

I had a mixed bag results as is seen from the graphs above. Let's dive in.

Original paper results:



My results:



As it is seen from the tables above, I got a significantly different result than the article authors. Most accuracy changes from incorporating a second dataset were positive (most impacting in the CXP case), and the negative ones were not as bad as the article stated.

## Ablation study

To make sure I am not missing anything, I have repeated the training and the study on the following dataset combinations for training, test and validation:

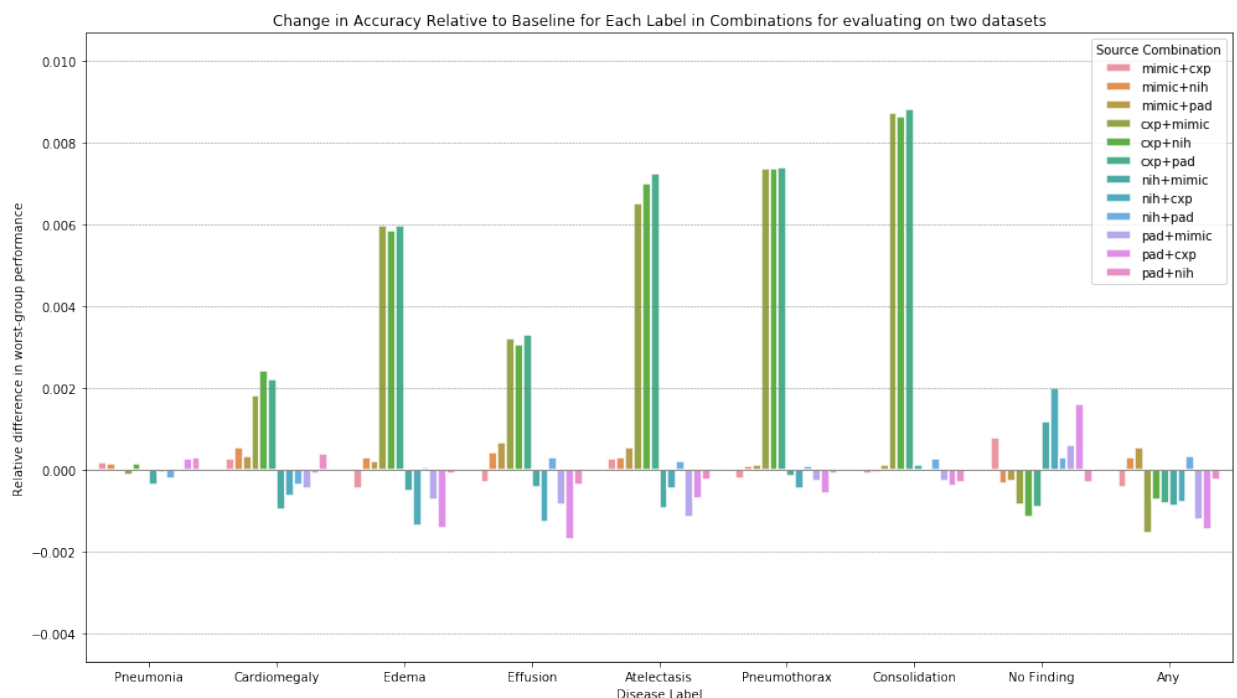
- Unbalanced dataset
  - train on combined (for example, CXP+NIH), validate on combined dataset (for example, CXP+NIH)
  - train on combined (for example, CXP+NIH), validate on single dataset (for example, just CXP)
- Balanced dataset
  - train on combined (for example, CXP+NIH), validate on combined dataset (for example, CXP+NIH)
  - train on combined (for example, CXP+NIH), validate on single dataset (for example, just CXP)

I also re-run the training on different sized results (skip every 20th entry vs every 30th), and in one case ran the study on the full amount of data for one combination.

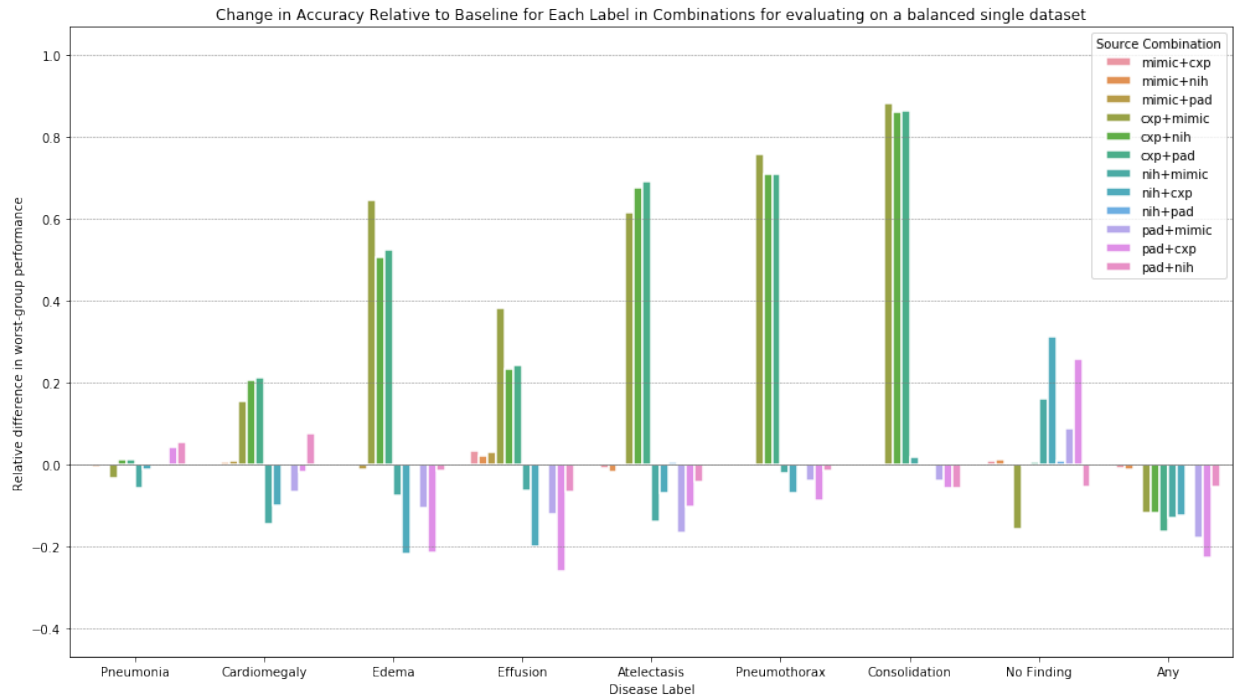
Balancing affected the results slightly, but I still didn't get the picture that the authors of the article got.

When comparing the balanced vs unbalanced dataset validation, the pattern stays the same.

Unbalanced dataset:



Balanced dataset:



When both training and validating on two datasets, the margin of error on accuracy was very small.


## Model comparison

## Discussion


### Is the paper reproducible?

Partially. My resulting visualization ended up being different from the one authors used in the article. However, looking at the results, especially when they are broken down by dataset, one can get to pretty much the same response as the authors of the article did: "It depends", and it looks like balancing the dataset, just as the article stated, doesn't always improve the outcome.

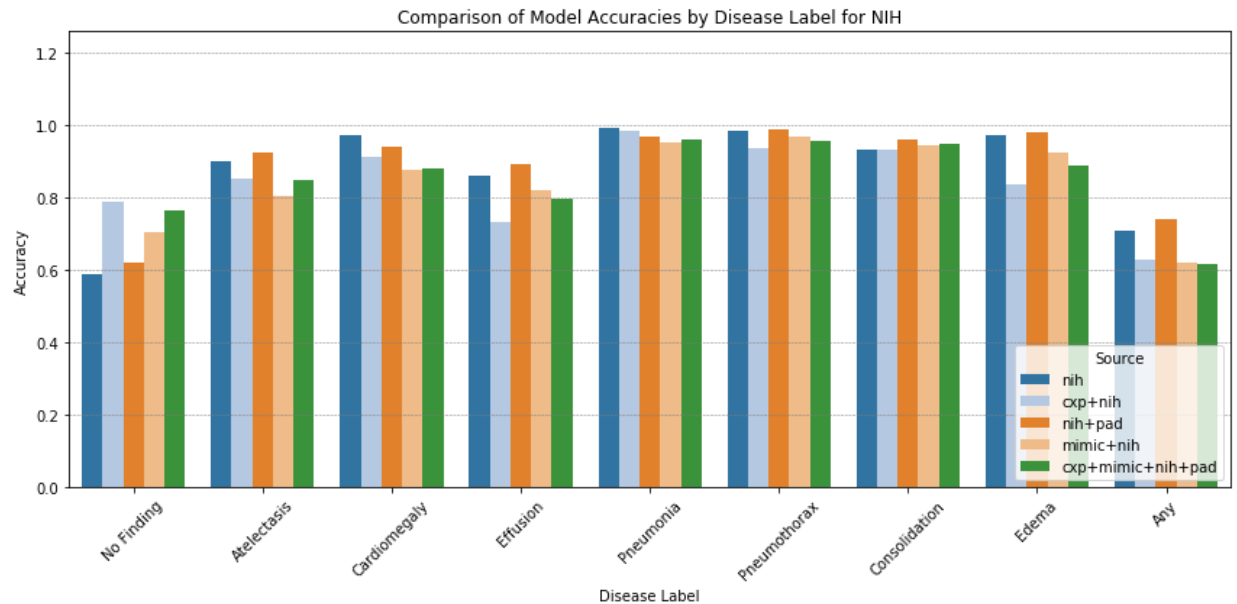
When you look at the results for individual datasets and the resulting combinations, the answer to the main question is very different depending on the dataset:




For PAD, the improvements from adding the second dataset were mainly for the "No finding" category, and the rest generally, even though very slightly, performed worse.



For MIMIC, the answer is "more-or-less", although we can see the very slight positive improvement for most labels.



For NIH, again, the performance mostly decreased.



One interesting exception, CXP almost always significantly benefitted from adding more datasets.

## If the paper is *not* reproducible, explain the results

A portion of the code provided by the article authors was runnable with minimal updates. I was able to reproduce the initial dataset statistics, but wasn't able to run the training code at all.

As a result, I followed the general guidelines given by the authors of the article and wrote my own code - so this may have affected the resulting data I got.

Another factor that might have influenced the resulting data was in data pre-processing. While trying to adapt the authors' code, I discovered an inconsistency in the dataset processing: in some cases the values for certain labels were `True/False`, in some cases numeric `0/1`, and in some cases other numeric value or even a `NaN`.

When adapting the code, I corrected the processing functions so they yield similar-looking results for ease of the combination in a single dataset.

## What was easy

The authors did a great job documenting some parts of the project, for example, access to data. Following the instructions was very easy, and while MIMIC-CXR-JPG dataset access took some time to get, overall the process was a breeze.

## What was difficult

Downloading the datasets is a hassle though, I ran out of space on my laptop, had to buy an external drive and restart the download process for MIMIC-CXR-JPG a few times. In the end, I got a message from my provider that my namely "unlimited" internet for the months was used 100% and I will be charged for each extra GB I use.

There are a few notebooks and standalone scripts provided to process the data. While it is possible to figure out what steps need to be done in what order, many of the parts of the process are not documented. 'pyproject.toml' did not run successfully for me, and I got stuck for a while, trying to figure out why and how to run it (I have a suspicion my processor architecture is not supported).

As a result, I opted to re-implement the training and model validation myself. There is code for training and validation in the project, which has a lot of comments (great!), but the process itself is not well documented, so the reproducer is left figuring out which steps in the code are needed and which are not, and how to adapt it to use for their experiment. The code is very general and there is a lot of it. There are some pointers in the readme, but they were not sufficient for reproducing things successfully without additional modification. wandb wasn't working for me either.

The data is not processed evenly / equally for each dataset, there are different values for the same labels (`NaN`, `True/False`, `1/0`, `1.1/0.0`). I had to write some processing code to make sure to mitigate those differences.

Additional complication was due to the fact that the amount of data is very large. Any training or processing takes a long time, the notebook kernel dies frequently and the overall process is frustrating. Downloading MIMIC took a week and ate all my provider's internet allowance for a month. Running the training on a full dataset proved it to be difficult due to



the whole different set of circumstances: I encountered out of memory issues, kernel panic, random automatic Windows updates, power outage, kids getting to the computer and switching the power supply off. In the end, the estimation to run the training on the full dataset was circa 16+ days, so I opted for running the code on the subset of the data instead.

I tried to avoid multiple separate files and scripts, and pulled many of the data preprocessing into my notebook. However, this increased the runtime of the notebook significantly. Additional factor affecting the runtime is the size of the input data, even when working on one dataset. I was never able to achieve the 8-minute runtime, in fact, my record was around 10 hours.

## Suggestions for the author

Trim the codebase leaving only relevant parts. Add documentation for the training and validation process. Add some background on why wandb is used and how to use it for this project correctly. Provide a suggested order of execution for the notebooks. Provide the instructions and code to plot the results.

## Future plans

I am still very interested in answering the question why my results, even though providing a similar answer, looked so different from the article authors'. I plan to do a few more things to try and figure this out:

- Try different models (for example, `LinearSVC` and other models tried by the authors of the article before arriving on their final model), and hyperparameters
- Try and incorporate more of the training and validation code provided by the article authors, and/or triple check mine
- Run the training on the full dataset (ensuring the computer is on backup power supply and is unreachable to anyone trying to switch it off)

# References

1. Rhys Compton; Lily Zhang; Aahlad Puli; Rajesh Ranganath, When More is Less: Incorporating Additional Datasets Can Hurt Performance By Introducing Spurious Correlations, arXiv preprint, 2023-08-09, Accepted at MLHC 2023, doi: [10.48550/arXiv.2308.04431](https://doi.org/10.48550/arXiv.2308.04431) (<https://doi.org/10.48550/arXiv.2308.04431>)
2. Haoran Zhang, Natalie Dullerud, Laleh Seyyed-Kalantari, Quaid Morris, Shalmali Joshi, and Marzyeh Ghassemi. An empirical framework for domain generalization in clinical settings. In Proceedings of the Conference on Health, Inference, and Learning, pages 279–290, 2021, doi: [10.48550/arXiv.2103.11163](https://doi.org/10.48550/arXiv.2103.11163) (<https://doi.org/10.48550/arXiv.2103.11163>)
3. Gao Huang, Zhuang Liu, Laurens Van Der Maaten, and Kilian Q Weinberger. Densely connected convolutional networks. In Proceedings of the IEEE conference on computer vision and pattern recognition, pages 4700–4708, 2017, doi: [10.48550/arXiv.1608.06993](https://doi.org/10.48550/arXiv.1608.06993) (<https://doi.org/10.48550/arXiv.1608.06993>)
4. Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In 2009 IEEE Conference on Computer Vision and Pattern Recognition, pages 248–255, 2009, doi: [10.1109/CVPR.2009.5206848](https://doi.org/10.1109/CVPR.2009.5206848) (<https://doi.org/10.1109/CVPR.2009.5206848>)
5. John R Zech, Marcus A Badgeley, Manway Liu, Anthony B Costa, Joseph J Titano, and Eric Karl Oermann. Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: a cross-sectional study. PLoS medicine, 15(11): e1002683, 2018, doi: [10.1371/journal.pmed.1002683](https://doi.org/10.1371/journal.pmed.1002683) (<https://doi.org/10.1371/journal.pmed.1002683>)

In [ ]: