

# Nina Stawski's (group 90) final project report

## Paper #39

### When More is Less: Incorporating Additional Datasets Can Hurt Performance By Introducing Spurious Correlations [10.48550/arXiv.2308.04431](https://doi.org/10.48550/arXiv.2308.04431) (<https://doi.org/10.48550/arXiv.2308.04431>)

Illinois ID: ninas2

- [VIDEO LINK \(https://youtu.be/ikw3XVGCM28\)](https://youtu.be/ikw3XVGCM28)
- [My Final project GitHub repo link \(https://github.com/nstawski/dlh-final-project\)](https://github.com/nstawski/dlh-final-project)
- [Original paper GitHub repo \(https://github.com/basedrhys/ood-generalization\)](https://github.com/basedrhys/ood-generalization)

## Introduction

### Background of the problem

#### Type of problem

This is a data preparation and processing problem. The authors of the article are testing a common belief that adding more data improves the resulting model performance. Their main hypothesis, which they subsequently prove, is that incorporating more data does not necessarily improve the model performance. It can introduce spurious correlations, and hurt the resulting model performance rather than helping it.

#### What is the importance/meaning of solving the problem

The paper is challenging a common belief, meaning a lot of researchers are likely trying to incorporate as much data as they can expecting it would improve the performance of their models. The outcome of this research would provide guidance on the possible pitfalls and the cases where you wouldn't want to add external data - so it could set a new standard of processing and incorporating data for everyone in the field.

## **The difficulty of the problem**

The problem is non-obvious and the paper is challenging the common belief held in the industry. The authors are putting a lot of state-of-the-art approaches to the test, and attempt to quantify the results as well as provide new standards and explanations. This is extremely hard to do so I believe the problem is difficult.

## **The state of the art methods and effectiveness**

The "industry standard" way of improving model performance is adding more data from additional datasets, which the authors of this article prove to not be effective, and even being harmful in many cases.

One of the main issues causing the model performance decrease when adding more data from other sources is spurious correlations, which in case of x-rays could be coming even from the scanner artifacts, or other hospital-specific data. One of the state-of-the-art ways to mitigate this is balancing a dataset to reduce the influence of hospital-specific factors. While balancing definitely improved the situation, the resulting model performance was still in many cases worse than with a single-hospital dataset.

## **Paper explanation**

### **What did the paper propose**

The paper used four most-used chest x-ray datasets - MIMIC-CXR-JPG, CheXpert, PadChest, ChestXray8 - to disprove a popular belief that adding more data always would improve the performance of your model. They postulate that, for the specific x-ray data, even the scanners themselves, the way hospitals produce data, or send specific patients to specific places to do their scan, can introduce spurious correlations which, in many cases, significantly affect the worst group performance.

### **What is the innovations of the method**

Existing research (for example, John R Zech, Marcus A Badgeley, Manway Liu, Anthony B Costa, Joseph J Titano, and Eric Karl Oermann. Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: a cross-sectional study. PLoS medicine, 15(11): e1002683, 2018.) proves that adding a second dataset improves the average per-group accuracy. In contrast, the paper I am reproducing focuses on the worst per-group accuracy.

### **How well the proposed method work (in its own metrics)**

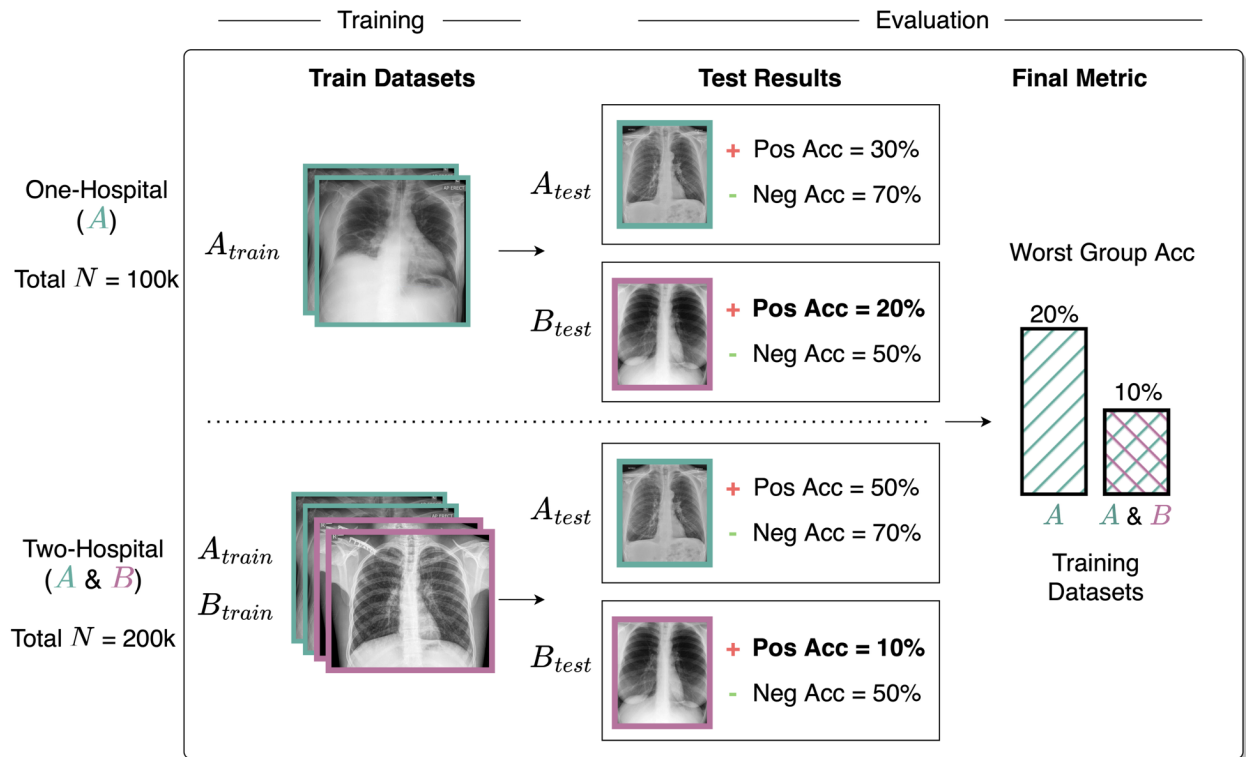
According to the article authors, their method works really well and proves that in nearly 50% of cases adding a second dataset, and even balancing it to reduce spurious correlations doesn't get the model to perform better than without that additional dataset. The models pick up on hospital-specific features even if those features weren't explicitly defined in the original data. They postulate that every CNN model, regardless of training disease or datasets, learns embeddings that can distinguish any of the hospital sources with near-perfect accuracy, even if the embeddings were trained via one or two hospitals' data.

**What is the contribution to the research regime (referring the Background above, how important the paper is to the problem).**

The article cautions against blindly adding more datasets, and provides a number of approaches you can take if you still decide to do so. The conclusion is adding more data shouldn't be done blindly. The authors of the article definitely discourage the researchers from the most common approach of throwing data at the problem to improve model performance.

# Scope of Reproducibility:

List hypotheses from the paper you will test and the corresponding experiments you will run.



## Hypothesis 1

In 43% of training dataset/disease tasks, adding data from an external source hurts worst-group performance.

## Hypothesis 2

Balancing the dataset to reduce spurious correlations is often beneficial, but in the scenarios where adding an additional data source hurts generalization performance, it does not always improve generalization; in some cases, training on a balanced dataset achieves lower worst-group accuracy than training on datasets from one or two hospitals.

# Methodology

This methodology is the core of your project. It consists of run-able codes with necessary annotations to show the experiment you executed for testing the hypotheses.

The methodology at least contains two subsections **data** and **model** in your experiment.

## Python environment and package versions

```
In [5]: # !pip install importlib  
# !pip install torch  
# !pip install torchvision  
# !pip install pandas  
# !pip install matplotlib  
# !pip install imblearn
```

```
In [1]: import numpy as np
import pandas as pd
from pathlib import Path
import os
from os.path import exists
import sys
import matplotlib.pyplot as plt
import seaborn as sns
from PIL import Image, ImageFile

import json
import random
from IPython.display import display
from datetime import datetime
from sklearn.metrics import f1_score, classification_report, accuracy_

import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset
from torchvision import datasets, models, transforms

from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler

import Data_Constants as Constants

#making sure all referenced files are reloaded
import importlib
importlib.reload(Constants)
```

```
Out[1]: <module 'Data_Constants' from 'C:\\Users\\Stan\\Documents\\GitHub\\dlh-final-project\\Data_Constants.py'>
```

I have originally set up the Jupyter Notebook and was running the code on my laptop. However, processing of the data took days and the training was promising to take months. My husband had a gaming computer with a powerful video card and Cuda available, so I ended up using his machine.

To keep the environment isolated, I installed `Anaconda` and created a separate environment for all my packages. I then ran `Jupyter` with access to local network and opened my notebook remotely from my laptop:

```
jupyter notebook --ip 192.168.x.xxx --port 8888
```

Original requirements listed a bit different versions, but due to package compatibility I had to update them. Most significant changes were: install a version of `torch` that supports Cuda, and upgrade `torchvision` to `0.9.1`

```
In [2]: env_df = pd.read_csv('env_packages.csv', header=None, names=['Package']
        styled_env_df = env_df.style.set_table_styles([
            {'selector': 'th, td',
             'props': [('min-width', '200px')]},
            {'selector': 'table',
             'props': [('width', '70%')]})
        ).set_properties(**{
            'background-color': 'white',
            'color': 'black',
            'border-color': 'black',
            'border-style': 'solid',
            'border-width': '1px'
        }).hide_index()

        styled_env_df
```

Out [2]:

Package	Version
python	3.6.13
pip	21.2.2
jupyter_core	4.8.1
imbalanced-learn	0.8.1
jupyter	1.0.0
matplotlib	2.2.2
numpy	1.19.5
pandas	1.1.0
pillow	8.4.0
scikit-learn	0.24.2
scipy	1.5.1
seaborn	0.11.2
torch	1.8.1+cu111
torchvision	0.9.1



It was a lot of debugging to make sure Cuda is available in the notebook, and it took me a few days to finally make it work. It made all processing and training code run a lot faster.

The next challenge I encountered was that the full source data did not fit into my laptop and I had to get an external storage to hold it - however, that external storage's speed wasn't keeping up. So instead, my husband got an upgrade to his internal storage with a very fast ssd.

```
In [3]: # os.environ['KMP_DUPLICATE_LIB_OK']='True'
# torch.set_default_device('cuda')

torch.set_default_tensor_type('torch.cuda.FloatTensor')
Tensor = torch.cuda.FloatTensor if torch.cuda.is_available() else torch.FloatTensor

print("Cuda is available:", torch.cuda.is_available())
```

Cuda is available: True

## Data

The study is using four datasets: MIMIC-CXR-JPG, CheXpert, PadChest, ChestXray8

The datasets are being filtered to include only frontal (PA/AP) images. Instances are labeled with one or more pathologies. Each dataset has a different set of diseases but they are preprocessed using code derived from ClinicalDG2 (Zhang et al., 2021) to extract the eight common labels and homogenize the datasets. Additionally, authors of the article created the Any label which indicates a positive label for any of the seven common disease labels, resulting in nine different binary labels. All experiments use the labels in a binary manner; a pathology is chosen as the target label, with an instance labeled 1 if the pathology of interest is present and 0 otherwise.

The authors apply an 80%/10%/10% subject-wise train/val/test split, with the same split used across seeds.

## MIMIC-CXR

1. [Obtain access \(https://mimic-cxr.mit.edu/about/access/\)](https://mimic-cxr.mit.edu/about/access/) to the MIMIC-CXR-JPG Database Database on PhysioNet and download the [dataset \(https://physionet.org/content/mimic-cxr-jpg/2.0.0/\)](https://physionet.org/content/mimic-cxr-jpg/2.0.0/). The best option is downloading from the GCP bucket:

```
gcloud auth login
mkdir MIMIC-CXR-JPG
gsutil -m rsync -d -r gs://mimic-cxr-jpg-2.0.0.physionet.org MIMIC-CXR-JPG
```

2. In order to obtain gender information for each patient, you will need to obtain access to [MIMIC-IV \(https://physionet.org/content/mimiciv/0.4/\)](https://physionet.org/content/mimiciv/0.4/). Download `core/patients.csv.gz` and place the file in the `MIMIC-CXR-JPG` directory.

## CheXpert

1. Sign up with your email address [here](https://stanfordmlgroup.github.io/competitions/chexpert/) (<https://stanfordmlgroup.github.io/competitions/chexpert/>).
2. Download either the original or the downsampled dataset (we recommend the downsampled version - `CheXpert-v1.0-small.zip`) and extract it.

## ChestX-ray8

1. Download the `images` folder and `Data_Entry_2017_v2020.csv` from the [NIH website \(https://nihcc.app.box.com/v/ChestXray-NIHCC\)](https://nihcc.app.box.com/v/ChestXray-NIHCC).
2. Unzip all of the files in the `images` folder.

## PadChest

1. The paper uses a resized version of PadChest, which can be downloaded [here](https://academictorrents.com/details/96ebb4f92b85929eadfb16761f310a6d04105797) (<https://academictorrents.com/details/96ebb4f92b85929eadfb16761f310a6d04105797>).
2. Unzip `images-224.tar`.
  - Statistics: include basic descriptive statistics of the dataset like size, cross validation split, label distribution, etc.
  - Data process: how do you manipulate the data, e.g., change the class labels, split the dataset to train/valid/test, refining the dataset.
  - Illustration: printing results, plotting figures for illustration.
  - You can upload your raw dataset to Google Drive and mount this Colab to the same directory. If your raw dataset is too large, you can upload the processed dataset and have a code to load the processed dataset.

## Data Processing

The original pre-processing for the article was done using the scripts outside of the Jupyter Notebook. Some of them didn't work for me, and the installation process didn't succeed despite multiple attempts either. Instead, I have adapted some of the original scripts to run in

the notebook (with some modifications so they actually work with my data), using the external "Constants.py" file that points to the location of the datasets.

1. In `./Data_Constants.py`, update `image_paths` to point to each of the four directories that you downloaded.
2. Run the next two cells to pre-process the data

## Validating

I am using the validation and pre-processing code provided by the authors of the article, with some modifications to make it run as expected.

```
In [4]: #making sure constants are up to date if they were changed
importlib.reload(Constants)

def validate_mimic():
    img_dir = Path(Constants.image_paths['MIMIC'])
    meta_dir = Path(Constants.meta_paths['MIMIC'])

    print('meta_dir', meta_dir, os.getcwd())
    print('meta_dir', meta_dir/'mimic-cxr-2.0.0-metadata.csv')
    assert (meta_dir/'mimic-cxr-2.0.0-metadata.csv').is_file()
    assert (meta_dir/'mimic-cxr-2.0.0-negbio.csv').is_file()
    assert (meta_dir/'patients.csv').is_file()
    # modified the file that's being checked since I don't have the fu
    # in the original script, the file in p19 was being checked.
    assert (img_dir/'p10/p10000032/s50414267/02aa804e-bde0afdd-112c0b3

def validate_cxp():
    img_dir = Path(Constants.image_paths['CXP'])
    if (img_dir/'CheXpert-v1.0').is_dir():
        cxp_subfolder = 'CheXpert-v1.0'
    else:
        cxp_subfolder = 'CheXpert-v1.0-small'
    assert (img_dir/cxp_subfolder/'train.csv').is_file()
    assert (img_dir/cxp_subfolder/'train/patient48822/study1/view1_fro
    assert (img_dir/cxp_subfolder/'valid/patient64636/study1/view1_fro

def validate_pad():
    img_dir = Path(Constants.image_paths['PAD'])
    meta_dir = Path(Constants.meta_paths['PAD'])
    assert (meta_dir/'PADCHEST_chest_x_ray_images_labels_160K_01.02.19
    assert (img_dir/'185566798805711692534207714722577525271_qb3lyn.pr

def validate_nih():
    img_dir = Path(Constants.image_paths['NIH'])
    meta_dir = Path(Constants.meta_paths['NIH'])
```

```

assert (meta_dir/'Data_Entry_2017.csv').is_file()
assert (img_dir/'images/00002072_003.png').is_file()

def validate_splits():
    for dataset in Constants.df_paths:
        for split in Constants.df_paths[dataset]:
            assert Path(Constants.df_paths[dataset][split]).is_file()

def validate_all():
    validate_mimic()
    validate_cxp()
    validate_nih()
    validate_pad()

```

## Data pre-processing setup

In [5]: *# making sure constants are up to date if they were changed after runn*  
importlib.reload(Constants)

```

def preprocess_mimic():
    img_dir = Path(Constants.image_paths['MIMIC'])
    meta_dir = Path(Constants.meta_paths['MIMIC'])
    out_folder = meta_dir/'clinicaldg'
    out_folder.mkdir(parents = True, exist_ok = True)

    patients = pd.read_csv(meta_dir/'patients.csv')
    labels = pd.read_csv(meta_dir/'mimic-cxr-2.0.0-negbio.csv')
    meta = pd.read_csv(meta_dir/'mimic-cxr-2.0.0-metadata.csv')

    df = meta.merge(patients, on = 'subject_id').merge(labels, on = ['
    df['age_decile'] = pd.cut(df['anchor_age'], bins = list(range(0, 1
    df['frontal'] = df.ViewPosition.isin(['AP', 'PA'])

    df['path'] = df.apply(lambda x: os.path.join(f'p{str(x["subject_id
    df.to_csv(out_folder/"preprocessed.csv", index=False)

def preprocess_pad():
    # I have modified this function from the original one, because I w
    img_dir = Path(Constants.image_paths['PAD'])
    meta_dir = Path(Constants.meta_paths['PAD'])
    out_folder = meta_dir/'clinicaldg'
    out_folder.mkdir(parents=True, exist_ok=True)

    dtype_spec = {
        'ImageID': str,
        'StudyID': str,

```

```

        'PatientID': str,
        'PatientBirth': str, # converting this to the integer later to
        'PatientSex_DICOM': str,
        'ViewPosition_DICOM': str,
        'Projection': str,
        'Labels': str,
        'WindowCenter_DICOM': str,
        'WindowWidth_DICOM': str
    }

df = pd.read_csv(meta_dir/'PADCHEST_chest_x_ray_images_labels_160K')
df = df[['ImageID', 'StudyID', 'PatientID', 'PatientBirth', 'PatientSex_DICOM', 'ViewPosition_DICOM', 'Projection', 'Labels', 'WindowCenter_DICOM', 'WindowWidth_DICOM']]
df = df[~df["Labels"].isnull()]
df = df[df["ImageID"].apply(lambda x: os.path.exists(os.path.join(meta_dir, x))) == True]
df = df[df.Projection.isin(['PA', 'L', 'AP_horizontal', 'AP'])]

df['frontal'] = ~(df['Projection'] == 'L')
df = df[~df['Labels'].apply(lambda x: 'exclude' in x or 'unchanged' in x)]

mapping = dict()
mapping['Effusion'] = ['hydropneumothorax', 'empyema', 'hemothorax']
mapping["Consolidation"] = ["air bronchogram"]
mapping['No Finding'] = ['normal']

for pathology in Constants.take_labels:
    mask = df["Labels"].str.contains(pathology.lower())
    if pathology in mapping:
        for syn in mapping[pathology]:
            mask |= df["Labels"].str.contains(syn.lower())
    df[pathology] = mask.astype(int)

df['PatientBirth'] = df['PatientBirth'].dropna().astype(float).astype(int)
df['Age'] = 2017 - df['PatientBirth']
df.reset_index(drop=True).to_csv(out_folder/"preprocessed.csv", index=False)

def preprocess_cxp():
    img_dir = Path(Constants.image_paths['CXP'])
    out_folder = img_dir/'clinicaldg'
    if (img_dir/'CheXpert-v1.0'/'train.csv').is_file():
        df = pd.concat([pd.read_csv(img_dir/'CheXpert-v1.0'/'train.csv'),
                        pd.read_csv(img_dir/'CheXpert-v1.0'/'valid.csv')],
                        ignore_index = True)
    elif (img_dir/'CheXpert-v1.0-small'/'train.csv').is_file():
        df = pd.concat([pd.read_csv(img_dir/'CheXpert-v1.0-small'/'train.csv'),
                        pd.read_csv(img_dir/'CheXpert-v1.0-small'/'valid.csv')],
                        ignore_index = True)
    elif (img_dir/'train.csv').is_file():
        raise ValueError('Please set Constants.image_paths["CXP"] to be the directory and rerun this script.')
    else:

```

```

else:
    raise ValueError("CheXpert files not found!")

out_folder.mkdir(parents = True, exist_ok = True)

df['subject_id'] = df['Path'].apply(lambda x: int(Path(x).parent.p
df['Path'] = df['Path'].apply(lambda x: str(x).replace("CheXpert-v
df.reset_index(drop = True).to_csv(out_folder/"preprocessed.csv",

def preprocess_nih():
    img_dir = Path(Constants.image_paths['NIH'])
    meta_dir = Path(Constants.meta_paths['NIH'])
    out_folder = meta_dir/'clinicaldg'
    out_folder.mkdir(parents = True, exist_ok = True)
    df = pd.read_csv(meta_dir/"Data_Entry_2017.csv")
    df['labels'] = df['Finding Labels'].apply(lambda x: x.split('|'))

    for label in Constants.take_labels:
        df[label] = df['labels'].apply(lambda x: label in x)
    df.reset_index(drop = True).to_csv(out_folder/"preprocessed.csv",

```

In [103]: %%script false --no-raise-error  
*# skipping this cell since I already ran this.*

```

if __name__ == '__main__':
    print("Validating paths...")
    validate_all()
    print("Preprocessing MIMIC-CXR...")
    preprocess_mimic()
    print("Preprocessing CheXpert...")
    preprocess_cxp()
    print("Preprocessing ChestX-ray8...")
    preprocess_nih()
    print("Preprocessing PadChest... This might take a few minutes...")
    preprocess_pad()
    print("Done.")

```

```

Validating paths...
meta_dir C:\Nina\e-root\data\mimic\physionet.org\files\mimic-cxr-jpg\
2.0.0 C:\Users\Stan\Documents\GitHub\dlh-final-project
meta_dir C:\Nina\e-root\data\mimic\physionet.org\files\mimic-cxr-jpg\
2.0.0\mimic-cxr-2.0.0-metadata.csv
Preprocessing MIMIC-CXR...
Preprocessing CheXpert...
Preprocessing ChestX-ray8...
Preprocessing PadChest... This might take a few minutes...
Done.

```

## Next, we need to resize and process the data.

I am using the code provided by the authors of the article to do this.

```
In [6]: def process_MIMIC(split, only_frontal):
    copy_subjectid = split['subject_id']
    split = split.drop(columns = ['subject_id']).replace(
        [[None], -1, "[False]", "[True]", "[ True]", 'UNABLE TO OE
        'DIVORCED', 'SEPARATED', '0-10', '10-20', '20-30', '30-40
        '>=90'],
        [0, 0, 0, 1, 1, 0, 0, 'MARRIED/LIFE PARTNER', 'MARRIED/LIF
        'DIVORCED/SEPARATED', '0-20', '0-20', '20-40', '20-40', '

    split['subject_id'] = copy_subjectid.astype(str)
    split['study_id'] = split['study_id'].astype(str)
    split['Age'] = split["age_decile"]
    split['Sex'] = split["gender"]
    split = split.rename(
        columns = {
            'Pleural Effusion': 'Effusion',
        })
    split['path'] = split['path'].astype(str).apply(lambda x: os.path.
    if only_frontal:
        split = split[split.frontal]

    split['env'] = 'MIMIC'
    split.loc[split.Age == 0, 'Age'] = '0-20'

    return split[['subject_id', 'path', 'Sex', "Age", 'env', 'frontal', '

def process_NIH(split, only_frontal = True):
    split['Patient Age'] = np.where(split['Patient Age'].between(0,19)
    split['Patient Age'] = np.where(split['Patient Age'].between(20,39
    split['Patient Age'] = np.where(split['Patient Age'].between(40,59
    split['Patient Age'] = np.where(split['Patient Age'].between(60,79
    split['Patient Age'] = np.where(split['Patient Age']>=80, 81, spli

    copy_subjectid = split['Patient ID']

    split = split.drop(columns = ['Patient ID']).replace([[None], -1,
        [0, 0, 0, 1, 1, "0-20", "20-40", "40-60",

    split['subject_id'] = copy_subjectid.astype(str)
    split['Sex'] = split['Patient Gender']
    split['Age'] = split['Patient Age']
    split = split.drop(columns=["Patient Gender", 'Patient Age'])
    split['path'] = split['Image Index'].astype(str).apply(lambda x: c
    split['env'] = 'NIH'
```

```

split['frontal'] = True
split['study_id'] = split['subject_id'].astype(str)
return split[['subject_id', 'path', 'Sex', "Age", 'env', 'frontal', 's

def process_CXP(split, only_frontal):
    split['Age'] = np.where(split['Age'].between(0,19), 19, split['Age']
    split['Age'] = np.where(split['Age'].between(20,39), 39, split['Age']
    split['Age'] = np.where(split['Age'].between(40,59), 59, split['Age']
    split['Age'] = np.where(split['Age'].between(60,79), 79, split['Age']
    split['Age'] = np.where(split['Age']>=80, 81, split['Age'])

    copy_subjectid = split['subject_id']
    split = split.drop(columns = ['subject_id']).replace([[None], -1,
                                                         [0, 0, 0, 1, 1, "0-20", "20-40", "40-60",

    split['subject_id'] = copy_subjectid.astype(str)
    split['Sex'] = np.where(split['Sex']=='Female', 'F', split['Sex'])
    split['Sex'] = np.where(split['Sex']=='Male', 'M', split['Sex'])
    split = split.rename(
        columns = {
            'Pleural Effusion': 'Effusion',
            'Lung Opacity': 'Airspace Opacity'
        })
    split['path'] = split['Path'].astype(str).apply(lambda x: os.path.
    split['frontal'] = (split['Frontal/Lateral'] == 'Frontal')
    if only_frontal:
        split = split[split['frontal']]
    split['env'] = 'CXP'
    split['study_id'] = split['path'].apply(lambda x: x[x.index('patie
    return split[['subject_id', 'path', 'Sex', "Age", 'env', 'frontal', 's

def process_PAD(split, only_frontal):
    split['Age'] = np.where(split['Age'].between(0,19), 19, split['Age']
    split['Age'] = np.where(split['Age'].between(20,39), 39, split['Age']
    split['Age'] = np.where(split['Age'].between(40,59), 59, split['Age']
    split['Age'] = np.where(split['Age'].between(60,79), 79, split['Age']
    split['Age'] = np.where(split['Age']>=80, 81, split['Age'])

    split = split.replace([[None], -1, "[False]", "[True]", "[ True]",
                               [0, 0, 0, 1, 1, "0-20", "20-40", "40-60",

    split.loc[split['Age'] == 0.0, 'Age'] = '0-20'
    split.loc[split['Age'].isnull(), 'Age'] = '0-20'
    split = split.rename(columns = {
        'PatientID': 'subject_id',
        'StudyID': 'study_id',
        'PatientSex_DICOM' : 'Sex'
    })

```



```

split.loc[~split['Sex'].isin(['M', 'F', 'O']), 'Sex'] = 'O'
split['path'] = split['ImageID'].astype(str).apply(lambda x: os.p
if only_frontal:
    split = split[split['frontal']]
split['env'] = 'PAD'
return split[['subject_id', 'path', 'Sex', "Age", 'env', 'frontal', 's

def split(df, split_portions = (0.8, 0.9), seed=0):
    # We don't want the data splits to be affected by seed
    # So lets temporarily set the seed to a static value...
    rand_state = np.random.get_state()
    np.random.seed(seed)

    # Split our data (irrespective of the random seed provided in train)
    subject_df = pd.DataFrame({'subject_id': np.sort(df['subject_id']).
    subject_df['random_number'] = np.random.uniform(size=len(subject_d

    train_id = subject_df[subject_df['random_number'] <= split_portion
    valid_id = subject_df[(subject_df['random_number'] > split_portion
    test_id = subject_df[subject_df['random_number'] > split_portions[

    train_df = df[df.subject_id.isin(train_id.subject_id)]
    valid_df = df[df.subject_id.isin(valid_id.subject_id)]
    test_df = df[df.subject_id.isin(test_id.subject_id)]

    # ...then return the random state back to what it was
    np.random.set_state(rand_state)

    return train_df, valid_df, test_df

def get_process_func(env):
    if env == 'MIMIC':
        return process_MIMIC
    elif env == 'NIH':
        return process_NIH
    elif env == 'CXP':
        return process_CXP
    elif env == 'PAD':
        return process_PAD
    else:
        raise NotImplementedError

```

```
In [7]: # show data paths from constants
Constants.df_paths

def img_exists(path):
    return exists(path)

def is_diseased(row):
    # diseases = Constants.take_labels[1:]
    return int((row[Constants.take_labels[1:]]).sum() > 0)
```

**The following cell is pre-processing the data and will take a long time to run**

The cell below needs to run once, after that everything is saved into the CSV file and can be loaded from there. this block of code needs to re-run only if the data changed.

```

In [106]: %%script false --no-raise-error
# skipping this cell since I already ran this.

# loads data with random splits
print('This might take a while.')

for data_env in Constants.df_paths:
    print('Processing:', data_env)
    func = get_process_func(data_env)
    print('Got processing function, filtering by only frontal...')
    df_env = func(pd.read_csv(Constants.df_paths[data_env]), only_frontal=True)
    print('Filtering out the data without images...')
    df_env["img_exists"] = df_env["path"].apply(img_exists)
    print(df_env["img_exists"].value_counts())
    df_env = df_env[df_env["img_exists"]]

    df_env = df_env.fillna(0)

    print('Adding "All" column...')
    df_env["All"] = df_env.apply(is_diseased, axis=1)

    print('Saving results...')
    df_env.to_csv(f'{Constants.base_path}\\processed\\{data_env}.csv',
                  index=False)

    display(df_env)

print("Done.")

```

Got processing function, filtering by only frontal...

Filtering out the data without images...

True 230693

False 18

Name: img\_exists, dtype: int64

Adding "All" column...

Saving results...

	subject_id	path
0	10000032	C:\Nina\e-root\data\mimic\physionet.org\files\...
2	10000032	C:\Nina\e-root\data\mimic\physionet.org\files\...
4	10000032	C:\Nina\e-root\data\mimic\physionet.org\files\...
5	10000032	C:\Nina\e-root\data\mimic\physionet.org\files\...

## Resample data

```
In [8]: dfs = {}
print('Processing the data, splitting to all, train, val and test...')
for env in Constants.df_paths:
    func = get_process_func(env)
    df_env = pd.read_csv(f"{Constants.base_path}/processed/{env}.csv")

    print('Source:', env)
    print('Data length:', len(df_env))

    train_df, valid_df, test_df = split(df_env)
    dfs[env] = {
        'all': df_env,
        'train': train_df,
        'val': valid_df,
        'test': test_df
    }
    print(f'{env}: done.')

print('All done.')
```

```
Processing the data, splitting to all, train, val and test...
Source: MIMIC
Data length: 230693
MIMIC: done.
Source: CXP
Data length: 191229
CXP: done.
Source: NIH
Data length: 112120
NIH: done.
Source: PAD
Data length: 99827
PAD: done.
All done.
```

## Balancing the dataset

```
In [9]: def get_prop(df, column="Pneumonia"):
    num_instances = len(df)
    num_diseased = df[df[column] == 1][column].count()
    return num_diseased / (num_instances - num_diseased)

def get_resample_class(orig_prop, new_prop, resample_method):
    if new_prop > orig_prop:
```

```

    .. new_prop < orig_prop:
        if resample_method == "over":
            return 1
        else:
            return 0
    if new_prop < orig_prop:
        if resample_method == "under":
            return 1
        else:
            return 0

def calculate_num_resample(df, orig_prop, new_prop, resample_method):
    pass

def balance_df_label(df, sampler, label_bal=0.05154780337262089, inverse
    target = df["Pneumonia"] == 1
    rus = sampler(random_state=0, sampling_strategy=label_bal if not i
    res_df, _ = rus.fit_resample(df, target)

    print(f"Previous pneumonia prop: {get_pneumonia_prop(df)} with {le
    print(f"Resampled pneumonia prop: {get_pneumonia_prop(res_df)} wit

    return res_df

def balance_proportion(orig_df, new_df, resample_method="over", column
    orig_df = orig_df.fillna(0.0)
    orig_prop = get_prop(orig_df, column)
    new_prop = get_prop(new_df, column)
    assert resample_method in ["over", "under"]
    resample_class = get_resample_class(orig_prop, new_prop, resample_
    print(f"Resampling '{column}' via '{resample_method}' on class {re

    # Estimate the number of items we'll need to resample
    df_diseased = orig_df[orig_df[column] == 1.0]
    df_normal = orig_df[orig_df[column] == 0.0]
    num_diseased = len(df_diseased)
    num_normal = len(df_normal)
    assert num_diseased + num_normal == len(orig_df)

    if resample_method == "over":
        if resample_class == 0:
            new_num_normal = int(num_diseased / new_prop)
            print(f"Resampling normal samples from {num_normal} to {ne
            df_normal_rs = df_normal.sample(new_num_normal, replace=Tr
            resampled_df = pd.concat([df_normal_rs, df_diseased])
        else:
            # Resample the pneumonia class
            # new_num_diseased = int(new_prop * num_normal)
            # print(f"Resampling diseased samples from {num_diseased}
            # df_diseased_rs = df_diseased.sample(new_num_diseased, re
            # resampled_df = pd.concat([df_normal, df_diseased_rs])

```

```
target = df["Pneumonia"] == 1
rus = RandomOverSampler(random_state=0, sampling_strategy=
resampled_df, _ = rus.fit_resample(df, target)

resampled_df.sort_index(inplace=True)
print(f"New df proportion: {get_prop(resampled_df, column)}")
return resampled_df

# balance_proportion(dfs["MIMIC"]["train"], dfs["MIMIC"]["test"])
```

In [15]: `dfs["CXP"]["train"]`

Out[15]:

	subject_id	path
0	1	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
1	2	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
2	2	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
3	3	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
4	4	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
...	...	...
191222	64734	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
191223	64735	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
191225	64737	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
191227	64739	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...
191228	64740	C:\Nina\e-root\data\CheXpert\CheXpert-v1.0-sma...

153411 rows × 17 columns

```
In [10]: def balance_df_label(df, sampler, label_bal=0.05154780337262089, invert
        target = df["Pneumonia"] == (1 if not invert else 0)
        rus = sampler(random_state=42, sampling_strategy=label_bal if not
        res_df, _ = rus.fit_resample(df, target)

        print(f"Previous pneumonia prop: {get_prop(df)} with {len(df)} ins
        print(f"Resampled pneumonia prop: {get_prop(res_df)} with {len(res

        return res_df

# # uncomment this code if you want a balanced dataset
# print('Balancing...')
# mimic_balanced = balance_df_label(dfs["MIMIC"]["train"], RandomOverS
# cxp_balanced = balance_df_label(dfs["CXP"]["train"], RandomOverSampl
# print('Done.')

# # Balance the size of the two datasets
# n = len(cxp_balanced)
# mimic_balanced = mimic_balanced.sample(n)
```

```
In [11]: import warnings
warnings.filterwarnings('ignore')
```

## Calculating stats

```
In [12]: stat_rows = []
num_instances = []

disease_labels = ["Pneumonia", "Cardiomegaly", "Edema", "Effusion", "A
target_labels = disease_labels + ["Any", "No Finding"]
all_labels = target_labels + ["Num Instances"]

dfs2 = {}

for env in dfs:
    df = dfs[env]['all']
    df['Any'] = (df[disease_labels] > 0).any(axis=1).astype(int)

    # keep only every 30th sample for the dataset to reduce the size o
    # I am keeping the full dataset coe commented out to be able to ea
    df2 = df[df.index % 30 == 0]
    dfs2[env] = {}
    dfs2[env]['all'] = df2

    train_df, valid_df, test_df = split(df2)
    dfs2[env] = {
        'all': df env,
```



```

        'train': train_df,
        'val': valid_df,
        'test': test_df
    }

    totals = {}
    totals['Dataset'] = env
#     totals['Num Instances'] = len(df)
    totals['Num Instances'] = len(df2)
    num_instances.append(totals['Num Instances'])

    for label in target_labels:
#         if label in df.columns:
#             totals[label] = df[label].sum() / len(df)
        if label in df2.columns:
            totals[label] = df2[label].sum() / len(df2)
        else:
            totals[label] = 0.0

    stat_rows.append(totals)

stat_df = pd.DataFrame(stat_rows)
stat_df.set_index('Dataset', inplace=True)

ordered_cols = all_labels
stat_df = stat_df[ordered_cols]

transposed_stat_df = stat_df.T

styled_transposed_stat_df = transposed_stat_df.style.apply(
    lambda x: ["background-color: lightblue" if x.name != 'Num Instances' else ''],
    axis=1
).background_gradient(cmap='Blues', subset=pd.IndexSlice[target_labels, :])
styled_transposed_stat_df = styled_transposed_stat_df.format("{:.2%}",
styled_transposed_stat_df = styled_transposed_stat_df.format("{:,.0f}")

styled_transposed_stat_df

```

Out[12]:

Dataset	MIMIC	CXP
Pneumonia	6.87%	2.68%
Cardiomegaly	16.62%	12.27%
Edema	11.70%	25.77%
Effusion	22.94%	40.20%
Atelectasis	20.42%	16.02%
Pneumothorax	4.23%	9.33%

<b>Consolidation</b>	4.68%	6.59%
<b>Any</b>	50.51%	71.14%
<b>No Finding</b>	35.01%	8.69%
<b>Num Instances</b>	7,690	6,375

Here is the table from the article for comparison:

Table 1: Total number of instances and disease prevalence in each dataset.

Target Label	MIMIC	CXP	NIH	PAD
Pneumonia	6.82%	2.43%	1.31%	4.84%
Cardiomegaly	17.05%	12.38%	2.51%	9.15%
Edema	11.83%	26.01%	2.11%	1.23%
Effusion	23.18%	40.28%	11.94%	5.99%
Atelectasis	20.11%	15.47%	10.33%	5.50%
Pneumothorax	4.19%	9.25%	4.66%	0.31%
Consolidation	4.67%	6.81%	4.19%	1.56%
Any	50.73%	70.35%	28.04%	23.03%
No Finding	34.76%	8.98%	53.65%	36.12%
Num Instances	243k	192k	113k	100k

Looks like the distribution of the labels in the original dataset, while not the same, still is close enough.

## Citation to the original paper

- Rhys Compton; Lily Zhang; Aahlad Puli; Rajesh Ranganath, When More is Less: Incorporating Additional Datasets Can Hurt Performance By Introducing Spurious Correlations, arXiv preprint, 2023-08-09, Accepted at MLHC 2023, doi: [10.48550/arXiv.2308.04431](https://doi.org/10.48550/arXiv.2308.04431) (<https://doi.org/10.48550/arXiv.2308.04431>)

## Original paper repo

- [ood-generalization](https://github.com/basedrhys/ood-generalization/tree/master) (<https://github.com/basedrhys/ood-generalization/tree/master>)

## Model

The model includes the model definition which usually is a class, model training, and other necessary parts.

## Model architecture

In the article, the authors use the same model architecture as Zhang et al. (2021): a **DenseNet-121** network (Huang et al., 2017) initialized with pre-trained weights from ImageNet (Deng et al., 2009). The final layer is replaced with a **two-output linear layer** (for binary classification). For simplicity, the authors only consider binary disease classification.

## Model Training

For training the network, all images are resized to **224 × 224** and normalized to the ImageNet (Deng et al., 2009) mean and standard deviation.

During training, the following image augmentations are applied:

- random horizontal flip
- random rotation up to 10 degrees
- a crop of random size (75% - 100%) and aspect ratio (3/4 to 4/3)

All runs use **Adam** with **lr = 1e-5** and **batch size = 128**, which was found to be a performant configuration in early tuning ((Zhang et al., 2021) use lr = 5e-4 and batch size = 32).

All test results are obtained using the optimal model found during training as measured by the highest validation macro-F1 score (following (Fiorillo et al., 2021; Berenguer et al., 2022)) as it gives a robust ranking of model performance under imbalanced labels.

In [13]: *This is the model defined and provided by the authors of the article. While they are using densenet 121 for the article, the provided model*

```

class EmbModel(nn.Module):
    # I had to add the num_labels parameter to reduce the resulting resp
    def __init__(self, emb_type, feature_size_override, pretrain, concat
        super().__init__()
        self.emb_type = emb_type
        self.pretrain = pretrain
        self.concat_features = concat_features
        self.num_labels = num_labels

        assert emb_type in ["densenet121", "densenet201", "resnet"], f"

        if emb_type == 'densenet121':

```

```

        model = models.densenet121()
        self.encoder = nn.Sequential(*list(model.children())[:-1]) #
        self.emb_dim = model.classifier.in_features
    elif emb_type == 'densenet201':
        model = models.densenet201()
        self.encoder = nn.Sequential(*list(model.children())[:-1]) #
        self.emb_dim = model.classifier.in_features
    elif emb_type == 'resnet':
        model = models.resnet50()
        self.encoder = nn.Sequential(*list(model.children())[:-1])
        self.emb_dim = list(model.children())[-1].in_features

    print("\nEmb Dim:")
    print(self.emb_dim)

    if feature_size_override:
        print(f"Manually setting output dim to {feature_size_override}")
        self.emb_dim = feature_size_override
        print(self.emb_dim)

    self.n_outputs = self.emb_dim + concat_features
    self.final_layer = nn.Linear(self.n_outputs, self.num_labels)

    nn.init.kaiming_normal_(self.final_layer.weight, mode='fan_out',

def forward(self, inp):
    if isinstance(inp, dict): # dict with image and additional featu
        x = inp['img']
        concat = inp['concat']
        assert(concat.shape[-1] == self.concat_features)
    else: # tensor image
        assert(self.concat_features == 0)
        x = inp

    x = self.encoder(x).squeeze(-1).squeeze(-1)
    if "densenet" in self.emb_type:
        x = F.relu(x)
        x = F.avg_pool2d(x, kernel_size = 7).view(x.size(0), -1)

    if isinstance(inp, dict):
        x = torch.cat([x, concat], dim = -1)

    x = self.final_layer(x)
    return x

```

# Training

I wasn't able to run the training code provided by the authors of the article - the setup didn't work for me neither on my MacBook Pro laptop, nor on my husband's Windows 10 gaming computer.

To proceed, I instead wrote my own training code using the standard approach learned in class and homeworks.

## Hyperparameters used

- Model: densenet121
- Number of epochs for each model trainig: 10
- Hidden size: 1024 since I am setting the `feature_size_override` to 1024
- Batch size: 128
- Learning rate:  $1e-5$
- Optimizer: Adam

# Computational requirements

It is possible to run this code on a CPU with minor modifications. However, since I moved to another computer with GPU, some portions of this notebook send the computation to Cuda directly (todo: rewrite so it checks for Cuda and sends to the appropriate device).

## Hardware and software

- AMD Ryzen 7 7800X4D 8-Core Processor (4.20 GHz)
- 64 GB RAM
- NVMe Samsung SSD 970 EVO Plus - 1TB
- Windows 10 64-bit

## Training requirements

I was not able to train on the full dataset since even on GPU one epoch of one model was running for 2-4 hours depending on the number of batches. This would require roughly 16 days to finish the whole training.

Initially, I attempted to run the training on the full dataset, but a number of circumstances (out of memory, kernel panic, random automatic Windows updates, power down, kids getting to the computer and switching the power supply off) proved that the expectation to run the training continuously for days to be completely unrealistic.

Instead, I modified my dataset to pick every 30th entry and discard the rest. As a result, I was able to run the training multiple times with different parameters when needed, both on balanced and unbalanced datasets.

- Average epoch running time: 6min
- Average time to complete all training: 12h
- Total number of attempts: 200+

## Creating a data loader

The authors of the article have a script to load the data in different configurations. I wasn't able to make it work because of the errors, so instead I am partially reusing it and creating my own Dataset class and a data loader.

In [14]:

```
ImageFile.LOAD_TRUNCATED_IMAGES = True # I was getting errors during t
```

In [15]:

```
class MultiEnvDataset(Dataset):
    def __init__(self, dataframes, subset='train', envs=None, transform=None):
        """
        Initializes the dataset with data from multiple environments and labels.
        :param dataframes: A dictionary with environment keys, each containing a DataFrame.
        :param subset: The subset to load ('train', 'val', or 'test').
        :param envs: A list of environment names to include. If None, all environments are used.
        :param transform: PyTorch transforms to apply to the images.
        """
        if envs is None:
            envs = list(dataframes.keys())

        self.data = pd.concat([dataframes[env][subset] for env in envs])

        self.label_columns = ["No Finding", "Atelectasis", "Cardiomegaly",
                              "Pneumothorax", "Consolidation", "Edema"]

        self.transform = transform

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        img_path = self.data.iloc[idx]['path']
        image = Image.open(img_path).convert('RGB') # Converts to RGB

        if self.transform:
            image = self.transform(image)

        labels = Tensor(self.data.iloc[idx][self.label_columns].values)
        if torch.isnan(labels).any():
            raise ValueError("NaN values found in labels")

        return image, labels
```

```
In [16]: transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
```

```
In [17]: envs_list = ["CXP"], ["MIMIC"], ["NIH"], ["PAD"], ["CXP", "NIH"], ["CX
env_list_map = {
    "cxp": 0,
    "mimic": 1,
    "nih": 2,
    "pad": 3,
    "cxp_nih": 4,
    "cxp_pad": 5,
    "mimic_cxp": 6,
    "mimic_nih": 7,
    "mimic_pad": 8,
    "nih_pad": 9,
    "cxp_mimic_nih_pad": 10,
}

# a few functions to simplify getting the right names of the dataset c
def get_dataset_index(env_name):
    return env_list_map[env_name]

def get_env_from_list(env_name):
    return envs_list[env_list_map[env_name]]
```

```
In [18]: datasets = []
for env in envs_list:
    elem = {
        "env": env
    }
    for subset in ["train", "val"]:
        elem[subset] = {}
#         elem[subset]["dataset"] = MultiEnvDataset(dfs, subset=subset
        elem[subset]["dataset"] = MultiEnvDataset(dfs2, subset=subset,
        elem[subset]["loader"] = torch.utils.data.DataLoader(elem[subs
    datasets.append(elem)
print("Done.")
```

Done.

## Metrics

```
In [19]: loss_func = nn.BCEWithLogitsLoss()

max_batches = 10

def calculate_accuacies(outputs, labels):
    predictions = torch.sigmoid(outputs) > 0.5
    predictions = predictions.to(labels.device)
    correct_pred = (predictions == labels)
```



```

        accuracies = correct_pred.float().mean(axis=0)
        return accuracies

def calculate_f1(outputs, labels):
    predictions = torch.sigmoid(outputs) > 0.5
    predictions = predictions.to(labels.device)

    predictions = predictions.detach().cpu().numpy()
    labels = labels.detach().cpu().numpy()

    f1 = f1_score(labels, predictions, average=None)
    return f1

def train_model_one_epoch(model, train_loader, loss_func, optimizer):
    print("Starting training...")
    start = datetime.now()
    prev = start
    model.train()
    running_loss = 0
    total_accuracy = []
    total_f1_scores = []

    print('number of batches:', len(train_loader))
    for batch, (inputs, labels) in enumerate(train_loader):
        inputs = inputs.cuda()

        optimizer.zero_grad()
        outputs = model(inputs)

        if torch.isnan(outputs).any():
            raise ValueError("NaN detected in model outputs")

        loss = loss_func(outputs, labels)
        if torch.isnan(loss).any():
            raise ValueError("NaN detected in loss computation")

        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.)
        optimizer.step()
        running_loss += loss.item() * inputs.size(0)

        accuracies = calculate_accuracies(outputs, labels)
        f1_scores = calculate_f1(outputs, labels)
        total_accuracy.append(accuracies)
        total_f1_scores.append(f1_scores)

        if batch % 100 == 0:
            mid = datetime.now()
            print("time passed from the beginning", mid-start)
            print('batch', batch + 1, 'time passed:', mid-prev)
            prev = mid

```

```

prev = min

epoch_loss = running_loss / len(train_loader.dataset)
end = datetime.now()
print("epoch done in", end-start, "number of batches:", batch)
epoch_accuracy = torch.stack(total_accuracy).mean(dim=0)
epoch_f1 = torch.tensor(total_f1_scores).mean(dim=0)
return epoch_loss, epoch_accuracy, epoch_f1

def validate_model(model, val_loader, loss_func):
    model.eval()
    running_loss = 0
    total_accuracy = []
    total_f1_scores = []
    with torch.no_grad():
        for inputs, labels in val_loader:
            inputs = inputs.cuda()
            outputs = model(inputs)
            loss = loss_func(outputs, labels)
            running_loss += loss.item() * inputs.size(0)

            accuracies = calculate_accuracies(outputs, labels)
            f1_scores = calculate_f1(outputs, labels)
            total_accuracy.append(accuracies)
            total_f1_scores.append(f1_scores)

    epoch_loss = running_loss / len(val_loader.dataset)
    epoch_accuracy = torch.stack(total_accuracy).mean(dim=0)
    epoch_f1 = torch.tensor(total_f1_scores).mean(dim=0)
    return epoch_loss, epoch_accuracy, epoch_f1

```

```

In [20]: def saveModel(model, env=None):
    now = datetime.now()
    dt_string = now.strftime("%d-%m-%Y-%H-%M-%S")

    model_file_name = "model/model-snapshot-"
    if (env):
        model_file_name += "env_" + "_".join(env) + "_"
    model_file_name += dt_string + ".pth"

    torch.save(model.state_dict(), model_file_name)

```

In [ ]:

```

%%script false --no-raise-error
# skipping the training since I already ran it in different variations

num_epoch = 10

metrics_df = pd.DataFrame(columns=["env", "epoch", "train_loss", "valid_loss", "train_accuracy", "valid_accuracy", "train_f1", "valid_f1"])

for dataset in datasets:
    model = EmbModel(emb_type="densenet121", feature_size_override=1024)
    model.cuda()
    model.train()
    optimizer = torch.optim.Adam(model.parameters(), lr=1e-5)
    print("Processing dataset env:", dataset["env"])
    for i in range(num_epoch):
        train_loss, train_accuracy, train_f1 = train_model_one_epoch(model, dataset["env"], optimizer)
        valid_loss, valid_accuracy, valid_f1 = validate_model(model, dataset["env"])

        print("Epoch: %.2f, Train Loss: %.2f, Validation Loss: %.2f" % (i+1, train_loss, valid_loss))

        # Convert tensors to CPU for DataFrame update
        train_accuracy = train_accuracy.cpu().numpy()
        valid_accuracy = valid_accuracy.cpu().numpy()
        train_f1 = train_f1.cpu().numpy()
        valid_f1 = valid_f1.cpu().numpy()

        # Append metrics to DataFrame
        metrics_df = metrics_df.append({
            "env": dataset["env"],
            "epoch": i + 1,
            "train_loss": train_loss,
            "valid_loss": valid_loss,
            "train_accuracy": np.mean(train_accuracy),
            "valid_accuracy": np.mean(valid_accuracy),
            "worst_train_accuracy": np.min(train_accuracy),
            "worst_valid_accuracy": np.min(valid_accuracy),
            "train_f1": np.mean(train_f1),
            "valid_f1": np.mean(valid_f1)
        }, ignore_index=True)

    saveModel(model, env=dataset["env"])
print("All done.")

```

Emb Dim:  
 1024  
 Manually setting output dim to 1024  
 1024  
 Processing dataset env: ['CXP']  
 Starting training...  
 number of batches: 40

```
time passed from the beginning 0:00:04.679461
batch 1 time passed: 0:00:04.679461
epoch done in 0:00:31.729330 number of batches: 39
Epoch: 1.00, Train Loss: 1.28, Validation Loss: 1.17
Starting training...
number of batches: 40
time passed from the beginning 0:00:00.599449
batch 1 time passed: 0:00:00.599449
epoch done in 0:00:23.961185 number of batches: 39
Epoch: 2.00, Train Loss: 0.61, Validation Loss: 0.41
Starting training...
number of batches: 40
time passed from the beginning 0:00:00.617424
batch 1 time passed: 0:00:00.617424
epoch done in 0:00:23.947628 number of batches: 39
Epoch: 3.00, Train Loss: 0.39, Validation Loss: 0.38
Starting training...
number of batches: 40
time passed from the beginning 0:00:00.588598
batch 1 time passed: 0:00:00.588598
epoch done in 0:00:23.979013 number of batches: 39
Epoch: 4.00, Train Loss: 0.35, Validation Loss: 0.37
Starting training...
number of batches: 40
time passed from the beginning 0:00:00.595298
batch 1 time passed: 0:00:00.595298
epoch done in 0:00:24.064484 number of batches: 39
Epoch: 5.00, Train Loss: 0.33, Validation Loss: 0.37
Starting training...
number of batches: 40
time passed from the beginning 0:00:00.590078
batch 1 time passed: 0:00:00.590078
epoch done in 0:00:23.999334 number of batches: 39
Epoch: 6.00, Train Loss: 0.32, Validation Loss: 0.36
Starting training...
number of batches: 40
time passed from the beginning 0:00:00.612733
batch 1 time passed: 0:00:00.612733
epoch done in 0:00:23.986761 number of batches: 39
Epoch: 7.00, Train Loss: 0.30, Validation Loss: 0.36
Starting training...
number of batches: 40
time passed from the beginning 0:00:00.608278
batch 1 time passed: 0:00:00.608278
epoch done in 0:00:24.037447 number of batches: 39
Epoch: 8.00, Train Loss: 0.29, Validation Loss: 0.36
Starting training...
number of batches: 40
time passed from the beginning 0:00:00.609474
batch 1 time passed: 0:00:00.609474
```

```
epoch done in 0:00:24.047149 number of batches: 39
Epoch: 9.00, Train Loss: 0.28, Validation Loss: 0.36
Starting training...
number of batches: 40
time passed from the beginning 0:00:00.603249
batch 1 time passed: 0:00:00.603249
epoch done in 0:00:24.086203 number of batches: 39
Epoch: 10.00, Train Loss: 0.27, Validation Loss: 0.36
```

```
Emb Dim:
1024
Manually setting output dim to 1024
1024
Processing dataset env: ['MIMIC']
Starting training...
number of batches: 49
time passed from the beginning 0:00:05.692402
batch 1 time passed: 0:00:05.692402
epoch done in 0:04:31.752617 number of batches: 48
Epoch: 1.00, Train Loss: 0.95, Validation Loss: 0.52
Starting training...
number of batches: 49
time passed from the beginning 0:00:05.497830
batch 1 time passed: 0:00:05.497830
epoch done in 0:04:21.522990 number of batches: 48
Epoch: 2.00, Train Loss: 0.43, Validation Loss: 0.38
Starting training...
number of batches: 49
time passed from the beginning 0:00:05.571090
batch 1 time passed: 0:00:05.571090
```

```
In [582]: %%script false --no-raise-error
# skipping this cell since I already ran this.

# Save DataFrame to CSV

df_now = datetime.now()
df_dt_string = df_now.strftime("%d-%m-%Y-%H-%M-%S")
metrics_df.to_csv(f"stats/{df_dt_string}_training_metrics.csv", index=

Couldn't find program: 'false'
```

## Validating the saved models and visualizing results

```

In [21]: def predict(model, val_loader, device='cuda'):
    model.eval()
    model.to(device)
    all_preds = []
    all_preds_raw = []
    all_labels = []

    print('Started prediction validation')
    print('Number of batches:', len(val_loader))
    predict_start_time = datetime.now()
    with torch.no_grad():
        for batch, (images, labels) in enumerate(val_loader):
            print("Batch number:", batch+1, "of", len(val_loader))
            images = images.to(device)
            labels = labels.to(device)

            outputs = model(images)
            probabilities = torch.sigmoid(outputs)

            preds = (probabilities > 0.5)

            any_disease = torch.any(preds[:, 1:], dim=1, keepdim=True)
            any_probability = torch.max(probabilities[:, 1:], dim=1, keepdim=True)
            any_label = torch.any(labels[:, 1:], dim=1, keepdim=True)

            preds = torch.cat((preds, any_disease), dim=1)
            probabilities = torch.cat((probabilities, any_probability), dim=1)
            labels = torch.cat((labels, any_label), dim=1)

            all_preds_raw.append(probabilities.cpu().numpy())
            all_preds.append(preds.cpu().numpy())
            all_labels.append(labels.cpu().numpy())

    all_preds = np.vstack(all_preds)
    all_preds_raw = np.vstack(all_preds_raw)
    all_labels = np.vstack(all_labels)

    predict_end_time = datetime.now()

    print('Done.')
    print('Prediction took:', predict_end_time-predict_start_time)

    return all_preds, all_preds_raw, all_labels

```

```

In [22]: def calculate_per_label_accuracy(predictions, labels):
    accuracies = {}
    num_labels = labels.shape[1]

    for i in range(num_labels):

```

```

        label_preds = predictions[:, i]
        label_true = labels[:, i]
        accuracies[target_labels[i]] = accuracy_score(label_true, label_preds)
    return accuracies

def calculate_extended_metrics(predictions, labels, target_labels):
    metrics = {}
    num_labels = labels.shape[1]

    for i in range(num_labels):
        label_preds = predictions[:, i]
        label_true = labels[:, i]
        cm = confusion_matrix(label_true, label_preds, labels=[0, 1])

        TP = TN = FP = FN = 0

        #check the shape of the confusion matrix
        if cm.shape == (2, 2):
            TN, FP, FN, TP = cm.ravel()
        elif cm.shape == (1, 1):
            if label_true[0] == 0: #all negatives
                TN = cm[0, 0]
            else: #all positives
                TP = cm[0, 0]

        TPR = TP / (TP + FN) if TP + FN != 0 else 0 # True positives
        SPC = TN / (TN + FP) if TN + FP != 0 else 0 # True negatives
        ACC = (TP + TN) / (TP + TN + FP + FN) if (TP + TN + FP + FN) != 0 else 0

        metrics[target_labels[i]] = {
            'Accuracy': ACC,
            'Positive Accuracy (TPR)': TPR,
            'Negative Accuracy (SPC)': SPC
        }

    return metrics

def calculate_stats(predictions, probabilities, labels, source):
    # Calculate overall accuracy
    accuracy = accuracy_score(labels, predictions)
    print(f"Overall Accuracy: {accuracy:.2%}")

    # Detailed classification report for each disease label
    report = classification_report(labels, predictions, target_names=target_labels,
                                   display_labels=labels)
    report_df = pd.DataFrame(report).transpose()
    report_df['source'] = source
    print("Classification Report:")
    print(report_df)

    extended_metrics = calculate_extended_metrics(predictions, labels, target_labels)

```

```

extended_metrics = calculate_extended_metrics(predictions, labels,
metrics_df = pd.DataFrame.from_dict(extended_metrics, orient='index')

accuracy_index = ['Accuracy', 'Positive Accuracy (TPR)', 'Negative

metrics_df.columns = ['Label'] + accuracy_index
metrics_df['source'] = source

metrics_df['Min Metric Value'] = metrics_df[accuracy_index].min(axis=1)
metrics_df['Min Metric Name'] = metrics_df[accuracy_index].idxmin(axis=1)
worst_label = metrics_df.loc[metrics_df['Min Metric Value'].idxmin()]
worst_metric_value = metrics_df.loc[metrics_df['Min Metric Value'].idxmin()]['Min Metric Value']
worst_metric_name = metrics_df.loc[metrics_df['Min Metric Value'].idxmin()]['Min Metric Name']

print(metrics_df)
print(f"Worst Performing Label: {worst_label} with a minimum metric value of {worst_metric_value}")

return report_df, metrics_df

```

```

In [23]: # repeating the labels code so I don't have to re-run the cell way above
target_labels = ["No Finding", "Atelectasis", "Cardiomegaly", "Effusion", "Pneumothorax"]

device = 'cuda' if torch.cuda.is_available() else 'cpu'

```

## Loading trained models

I separated the loading code into different cells rather than had a cycle in one, to be able to pick and choose which parts I run. The model names are hardcoded with the ones currently in the repo - if you are running the training code above, the new snapshots will be created and the file names should be updated below.

## Initializing data loaders



In [24]:

```
data_loader_MIMIC = datasets[get_dataset_index("mimic")]["val"]["loader"]
data_loader_MIMIC_PAD = datasets[get_dataset_index("mimic_pad")]["val"]["loader"]
data_loader_MIMIC_NIH = datasets[get_dataset_index("mimic_nih")]["val"]["loader"]
data_loader_MIMIC_CXP = datasets[get_dataset_index("mimic_cxp")]["val"]["loader"]
data_loader_CXP = datasets[get_dataset_index("cxp")]["val"]["loader"]
data_loader_CXP_NIH = datasets[get_dataset_index("cxp_nih")]["val"]["loader"]
data_loader_CXP_PAD = datasets[get_dataset_index("cxp_pad")]["val"]["loader"]
data_loader_NIH = datasets[get_dataset_index("nih")]["val"]["loader"]
data_loader_NIH_PAD = datasets[get_dataset_index("nih_pad")]["val"]["loader"]
data_loader_PAD = datasets[get_dataset_index("pad")]["val"]["loader"]
data_loader_CXP_MIMIC_NIH_PAD = datasets[get_dataset_index("cxp_mimic_nih_pad")]["val"]["loader"]
```

## MIMIC only

```
In [25]: print("Model trained on MIMIC")
model_MIMIC = EmbModel(emb_type="densenet121", feature_size_override=1
model_MIMIC.load_state_dict(torch.load("model/balanced/model-snapshot-
model_MIMIC.eval()
padding=(1, 1), bias=False)
)
(denselayer3): _DenseLayer(
  (norm1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=T
rue, track_running_stats=True)
  (relu1): ReLU(inplace=True)
  (conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1)
, bias=False)
  (norm2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=T
rue, track_running_stats=True)
  (relu2): ReLU(inplace=True)
  (conv2): Conv2d(128, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
)
(denselayer4): _DenseLayer(
  (norm1): BatchNorm2d(160, eps=1e-05, momentum=0.1, affine=T
rue, track_running_stats=True)
  (relu1): ReLU(inplace=True)
  (conv1): Conv2d(160, 128, kernel_size=(1, 1), stride=(1, 1)
. bias=False)
```

```
In [26]: mimic_predictions, mimic_probabilities, mimic_labels = predict(model_M
print('labels', mimic_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 9
Batch number: 1 of 9
Batch number: 2 of 9
Batch number: 3 of 9
Batch number: 4 of 9
Batch number: 5 of 9
Batch number: 6 of 9
Batch number: 7 of 9
Batch number: 8 of 9
Batch number: 9 of 9
Done.
Prediction took: 0:00:31.219591
labels [[1. 0. 0. ... 0. 0. 0.]
[0. 0. 0. ... 0. 0. 0.]
[0. 1. 0. ... 0. 0. 1.]
...
[1. 0. 0. ... 0. 0. 0.]
[1. 0. 0. ... 0. 0. 0.]
[1. 0. 0. ... 0. 0. 0.]
```

In [27]: `mimic_report, mimic_accuracy = calculate_stats(mimic_predictions, mimic_labels)`

Overall Accuracy: 20.97%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.661376	0.326371	0.437063	383.0	mimic
Atelectasis	0.269231	0.080000	0.123348	175.0	mimic
Cardiomegaly	0.104377	0.203947	0.138085	152.0	mimic
Effusion	0.466667	0.242775	0.319392	173.0	mimic
Pneumonia	0.076923	0.016393	0.027027	61.0	mimic
Pneumothorax	0.000000	0.000000	0.000000	29.0	mimic
Consolidation	0.000000	0.000000	0.000000	39.0	mimic
Edema	0.161290	0.067568	0.095238	74.0	mimic
Any	0.392111	0.380631	0.386286	444.0	mimic
micro avg	0.350861	0.252941	0.293961	1530.0	mimic
macro avg	0.236886	0.146409	0.169604	1530.0	mimic
weighted avg	0.384148	0.252941	0.291132	1530.0	mimic
samples avg	0.224625	0.213468	0.214888	1530.0	mimic

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy
0	No Finding	0.698502	0.326371	0
1	Atelectasis	0.813670	0.080000	0
2	Cardiomegaly	0.637640	0.203947	0
3	Effusion	0.832397	0.242775	0
4	Pneumonia	0.932584	0.016393	0
5	Pneumothorax	0.972846	0.000000	1
6	Consolidation	0.963483	0.000000	1
7	Edema	0.911049	0.067568	0
8	Any	0.497191	0.380631	0

	source	Min Metric Value	Min Metric Name
0	mimic	0.326371	Positive Accuracy (TPR)
1	mimic	0.080000	Positive Accuracy (TPR)
2	mimic	0.203947	Positive Accuracy (TPR)
3	mimic	0.242775	Positive Accuracy (TPR)
4	mimic	0.016393	Positive Accuracy (TPR)
5	mimic	0.000000	Positive Accuracy (TPR)
6	mimic	0.000000	Positive Accuracy (TPR)
7	mimic	0.067568	Positive Accuracy (TPR)
8	mimic	0.380631	Positive Accuracy (TPR)

Worst Performing Label: Pneumothorax with a minimum metric value of 0

.00% (Metric: Positive Accuracy (IPK))

```
In [28]: def combine_with_existing(combined_accuracy, combined_report, add_accu
        combined_accuracy = pd.concat([combined_accuracy, add_accuracy], i
        combined_report = pd.concat([combined_report, add_report], ignore_

        print(combined_accuracy)
        print(combined_report)

        return combined_accuracy, combined_report
```

## PAD only

```
In [29]: print("Model trained on PAD")
model_PAD = EmbModel(emb_type="densenet121", feature_size_override=1024)
model_PAD.load_state_dict(torch.load("model/balanced/model-snapshot-en
model_PAD.eval()
```

Model trained on PAD

Emb Dim:  
1024  
Manually setting output dim to 1024  
1024

```
Out[29]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True)
```

```
In [30]: pad_predictions, pad_probabilities, pad_labels = predict(model_PAD, da
print('labels', pad_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 6
Batch number: 1 of 6
Batch number: 2 of 6
Batch number: 3 of 6
Batch number: 4 of 6
Batch number: 5 of 6
Batch number: 6 of 6
Done.
Prediction took: 0:00:05.541071
labels [[1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 1. 0. 1.]
 [0. 0. 0. ... 0. 0. 0.]]
Done
```

```
In [31]: pad_report, pad_accuracy = calculate_stats(pad_predictions, pad_probab
```

```
Overall Accuracy: 30.59%
```

```
Classification Report:
```

	precision	recall	f1-score	support	source
No Finding	0.622807	0.223975	0.329466	317.0	pad
Atelectasis	0.107143	0.120000	0.113208	50.0	pad
Cardiomegaly	0.052632	0.102564	0.069565	39.0	pad
Effusion	0.250000	0.083333	0.125000	72.0	pad
Pneumonia	0.000000	0.000000	0.000000	21.0	pad
Pneumothorax	0.000000	0.000000	0.000000	6.0	pad
Consolidation	0.000000	0.000000	0.000000	26.0	pad
Edema	0.000000	0.000000	0.000000	12.0	pad
Any	0.363014	0.291209	0.323171	182.0	pad
micro avg	0.303030	0.193103	0.235889	725.0	pad
macro avg	0.155066	0.091231	0.106712	725.0	pad
weighted avg	0.398494	0.193103	0.249147	725.0	pad
samples avg	0.123425	0.146825	0.127541	725.0	pad

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy
y (SPC) \				
0	No Finding	0.582973	0.223975	0
.885638				
1	Atelectasis	0.864358	0.120000	0
.922240				
2	Cardiomegaly	0.845599	0.102564	0
.889908				

3	Effusion	0.878788	0.083333	0
.971014				
4	Pneumonia	0.963925	0.000000	0
.994048				
5	Pneumothorax	0.936508	0.000000	0
.944687				
6	Consolidation	0.958153	0.000000	0
.995502				
7	Edema	0.981241	0.000000	0
.998532				
8	Any	0.679654	0.291209	0
.818004				

	source	Min Metric Value	Min Metric Name
0	pad	0.223975	Positive Accuracy (TPR)
1	pad	0.120000	Positive Accuracy (TPR)
2	pad	0.102564	Positive Accuracy (TPR)
3	pad	0.083333	Positive Accuracy (TPR)
4	pad	0.000000	Positive Accuracy (TPR)
5	pad	0.000000	Positive Accuracy (TPR)
6	pad	0.000000	Positive Accuracy (TPR)
7	pad	0.000000	Positive Accuracy (TPR)
8	pad	0.291209	Positive Accuracy (TPR)

Worst Performing Label: Pneumonia with a minimum metric value of 0.00 % (Metric: Positive Accuracy (TPR))

In [32]: combined\_accuracy, combined\_report = combine\_with\_existing(mimic\_report

	precision	recall	f1-score	support	source
0	0.661376	0.326371	0.437063	383.0	mimic
1	0.269231	0.080000	0.123348	175.0	mimic
2	0.104377	0.203947	0.138085	152.0	mimic
3	0.466667	0.242775	0.319392	173.0	mimic
4	0.076923	0.016393	0.027027	61.0	mimic
5	0.000000	0.000000	0.000000	29.0	mimic
6	0.000000	0.000000	0.000000	39.0	mimic
7	0.161290	0.067568	0.095238	74.0	mimic
8	0.392111	0.380631	0.386286	444.0	mimic
9	0.350861	0.252941	0.293961	1530.0	mimic
10	0.236886	0.146409	0.169604	1530.0	mimic
11	0.384148	0.252941	0.291132	1530.0	mimic
12	0.224625	0.213468	0.214888	1530.0	mimic
13	0.622807	0.223975	0.329466	317.0	pad
14	0.107143	0.120000	0.113208	50.0	pad
15	0.052632	0.102564	0.069565	39.0	pad
16	0.250000	0.083333	0.125000	72.0	pad
17	0.000000	0.000000	0.000000	21.0	pad
18	0.000000	0.000000	0.000000	6.0	pad
19	0.000000	0.000000	0.000000	26.0	pad
20	0.000000	0.000000	0.000000	12.0	pad

21	0.363014	0.291209	0.323171	182.0	pad
22	0.303030	0.193103	0.235889	725.0	pad
23	0.155066	0.091231	0.106712	725.0	pad
24	0.398494	0.193103	0.249147	725.0	pad
25	0.123425	0.146825	0.127541	725.0	pad
	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy	
cy (SPC) \					
0	No Finding	0.698502	0.326371		
0.906569					
1	Atelectasis	0.813670	0.080000		
0.957447					
2	Cardiomegaly	0.637640	0.203947		
0.709607					
3	Effusion	0.832397	0.242775		
0.946369					
4	Pneumonia	0.932584	0.016393		
0.988083					
5	Pneumothorax	0.972846	0.000000		
1.000000					
6	Consolidation	0.963483	0.000000		
1.000000					
7	Edema	0.911049	0.067568		
0.973843					
8	Any	0.497191	0.380631		
0.580128					
9	No Finding	0.582973	0.223975		
0.885638					
10	Atelectasis	0.864358	0.120000		
0.922240					
11	Cardiomegaly	0.845599	0.102564		
0.889908					
12	Effusion	0.878788	0.083333		
0.971014					
13	Pneumonia	0.963925	0.000000		
0.994048					
14	Pneumothorax	0.936508	0.000000		
0.944687					
15	Consolidation	0.958153	0.000000		
0.995502					
16	Edema	0.981241	0.000000		
0.998532					
17	Any	0.679654	0.291209		
0.818004					

	source	Min Metric Value	Min Metric Name
0	mimic	0.326371	Positive Accuracy (TPR)
1	mimic	0.080000	Positive Accuracy (TPR)
2	mimic	0.203947	Positive Accuracy (TPR)
3	mimic	0.242775	Positive Accuracy (TPR)
4	mimic	0.016393	Positive Accuracy (TPR)
5	mimic	0.000000	Positive Accuracy (TPR)
6	mimic	0.000000	Positive Accuracy (TPR)
7	mimic	0.067568	Positive Accuracy (TPR)
8	mimic	0.380631	Positive Accuracy (TPR)
9	mimic	0.223975	Positive Accuracy (TPR)
10	mimic	0.120000	Positive Accuracy (TPR)
11	mimic	0.102564	Positive Accuracy (TPR)
12	mimic	0.083333	Positive Accuracy (TPR)
13	mimic	0.000000	Positive Accuracy (TPR)
14	mimic	0.000000	Positive Accuracy (TPR)
15	mimic	0.000000	Positive Accuracy (TPR)
16	mimic	0.000000	Positive Accuracy (TPR)
17	mimic	0.291209	Positive Accuracy (TPR)

5	mimic	0.000000	Positive Accuracy (IPR)
6	mimic	0.000000	Positive Accuracy (TPR)
7	mimic	0.067568	Positive Accuracy (TPR)
8	mimic	0.380631	Positive Accuracy (TPR)
9	pad	0.223975	Positive Accuracy (TPR)
10	pad	0.120000	Positive Accuracy (TPR)
11	pad	0.102564	Positive Accuracy (TPR)
12	pad	0.083333	Positive Accuracy (TPR)
13	pad	0.000000	Positive Accuracy (TPR)
14	pad	0.000000	Positive Accuracy (TPR)
15	pad	0.000000	Positive Accuracy (TPR)
16	pad	0.000000	Positive Accuracy (TPR)
17	pad	0.291209	Positive Accuracy (TPR)

## CXP only

```
In [33]: print("Model trained on CXP")
model_CXP = EmbModel(emb_type="densenet121", feature_size_override=1024)
model_CXP.load_state_dict(torch.load("model/balanced/model-snapshot-emb.pkl"))
model_CXP.eval()
```

Model trained on CXP

Emb Dim:

1024

Manually setting output dim to 1024

1024

```
Out[33]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True)
```



```
In [34]: exp_predictions, exp_probabilities, exp_labels = predict(model_CXP, da

print('predictions', exp_predictions)
print('labels', exp_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 8
Batch number: 1 of 8
Batch number: 2 of 8
Batch number: 3 of 8
Batch number: 4 of 8
Batch number: 5 of 8
Batch number: 6 of 8
Batch number: 7 of 8
Batch number: 8 of 8
Done.
Prediction took: 0:00:02.607437
predictions [[False False False ... False False False]
 [False  True False ... False False  True]
 [False  True False ... False False  True]
 ...
 [False False False ... False False False]
 [False False False ... False  True  True]
 [False False False ... False False False]]
labels [[1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 1.]
 ...
 [0. 0. 0. ... 0. 0. 1.]
 [0. 0. 0. ... 0. 1. 1.]
 [0. 0. 0. ... 0. 0. 1.]]
Done
```

```
In [35]: exp_report, exp_accuracy = calculate_stats(exp_predictions, exp_probab
```

```
Overall Accuracy: 19.10%
Classification Report:
```

	precision	recall	f1-score	support	source
No Finding	0.000000	0.000000	0.000000	184.0	cxp
Atelectasis	0.059633	0.115044	0.078550	113.0	cxp
Cardiomegaly	0.047619	0.009615	0.016000	104.0	cxp
Effusion	0.631016	0.422939	0.506438	279.0	cxp
Pneumonia	0.000000	0.000000	0.000000	31.0	cxp
Pneumothorax	0.000000	0.000000	0.000000	55.0	cxp
Consolidation	0.000000	0.000000	0.000000	46.0	cxp
Edema	0.346939	0.349315	0.348123	146.0	cxp
Any	0.521830	0.483622	0.502000	519.0	cxp
micro avg	0.406748	0.293839	0.341195	1477.0	cxp
macro avg	0.178560	0.153393	0.161234	1477.0	cxp

```

metric avg      0.178333  0.179958  0.183801  1477.0  cxp
weighted avg      0.344771  0.293839  0.313609  1477.0  cxp

samples avg      0.198069  0.179958  0.183801  1477.0  cxp
Label Accuracy Positive Accuracy (TPR) Negative Accuracy (SPC) \
0 No Finding 0.803758 0.000000 0.000000 0
.994832
1 Atelectasis 0.681628 0.115044 0.000000 0
.757396
2 Cardiomegaly 0.871608 0.009615 0.000000 0
.976581
3 Effusion 0.759916 0.422939 0.000000 0
.898380
4 Pneumonia 0.966597 0.000000 0.000000 0
.998921
5 Pneumothorax 0.936326 0.000000 0.000000 0
.993355
6 Consolidation 0.949896 0.000000 0.000000 0
.997807
7 Edema 0.800626 0.349315 0.000000 0
.881773
8 Any 0.480167 0.483622 0.000000 0
.476082

```

```

source Min Metric Value Min Metric Name
0 cxp 0.000000 Positive Accuracy (TPR)
1 cxp 0.115044 Positive Accuracy (TPR)
2 cxp 0.009615 Positive Accuracy (TPR)
3 cxp 0.422939 Positive Accuracy (TPR)
4 cxp 0.000000 Positive Accuracy (TPR)
5 cxp 0.000000 Positive Accuracy (TPR)
6 cxp 0.000000 Positive Accuracy (TPR)
7 cxp 0.349315 Positive Accuracy (TPR)
8 cxp 0.476082 Negative Accuracy (SPC)

```

Worst Performing Label: No Finding with a minimum metric value of 0.00% (Metric: Positive Accuracy (TPR))

In [36]: `combined_accuracy, combined_report = combine_with_existing(combined_ac`

```

precision recall f1-score support source
0 0.661376 0.326371 0.437063 383.0 mimic
1 0.269231 0.080000 0.123348 175.0 mimic
2 0.104377 0.203947 0.138085 152.0 mimic
3 0.466667 0.242775 0.319392 173.0 mimic
4 0.076923 0.016393 0.027027 61.0 mimic
5 0.000000 0.000000 0.000000 29.0 mimic
6 0.000000 0.000000 0.000000 39.0 mimic
7 0.161290 0.067568 0.095238 74.0 mimic
8 0.392111 0.380631 0.386286 444.0 mimic
9 0.350861 0.252941 0.293961 1530.0 mimic

```

10	0.236886	0.146409	0.169604	1530.0	mimic
11	0.384148	0.252941	0.291132	1530.0	mimic
12	0.224625	0.213468	0.214888	1530.0	mimic
13	0.622807	0.223975	0.329466	317.0	pad
14	0.107143	0.120000	0.113208	50.0	pad
15	0.052632	0.102564	0.069565	39.0	pad
16	0.250000	0.083333	0.125000	72.0	pad
17	0.000000	0.000000	0.000000	21.0	pad
18	0.000000	0.000000	0.000000	6.0	pad
19	0.000000	0.000000	0.000000	26.0	pad
20	0.000000	0.000000	0.000000	12.0	pad
21	0.363014	0.291209	0.323171	182.0	pad
22	0.303030	0.193103	0.235889	725.0	pad
23	0.155066	0.091231	0.106712	725.0	pad
24	0.398494	0.193103	0.249147	725.0	pad
25	0.123425	0.146825	0.127541	725.0	pad
26	0.000000	0.000000	0.000000	184.0	cxp
27	0.059633	0.115044	0.078550	113.0	cxp
28	0.047619	0.009615	0.016000	104.0	cxp
29	0.631016	0.422939	0.506438	279.0	cxp
30	0.000000	0.000000	0.000000	31.0	cxp
31	0.000000	0.000000	0.000000	55.0	cxp
32	0.000000	0.000000	0.000000	46.0	cxp
33	0.346939	0.349315	0.348123	146.0	cxp
34	0.521830	0.483622	0.502000	519.0	cxp
35	0.406748	0.293839	0.341195	1477.0	cxp
36	0.178560	0.153393	0.161234	1477.0	cxp
37	0.344771	0.293839	0.313609	1477.0	cxp
38	0.198069	0.179958	0.183801	1477.0	cxp

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy
cy (SPC) \				
0	No Finding	0.698502		0.326371
0.906569				
1	Atelectasis	0.813670		0.080000
0.957447				
2	Cardiomegaly	0.637640		0.203947
0.709607				
3	Effusion	0.832397		0.242775
0.946369				
4	Pneumonia	0.932584		0.016393
0.988083				
5	Pneumothorax	0.972846		0.000000
1.000000				
6	Consolidation	0.963483		0.000000
1.000000				
7	Edema	0.911049		0.067568
0.973843				
8	Any	0.497191		0.380631
0.580128				
9	No Finding	0.582973		0.223975
0.000000				

```

0.885638
10 Atelectasis 0.864358 0.120000
0.922240
11 Cardiomegaly 0.845599 0.102564
0.889908
12 Effusion 0.878788 0.083333
0.971014
13 Pneumonia 0.963925 0.000000
0.994048
14 Pneumothorax 0.936508 0.000000
0.944687
15 Consolidation 0.958153 0.000000
0.995502
16 Edema 0.981241 0.000000
0.998532
17 Any 0.679654 0.291209
0.818004
18 No Finding 0.803758 0.000000
0.994832
19 Atelectasis 0.681628 0.115044
0.757396
20 Cardiomegaly 0.871608 0.009615
0.976581
21 Effusion 0.759916 0.422939
0.898380
22 Pneumonia 0.966597 0.000000
0.998921
23 Pneumothorax 0.936326 0.000000
0.993355
24 Consolidation 0.949896 0.000000
0.997807
25 Edema 0.800626 0.349315
0.881773
26 Any 0.480167 0.483622
0.476082

```

	source	Min Metric Value	Min Metric Name
0	mimic	0.326371	Positive Accuracy (TPR)
1	mimic	0.080000	Positive Accuracy (TPR)
2	mimic	0.203947	Positive Accuracy (TPR)
3	mimic	0.242775	Positive Accuracy (TPR)
4	mimic	0.016393	Positive Accuracy (TPR)
5	mimic	0.000000	Positive Accuracy (TPR)
6	mimic	0.000000	Positive Accuracy (TPR)
7	mimic	0.067568	Positive Accuracy (TPR)
8	mimic	0.380631	Positive Accuracy (TPR)
9	pad	0.223975	Positive Accuracy (TPR)
10	pad	0.120000	Positive Accuracy (TPR)
11	pad	0.102564	Positive Accuracy (TPR)
12	pad	0.083333	Positive Accuracy (TPR)

## NIH only

```
Manually setting output dim to 1024
1024
```

Page 53 of 103

```
In [38]: nih_predictions, nih_probabilities, nih_labels = predict(model_NIH, da

print('predictions', nih_predictions)
print('labels', nih_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 8
Batch number: 1 of 8
Batch number: 2 of 8
Batch number: 3 of 8
Batch number: 4 of 8
Batch number: 5 of 8
Batch number: 6 of 8
Batch number: 7 of 8
Batch number: 8 of 8
Done.
Prediction took: 0:00:06.114199
predictions [[ True False False ... False False False]
 [ True False False ... False False False]
 [ True False False ... False False False]
 ...
 [False False False ... False False False]
 [ True False False ... False False False]
 [ True False False ... False False False]]
labels [[0. 1. 0. ... 0. 0. 1.]
 [1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 1. 1.]
 ...
 [0. 0. 1. ... 0. 1. 1.]
 [0. 0. 0. ... 0. 0. 1.]
 [1. 0. 0. ... 0. 0. 0.]]
Done
```

```
In [39]: nih_report, nih_accuracy = calculate_stats(nih_predictions, nih_probab
```

```
Overall Accuracy: 27.84%
Classification Report:
```

	precision	recall	f1-score	support	source
No Finding	0.411215	0.695652	0.516887	253.0	nih
Atelectasis	0.000000	0.000000	0.000000	131.0	nih
Cardiomegaly	0.500000	0.012048	0.023529	83.0	nih
Effusion	0.500000	0.009836	0.019293	305.0	nih
Pneumonia	0.000000	0.000000	0.000000	14.0	nih
Pneumothorax	0.000000	0.000000	0.000000	61.0	nih
Consolidation	0.000000	0.000000	0.000000	66.0	nih
Edema	0.000000	0.000000	0.000000	152.0	nih
Any	0.600000	0.016393	0.031915	549.0	nih
micro avg	0.412664	0.117100	0.182432	1614.0	nih
macro avg	0.223468	0.081548	0.065736	1614.0	nih

metrics avg	0.225100	0.183250	0.180737	1614.0	nih
weighted avg	0.388747	0.117100	0.096735	1614.0	nih
samples avg	0.180737	0.183250	0.181441	1614.0	nih
Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy		
y (SPC) \					
0 No Finding	0.669347	0.695652		0	
.660377					
1 Atelectasis	0.866332	0.000000		0	
.997685					
2 Cardiomegaly	0.916583	0.012048		0	
.998904					
3 Effusion	0.693467	0.009836		0	
.995652					
4 Pneumonia	0.985930	0.000000		1	
.000000					
5 Pneumothorax	0.938693	0.000000		1	
.000000					
6 Consolidation	0.931658	0.000000		0	
.997847					
7 Edema	0.844221	0.000000		0	
.996441					
8 Any	0.451256	0.016393		0	
.986547					

source	Min Metric Value	Min Metric Name
0 nih	0.660377	Negative Accuracy (SPC)
1 nih	0.000000	Positive Accuracy (TPR)
2 nih	0.012048	Positive Accuracy (TPR)
3 nih	0.009836	Positive Accuracy (TPR)
4 nih	0.000000	Positive Accuracy (TPR)
5 nih	0.000000	Positive Accuracy (TPR)
6 nih	0.000000	Positive Accuracy (TPR)
7 nih	0.000000	Positive Accuracy (TPR)
8 nih	0.016393	Positive Accuracy (TPR)

Worst Performing Label: Atelectasis with a minimum metric value of 0.00% (Metric: Positive Accuracy (TPR))

In [40]: combined\_accuracy, combined\_report = combine\_with\_existing(combined\_ac

	precision	recall	f1-score	support	source
0	0.661376	0.326371	0.437063	383.0	mimic
1	0.269231	0.080000	0.123348	175.0	mimic
2	0.104377	0.203947	0.138085	152.0	mimic
3	0.466667	0.242775	0.319392	173.0	mimic
4	0.076923	0.016393	0.027027	61.0	mimic
5	0.000000	0.000000	0.000000	29.0	mimic
6	0.000000	0.000000	0.000000	39.0	mimic
7	0.161290	0.067568	0.095238	74.0	mimic
8	0.392111	0.380631	0.386286	444.0	mimic
9	0.350861	0.252941	0.293961	1530.0	mimic
10	0.236886	0.146409	0.169604	1530.0	mimic
11	0.384148	0.252941	0.291132	1530.0	mimic
12	0.224625	0.213468	0.214888	1530.0	mimic
13	0.622807	0.223975	0.329466	317.0	pad
14	0.107143	0.120000	0.113208	50.0	pad
15	0.052632	0.102564	0.069565	39.0	pad
16	0.250000	0.083333	0.125000	72.0	pad
17	0.000000	0.000000	0.000000	21.0	pad
18	0.000000	0.000000	0.000000	6.0	pad

## CXP and NIH



```

In [41]: print("Loading a model trained on both CXP and NIH")

model_CXP_NIH = EmbModel(emb_type="densenet121", feature_size_override
model_CXP_NIH.load_state_dict(torch.load("model/balanced/model-snapsho
model_CXP_NIH.eval()

(conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
(norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(rel0): ReLU(inplace=True)
(pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
(denseblock1): _DenseBlock(
  (denselayer1): _DenseLayer(
    (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu1): ReLU(inplace=True)
    (conv1): Conv2d(64, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (norm2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu2): ReLU(inplace=True)
    (conv2): Conv2d(128, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  )
)

```

```
In [42]: exp_nih_predictions, exp_nih_probabilities, exp_nih_labels = predict(m)

print('predictions', exp_nih_predictions)
print('labels', exp_nih_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 8
Batch number: 1 of 8
Batch number: 2 of 8
Batch number: 3 of 8
Batch number: 4 of 8
Batch number: 5 of 8
Batch number: 6 of 8
Batch number: 7 of 8
Batch number: 8 of 8
Done.
Prediction took: 0:00:02.665518
predictions [[False False False ... False True True]
 [False False False ... False False False]
 [False False False ... False False False]
 ...
 [False False False ... False False False]
 [False False False ... False False False]
 [False True False ... False True True]]
labels [[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 1.]
 [1. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]]
Done
```

In [43]: `cxp_nih_report, cxp_nih_accuracy = calculate_stats(cxp_nih_predictions`

Overall Accuracy: 16.18%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.166667	0.005435	0.010526	184.0	cxp+nih
Atelectasis	0.057018	0.115044	0.076246	113.0	cxp+nih
Cardiomegaly	0.138889	0.048077	0.071429	104.0	cxp+nih
Effusion	0.331140	0.541219	0.410884	279.0	cxp+nih
Pneumonia	0.043478	0.064516	0.051948	31.0	cxp+nih
Pneumothorax	0.000000	0.000000	0.000000	55.0	cxp+nih
Consolidation	0.000000	0.000000	0.000000	46.0	cxp+nih
Edema	0.118357	0.335616	0.175000	146.0	cxp+nih
Any	0.500861	0.560694	0.529091	519.0	cxp+nih
micro avg	0.266112	0.346649	0.301088	1477.0	cxp+nih
macro avg	0.150712	0.185622	0.147236	1477.0	cxp+nih
weighted avg	0.286064	0.346649	0.294094	1477.0	cxp+nih
samples avg	0.208246	0.216058	0.203544	1477.0	cxp+nih

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy
0	No Finding	0.803758	0.005435	0.005435

In [44]: `combined_accuracy, combined_report = combine_with_existing(combined_ac`

	precision	recall	f1-score	support	source
0	0.661376	0.326371	0.437063	383.0	mimic
1	0.269231	0.080000	0.123348	175.0	mimic
2	0.104377	0.203947	0.138085	152.0	mimic
3	0.466667	0.242775	0.319392	173.0	mimic
4	0.076923	0.016393	0.027027	61.0	mimic
..	...	...	...	...	...
60	0.500861	0.560694	0.529091	519.0	cxp+nih
61	0.266112	0.346649	0.301088	1477.0	cxp+nih
62	0.150712	0.185622	0.147236	1477.0	cxp+nih
63	0.286064	0.346649	0.294094	1477.0	cxp+nih
64	0.208246	0.216058	0.203544	1477.0	cxp+nih

[65 rows x 5 columns]

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy
0	No Finding	0.698502	0.326371	0.906569
1	Atelectasis	0.813670	0.080000	0.057447

In [45]:

```
# nih_cxp_report, nih_cxp_accuracy = calculate_stats(cxp_nih_prediction)
# combined_accuracy, combined_report = combine_with_existing(combined_
```

## CXP and PAD

In [46]:

```
print("Loading a model trained on both CXP and PAD")

model_CXP_PAD = EmbModel(emb_type="densenet121", feature_size_override=1024)
model_CXP_PAD.load_state_dict(torch.load("model/balanced/model-snapshot.pth"))
model_CXP_PAD.eval()
```

Loading a model trained on both CXP and PAD

Emb Dim:

1024

Manually setting output dim to 1024

1024

Out[46]:

```
EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
```

```
In [47]: exp_pad_predictions, exp_pad_probabilities, exp_pad_labels = predict(m

print('predictions', exp_pad_predictions)
print('labels', exp_pad_labels)
print("Done")
```

Started prediction validation

Number of batches: 11

Batch number: 1 of 11

Batch number: 2 of 11

Batch number: 3 of 11

Batch number: 4 of 11

Batch number: 5 of 11

Batch number: 6 of 11

Batch number: 7 of 11

Batch number: 8 of 11

Batch number: 9 of 11

Batch number: 10 of 11

Batch number: 11 of 11

Done.

Prediction took: 0:00:31.718914

predictions [[False False False ... False True True]

[False False False ... False False False]

[False False False ... False False False]

...

[False False False ... False True True]

[False False False ... False False False]

[False False False ... False False False]]

labels [[0. 1. 0. ... 0. 0. 1.]

[1. 0. 0. ... 0. 0. 0.]

[1. 0. 0. ... 0. 0. 0.]

...

[0. 1. 0. ... 0. 1. 1.]

[0. 1. 1. ... 0. 1. 1.]

[0. 0. 0. ... 0. 1. 1.]]

Done

```
In [48]: exp_pad_report, exp_pad_accuracy = calculate_stats(exp_pad_predictions
```

Overall Accuracy: 18.54%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.600000	0.018809	0.036474	319.0	cxp+pad
Atelectasis	0.000000	0.000000	0.000000	256.0	cxp+pad
Cardiomegaly	0.000000	0.000000	0.000000	196.0	cxp+pad
Effusion	0.523466	0.357143	0.424597	406.0	cxp+pad
Pneumonia	0.000000	0.000000	0.000000	54.0	cxp+pad
Pneumothorax	0.000000	0.000000	0.000000	84.0	cxp+pad
Consolidation	0.000000	0.000000	0.000000	79.0	cxp+pad
Edema	0.314050	0.355140	0.333333	214.0	cxp+pad

Edema	0.782082	0.398274	0.527778	811.0	cxp+pad
Any					
micro avg	0.577731	0.227367	0.326313	2419.0	cxp+pad
macro avg	0.246622	0.125485	0.146909	2419.0	cxp+pad
weighted avg	0.456967	0.227367	0.282506	2419.0	cxp+pad
samples avg	0.178224	0.155401	0.162249	2419.0	cxp+pad

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy
y (SPC) \				
0	No Finding	0.768613	0.018809	0
		.996194		
1	Atelectasis	0.808759	0.000000	0
		.994614		
2	Cardiomegaly	0.855474	0.000000	0
		.998296		
3	Effusion	0.713139	0.357143	0
		.863071		
4	Pneumonia	0.960584	0.000000	1
		.000000		
5	Pneumothorax	0.937226	0.000000	0
		.998445		
6	Consolidation	0.942336	0.000000	1
		.000000		
7	Edema	0.778102	0.355140	0
		.856401		
8	Any	0.578102	0.398274	0
		.838998		

	source	Min Metric Value	Min Metric Name
0	cxp+pad	0.018809	Positive Accuracy (TPR)
1	cxp+pad	0.000000	Positive Accuracy (TPR)
2	cxp+pad	0.000000	Positive Accuracy (TPR)
3	cxp+pad	0.357143	Positive Accuracy (TPR)
4	cxp+pad	0.000000	Positive Accuracy (TPR)
5	cxp+pad	0.000000	Positive Accuracy (TPR)
6	cxp+pad	0.000000	Positive Accuracy (TPR)
7	cxp+pad	0.355140	Positive Accuracy (TPR)
8	cxp+pad	0.398274	Positive Accuracy (TPR)

Worst Performing Label: Atelectasis with a minimum metric value of 0.00% (Metric: Positive Accuracy (TPR))

In [49]: `combined_accuracy, combined_report = combine_with_existing(combined_ac`

	precision	recall	f1-score	support	source
0	0.661376	0.326371	0.437063	383.0	mimic
1	0.269231	0.080000	0.123348	175.0	mimic
2	0.104377	0.203947	0.138085	152.0	mimic
3	0.466667	0.242775	0.319392	173.0	mimic
4	0.076923	0.016393	0.027027	61.0	mimic
..	...	...	...	...	...
73	0.782082	0.398274	0.527778	811.0	cxp+pad
74	0.577731	0.227367	0.326313	2419.0	cxp+pad
75	0.246622	0.125485	0.146909	2419.0	cxp+pad
76	0.456967	0.227367	0.282506	2419.0	cxp+pad
77	0.178224	0.155401	0.162249	2419.0	cxp+pad

[78 rows x 5 columns]

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy
cy (SPC) \				
0	No Finding	0.698502	0.326371	0.906569
1	Atelectasis	0.813670	0.080000	0.057447

In [50]: `# pad_cxp_report, pad_cxp_accuracy = calculate_stats(cxp_pad_prediction)`  
`# combined_accuracy, combined_report = combine_with_existing(combined_ac`

## NIH and PAD

```
In [51]: print("Loading a model trained on both NIH and PAD")

model_NIH_PAD = EmbModel(emb_type="densenet121", feature_size_override
model_NIH_PAD.load_state_dict(torch.load("model/balanced/model-snapsho
model_NIH_PAD.eval()
```

Loading a model trained on both NIH and PAD

Emb Dim:

1024

Manually setting output dim to 1024

1024

```
Out[51]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=Tr
```



```
In [52]: nih_pad_predictions, nih_pad_probabilities, nih_pad_labels = predict(m

print('predictions', nih_pad_predictions)
print('labels', nih_pad_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 9
Batch number: 1 of 9
Batch number: 2 of 9
Batch number: 3 of 9
Batch number: 4 of 9
Batch number: 5 of 9
Batch number: 6 of 9
Batch number: 7 of 9
Batch number: 8 of 9
Batch number: 9 of 9
Done.
Prediction took: 0:00:33.704328
predictions [[False False False ... False False False]
 [False False False ... False False False]
 [False False False ... False False False]
 ...
 [ True False False ... False False False]
 [ True False False ... False False False]
 [False False False ... False False False]]
labels [[1. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 ...
 [1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 1. ... 0. 1. 1.]]
Done
```

```
In [53]: nih_pad_report, nih_pad_accuracy = calculate_stats(nih_pad_predictions
```

Overall Accuracy: 35.48%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.499234	0.721239	0.590045	452.0	nih+pad
Atelectasis	1.000000	0.005181	0.010309	193.0	nih+pad
Cardiomegaly	0.000000	0.000000	0.000000	131.0	nih+pad
Effusion	0.294118	0.025126	0.046296	199.0	nih+pad
Pneumonia	0.000000	0.000000	0.000000	44.0	nih+pad
Pneumothorax	0.000000	0.000000	0.000000	35.0	nih+pad
Consolidation	0.000000	0.000000	0.000000	59.0	nih+pad
Edema	0.000000	0.000000	0.000000	80.0	nih+pad
Any	0.608696	0.029536	0.056338	474.0	nih+pad
micro avg	0.494286	0.207558	0.292353	1667.0	nih+pad

micro avg	0.151200	0.207558	0.120200	1667.0	nih+pad
macro avg	0.266894	0.086787	0.078110	1667.0	nih+pad
weighted avg	0.459331	0.207558	0.182728	1667.0	nih+pad
samples avg	0.302112	0.301629	0.301129	1667.0	nih+pad

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy (SPC)
0	No Finding	0.590045	0.721239	0.499234
1	Atelectasis	0.826244	0.005181	0.000000
2	Cardiomegaly	0.881448	0.000000	0.000000
3	Effusion	0.813575	0.025126	0.986755
4	Pneumonia	0.960181	0.000000	0.000000
5	Pneumothorax	0.968326	0.000000	0.000000
6	Consolidation	0.944796	0.000000	0.998088
7	Edema	0.923982	0.000000	0.996098
8	Any	0.575566	0.029536	0.985737

	source	Min Metric Value	Min Metric Name
0	nih+pad	0.499234	Negative Accuracy (SPC)
1	nih+pad	0.005181	Positive Accuracy (TPR)
2	nih+pad	0.000000	Positive Accuracy (TPR)
3	nih+pad	0.025126	Positive Accuracy (TPR)
4	nih+pad	0.000000	Positive Accuracy (TPR)
5	nih+pad	0.000000	Positive Accuracy (TPR)
6	nih+pad	0.000000	Positive Accuracy (TPR)
7	nih+pad	0.000000	Positive Accuracy (TPR)
8	nih+pad	0.029536	Positive Accuracy (TPR)

Worst Performing Label: Cardiomegaly with a minimum metric value of 0.00% (Metric: Positive Accuracy (TPR))

In [54]: combined\_accuracy, combined\_report = combine\_with\_existing(combined\_ac

	precision	recall	f1-score	support	source
0	0.661376	0.326371	0.437063	383.0	mimic
1	0.269231	0.080000	0.123348	175.0	mimic
2	0.104377	0.203947	0.138085	152.0	mimic
3	0.466667	0.242775	0.319392	173.0	mimic
4	0.076923	0.016393	0.027027	61.0	mimic
..	...	...	...	...	...
86	0.608696	0.029536	0.056338	474.0	nih+pad
87	0.494286	0.207558	0.292353	1667.0	nih+pad
88	0.266894	0.086787	0.078110	1667.0	nih+pad

```

89  0.459331  0.207558  0.182728  1667.0  nih+pad
90  0.302112  0.301629  0.301129  1667.0  nih+pad

```

```
[91 rows x 5 columns]
```

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy
0	No Finding	0.698502	0.326371	
1	Atelectasis	0.813670	0.080000	
2	Cardiomegaly	0.637640	0.203947	
3	Effusion	0.832397	0.242775	
4	Pneumonia	0.932584	0.016393	
...	...	...	...	
58	Pneumonia	0.960181	0.000000	
59	Pneumothorax	0.968326	0.000000	
60	Consolidation	0.944796	0.000000	
61	Edema	0.923982	0.000000	
62	Any	0.575566	0.029536	

	source	Min Metric Value	Min Metric Name
0	mimic	0.326371	Positive Accuracy (TPR)
1	mimic	0.080000	Positive Accuracy (TPR)
2	mimic	0.203947	Positive Accuracy (TPR)
3	mimic	0.242775	Positive Accuracy (TPR)
4	mimic	0.016393	Positive Accuracy (TPR)
...	...	...	...
58	nih+pad	0.000000	Positive Accuracy (TPR)
59	nih+pad	0.000000	Positive Accuracy (TPR)
60	nih+pad	0.000000	Positive Accuracy (TPR)
61	nih+pad	0.000000	Positive Accuracy (TPR)
62	nih+pad	0.029536	Positive Accuracy (TPR)

```
[63 rows x 7 columns]
```

```
In [55]: # pad_nih_report, pad_nih_accuracy = calculate_stats(nih_pad_predictions,
# combined_accuracy, combined_report = combine_with_existing(combined_report,
```

## MIMIC and PAD

```
In [56]: print("Loading a model trained on both MIMIC and PAD")

model_MIMIC_PAD = EmbModel(emb_type="densenet121", feature_size_overri
model_MIMIC_PAD.load_state_dict(torch.load("model/balanced/model-snaps
model_MIMIC_PAD.eval()
```

```
Emb Dim:
1024
Manually setting output dim to 1024
1024
```

```
Out[56]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
          (relu1): ReLU(inplace=True)
```

```
In [57]: mimic_pad_predictions, mimic_pad_probabilities, mimic_pad_labels = pre

print('predictions', mimic_pad_predictions)
print('labels', mimic_pad_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 11
Batch number: 1 of 11
Batch number: 2 of 11
Batch number: 3 of 11
Batch number: 4 of 11
Batch number: 5 of 11
Batch number: 6 of 11
Batch number: 7 of 11
Batch number: 8 of 11
Batch number: 9 of 11
Batch number: 10 of 11
Batch number: 11 of 11
Done.
Prediction took: 0:00:30.902934
predictions [[False False False ... False False False]
 [False False False ... False False False]
 [False False False ... False False False]
 ...
 [False False False ... False False True]
 [ True False False ... False False False]
 [False False False ... False False False]]
labels [[0. 0. 0. ... 0. 0. 0.]
 [0. 1. 0. ... 0. 0. 1.]
 [1. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 1.]
 [0. 0. 0. ... 0. 0. 1.]
 [0. 0. 0. ... 0. 0. 1.]]
Done
```

```
In [58]: mimic_pad_report, mimic_pad_accuracy = calculate_stats(mimic_pad_predi
```

Overall Accuracy: 27.01%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.581081	0.539185	0.559350	319.0	mimic+pad
Atelectasis	0.137931	0.015625	0.028070	256.0	mimic+pad
Cardiomegaly	0.243243	0.045918	0.077253	196.0	mimic+pad
Effusion	0.579186	0.315271	0.408293	406.0	mimic+pad
Pneumonia	0.000000	0.000000	0.000000	54.0	mimic+pad
Pneumothorax	0.000000	0.000000	0.000000	84.0	mimic+pad
Consolidation	0.000000	0.000000	0.000000	79.0	mimic+pad
Edema	0.285714	0.028037	0.051064	214.0	mimic+pad

Any	0.823529	0.258940	0.393996	811.0	mimic+pad
micro avg	0.615832	0.218685	0.322758	2419.0	mimic+pad
macro avg	0.294521	0.133664	0.168670	2419.0	mimic+pad
weighted avg	0.509519	0.218685	0.288130	2419.0	mimic+pad
samples avg	0.245560	0.222567	0.229755	2419.0	mimic+pad

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy
y (SPC) \				
0	No Finding	0.802190	0.539185	0
1	Atelectasis	0.797810	0.015625	0
2	Cardiomegaly	0.843066	0.045918	0
3	Effusion	0.729197	0.315271	0
4	Pneumonia	0.960584	0.000000	1
5	Pneumothorax	0.938686	0.000000	1
6	Consolidation	0.942336	0.000000	1
7	Edema	0.837226	0.028037	0
8	Any	0.528467	0.258940	0

	source	Min Metric Value	Min Metric Name
0	mimic+pad	0.539185	Positive Accuracy (TPR)
1	mimic+pad	0.015625	Positive Accuracy (TPR)
2	mimic+pad	0.045918	Positive Accuracy (TPR)
3	mimic+pad	0.315271	Positive Accuracy (TPR)
4	mimic+pad	0.000000	Positive Accuracy (TPR)
5	mimic+pad	0.000000	Positive Accuracy (TPR)
6	mimic+pad	0.000000	Positive Accuracy (TPR)
7	mimic+pad	0.028037	Positive Accuracy (TPR)
8	mimic+pad	0.258940	Positive Accuracy (TPR)

Worst Performing Label: Pneumonia with a minimum metric value of 0.00 % (Metric: Positive Accuracy (TPR))

In [59]: combined\_accuracy, combined\_report = combine\_with\_existing(combined\_ac

	precision	recall	f1-score	support	source
0	0.661376	0.326371	0.437063	383.0	mimic
1	0.269231	0.080000	0.123348	175.0	mimic
2	0.104377	0.203947	0.138085	152.0	mimic
3	0.466667	0.242775	0.319392	173.0	mimic
4	0.076923	0.016393	0.027027	61.0	mimic
..	...	...	...	...	...
99	0.823529	0.258940	0.393996	811.0	mimic+pad

```

100    0.615832    0.218685    0.322758    2419.0    mimic+pad
101    0.294521    0.133664    0.168670    2419.0    mimic+pad
102    0.509519    0.218685    0.288130    2419.0    mimic+pad
103    0.245560    0.222567    0.229755    2419.0    mimic+pad

```

[104 rows x 5 columns]

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy (SPC)
0	No Finding	0.698502	0.326371	0.906569
1	Atelectasis	0.813670	0.080000	0.957447
2	Cardiomegaly	0.637640	0.203947	0.709607
3	Effusion	0.832397	0.242775	0.946369
4	Pneumonia	0.932584	0.016393	0.988083
..	...	...	...	...
67	Pneumonia	0.960584	0.000000	1.000000
68	Pneumothorax	0.938686	0.000000	1.000000
69	Consolidation	0.942336	0.000000	1.000000
70	Edema	0.837226	0.028037	0.987024
71	Any	0.528467	0.258940	0.919499

	source	Min Metric Value	Min Metric Name
0	mimic	0.326371	Positive Accuracy (TPR)
1	mimic	0.080000	Positive Accuracy (TPR)
2	mimic	0.203947	Positive Accuracy (TPR)
3	mimic	0.242775	Positive Accuracy (TPR)
4	mimic	0.016393	Positive Accuracy (TPR)
..	...	...	...
67	mimic+pad	0.000000	Positive Accuracy (TPR)
68	mimic+pad	0.000000	Positive Accuracy (TPR)
69	mimic+pad	0.000000	Positive Accuracy (TPR)
70	mimic+pad	0.028037	Positive Accuracy (TPR)
71	mimic+pad	0.258940	Positive Accuracy (TPR)

[72 rows x 7 columns]

In [60]:

```
# pad_mimic_report, pad_mimic_accuracy = calculate_stats(mimic_pad_pre
# combined_accuracy, combined_report = combine_with_existing(combined_
```

## MIMIC and NIH

In [61]:

```
print("Loading a model trained on both MIMIC and NIH")
```

```
model_MIMIC_NIH = EmbModel(emb_type="densenet121", feature_size_overri
model_MIMIC_NIH.load_state_dict(torch.load("model/balanced/model-snaps
model_MIMIC_NIH.eval()
```

Loading a model trained on both MIMIC and NIH

Emb Dim:

1024

Manually setting output dim to 1024

1024

Out[61]:

```
EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=Tr
```



```
In [62]: mimic_nih_predictions, mimic_nih_probabilities, mimic_nih_labels = pre

print('predictions', mimic_nih_predictions)
print('labels', mimic_nih_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 9
Batch number: 1 of 9
Batch number: 2 of 9
Batch number: 3 of 9
Batch number: 4 of 9
Batch number: 5 of 9
Batch number: 6 of 9
Batch number: 7 of 9
Batch number: 8 of 9
Batch number: 9 of 9
Done.
Prediction took: 0:00:30.296571
predictions [[False False False ... False False  True]
 [False False False ... False False False]
 [ True False False ... False False False]
 ...
 [ True False False ... False False False]
 [ True False False ... False  True  True]
 [False False False ... False False False]]
labels [[0. 0. 0. ... 0. 0. 1.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 1. 0. 1.]
 ...
 [1. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]
 [1. 0. 0. ... 0. 0. 0.]]
Done
```

```
In [63]: mimic_nih_report, mimic_nih_accuracy = calculate_stats(mimic_nih_predi
```

Overall Accuracy: 18.45%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.578571	0.422977	0.488688	383.0	mimic+nih
Atelectasis	0.148148	0.137143	0.142433	175.0	mimic+nih
Cardiomegaly	0.219512	0.059211	0.093264	152.0	mimic+nih
Effusion	0.280000	0.485549	0.355180	173.0	mimic+nih
Pneumonia	0.040590	0.180328	0.066265	61.0	mimic+nih
Pneumothorax	0.000000	0.000000	0.000000	29.0	mimic+nih
Consolidation	0.000000	0.000000	0.000000	39.0	mimic+nih
Edema	0.021212	0.094595	0.034653	74.0	mimic+nih
Any	0.437838	0.547297	0.486486	444.0	mimic+nih
micro avg	0.274669	0.352941	0.308924	1530.0	mimic+nih

macro avg	0.191764	0.214122	0.185219	1530.0	mimic+nih
weighted avg	0.344948	0.352941	0.333544	1530.0	mimic+nih
samples avg	0.248892	0.289341	0.250638	1530.0	mimic+nih

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy (SPC)
0	No Finding	0.682584	0.422977	0.827737
1	Atelectasis	0.729401	0.137143	0.845465
2	Cardiomegaly	0.836142	0.059211	0.965066
3	Effusion	0.714419	0.485549	0.758659
4	Pneumonia	0.709738	0.180328	0.741807
5	Pneumothorax	0.970974	0.000000	0.998075
6	Consolidation	0.940075	0.000000	0.975705
7	Edema	0.634831	0.094595	0.675050
8	Any	0.519663	0.547297	0.500000

	source	Min Metric Value	Min Metric Name
0	mimic+nih	0.422977	Positive Accuracy (TPR)
1	mimic+nih	0.137143	Positive Accuracy (TPR)
2	mimic+nih	0.059211	Positive Accuracy (TPR)
3	mimic+nih	0.485549	Positive Accuracy (TPR)
4	mimic+nih	0.180328	Positive Accuracy (TPR)
5	mimic+nih	0.000000	Positive Accuracy (TPR)
6	mimic+nih	0.000000	Positive Accuracy (TPR)
7	mimic+nih	0.094595	Positive Accuracy (TPR)
8	mimic+nih	0.500000	Negative Accuracy (SPC)

Worst Performing Label: Pneumothorax with a minimum metric value of 0.00% (Metric: Positive Accuracy (TPR))

In [64]: combined\_accuracy, combined\_report = combine\_with\_existing(combined\_ac

	precision	recall	f1-score	support	source
0	0.661376	0.326371	0.437063	383.0	mimic
1	0.269231	0.080000	0.123348	175.0	mimic
2	0.104377	0.203947	0.138085	152.0	mimic
3	0.466667	0.242775	0.319392	173.0	mimic
4	0.076923	0.016393	0.027027	61.0	mimic
...	...	...	...	...	...
112	0.437838	0.547297	0.486486	444.0	mimic+nih
113	0.274669	0.352941	0.308924	1530.0	mimic+nih
114	0.191764	0.214122	0.185219	1530.0	mimic+nih

```

115    0.344948  0.352941  0.333544  1530.0  mimic+nih
116    0.248892  0.289341  0.250638  1530.0  mimic+nih

```

```
[117 rows x 5 columns]
```

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy (SPC)
0	No Finding	0.698502	0.326371	0.906569
1	Atelectasis	0.813670	0.080000	0.957447
2	Cardiomegaly	0.637640	0.203947	0.709607
3	Effusion	0.832397	0.242775	0.946369
4	Pneumonia	0.932584	0.016393	0.988083
...	...	...	...	...
76	Pneumonia	0.709738	0.180328	0.741807
77	Pneumothorax	0.970974	0.000000	0.998075
78	Consolidation	0.940075	0.000000	0.975705
79	Edema	0.634831	0.094595	0.675050
80	Any	0.519663	0.547297	0.500000

	source	Min Metric Value	Min Metric Name
0	mimic	0.326371	Positive Accuracy (TPR)
1	mimic	0.080000	Positive Accuracy (TPR)
2	mimic	0.203947	Positive Accuracy (TPR)
3	mimic	0.242775	Positive Accuracy (TPR)
4	mimic	0.016393	Positive Accuracy (TPR)
...	...	...	...
76	mimic+nih	0.180328	Positive Accuracy (TPR)
77	mimic+nih	0.000000	Positive Accuracy (TPR)
78	mimic+nih	0.000000	Positive Accuracy (TPR)
79	mimic+nih	0.094595	Positive Accuracy (TPR)
80	mimic+nih	0.500000	Negative Accuracy (SPC)

```
[81 rows x 7 columns]
```

```
In [65]: # nih_mimic_report, nih_mimic_accuracy = calculate_stats(mimic_nih_pre
# combined_accuracy, combined_report = combine_with_existing(combined_
```

## MIMIC and CXP

```
In [66]: print("Loading a model trained on both MIMIC and CXP")

model_MIMIC_CXP = EmbModel(emb_type="densenet121", feature_size_overri
model_MIMIC_CXP.load_state_dict(torch.load("model/balanced/model-snaps
model_MIMIC_CXP.eval()
```

Loading a model trained on both MIMIC and CXP

Emb Dim:  
1024  
Manually setting output dim to 1024  
1024

```
Out[66]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=Tr
```

```
In [67]: mimic_cxp_predictions, mimic_cxp_probabilities, mimic_cxp_labels = pre

print('predictions', mimic_cxp_predictions)
print('labels', mimic_cxp_labels)
print("Done")
```

```
Started prediction validation
Number of batches: 9
Batch number: 1 of 9
Batch number: 2 of 9
Batch number: 3 of 9
Batch number: 4 of 9
Batch number: 5 of 9
Batch number: 6 of 9
Batch number: 7 of 9
Batch number: 8 of 9
Batch number: 9 of 9
Done.
Prediction took: 0:00:30.373398
predictions [[False True False ... False False True]
 [False False False ... False False False]
 [False False False ... False False False]
 ...
 [False False False ... False False False]
 [False False False ... False False True]
 [False False False ... False False True]]
labels [[1. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 1. 0. 1.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [0. 1. 0. ... 0. 0. 1.]
 [0. 0. 0. ... 0. 1. 1.]
 [0. 0. 0. ... 1. 0. 1.]]
Done
```

```
In [68]: mimic_cxp_report, mimic_cxp_accuracy = calculate_stats(mimic_cxp_predi
```

Overall Accuracy: 17.98%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.741259	0.276762	0.403042	383.0	mimic+cxp
Atelectasis	0.087542	0.148571	0.110169	175.0	mimic+cxp
Cardiomegaly	0.214286	0.019737	0.036145	152.0	mimic+cxp
Effusion	0.182022	0.468208	0.262136	173.0	mimic+cxp
Pneumonia	0.285714	0.032787	0.058824	61.0	mimic+cxp
Pneumothorax	0.000000	0.000000	0.000000	29.0	mimic+cxp
Consolidation	0.000000	0.000000	0.000000	39.0	mimic+cxp
Edema	0.176471	0.040541	0.065934	74.0	mimic+cxp
Any	0.413793	0.459459	0.435432	444.0	mimic+cxp
micro avg	0.298246	0.277778	0.287648	1530.0	mimic+cxp

macro avg	0.233454	0.160674	0.152409	1530.0	mimic+cxp
weighted avg	0.377448	0.277778	0.278619	1530.0	mimic+cxp
samples avg	0.222519	0.217868	0.216000	1530.0	mimic+cxp
	Label Accuracy	Positive Accuracy (TPR)	Negative Accuracy		
y (SPC) \					
0 No Finding	0.705993	0.276762	0		
.945985					
1 Atelectasis	0.606742	0.148571	0		
.696529					
2 Cardiomegaly	0.850187	0.019737	0		
.987991					
3 Effusion	0.573034	0.468208	0		
.593296					
4 Pneumonia	0.940075	0.032787	0		
.995035					
5 Pneumothorax	0.972846	0.000000	1		
.000000					
6 Consolidation	0.955056	0.000000	0		
.991254					
7 Edema	0.920412	0.040541	0		
.985915					
8 Any	0.504682	0.459459	0		
.536859					

	source	Min Metric Value	Min Metric Name
0	mimic+cxp	0.276762	Positive Accuracy (TPR)
1	mimic+cxp	0.148571	Positive Accuracy (TPR)
2	mimic+cxp	0.019737	Positive Accuracy (TPR)
3	mimic+cxp	0.468208	Positive Accuracy (TPR)
4	mimic+cxp	0.032787	Positive Accuracy (TPR)
5	mimic+cxp	0.000000	Positive Accuracy (TPR)
6	mimic+cxp	0.000000	Positive Accuracy (TPR)
7	mimic+cxp	0.040541	Positive Accuracy (TPR)
8	mimic+cxp	0.459459	Positive Accuracy (TPR)

Worst Performing Label: Pneumothorax with a minimum metric value of 0.00% (Metric: Positive Accuracy (TPR))

In [69]: combined\_accuracy, combined\_report = combine\_with\_existing(combined\_ac

	precision	recall	f1-score	support	source
0	0.661376	0.326371	0.437063	383.0	mimic
1	0.269231	0.080000	0.123348	175.0	mimic
2	0.104377	0.203947	0.138085	152.0	mimic
3	0.466667	0.242775	0.319392	173.0	mimic
4	0.076923	0.016393	0.027027	61.0	mimic
..	...	...	...	...	...
125	0.413793	0.459459	0.435432	444.0	mimic+cxp
126	0.298246	0.277778	0.287648	1530.0	mimic+cxp
127	0.233454	0.160674	0.152409	1530.0	mimic+cxp

```

128    0.377448    0.277778    0.278619    1530.0    mimic+cxp
129    0.222519    0.217868    0.216000    1530.0    mimic+cxp

```

```
[130 rows x 5 columns]
```

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy (SPC)
0	No Finding	0.698502	0.326371	0.906569
1	Atelectasis	0.813670	0.080000	0.957447
2	Cardiomegaly	0.637640	0.203947	0.709607
3	Effusion	0.832397	0.242775	0.946369
4	Pneumonia	0.932584	0.016393	0.988083
..	...	...	...	...
85	Pneumonia	0.940075	0.032787	0.995035
86	Pneumothorax	0.972846	0.000000	1.000000
87	Consolidation	0.955056	0.000000	0.991254
88	Edema	0.920412	0.040541	0.985915
89	Any	0.504682	0.459459	0.536859

	source	Min Metric Value	Min Metric Name
0	mimic	0.326371	Positive Accuracy (TPR)
1	mimic	0.080000	Positive Accuracy (TPR)
2	mimic	0.203947	Positive Accuracy (TPR)
3	mimic	0.242775	Positive Accuracy (TPR)
4	mimic	0.016393	Positive Accuracy (TPR)
..	...	...	...
85	mimic+cxp	0.032787	Positive Accuracy (TPR)
86	mimic+cxp	0.000000	Positive Accuracy (TPR)
87	mimic+cxp	0.000000	Positive Accuracy (TPR)
88	mimic+cxp	0.040541	Positive Accuracy (TPR)
89	mimic+cxp	0.459459	Positive Accuracy (TPR)

```
[90 rows x 7 columns]
```

```
In [70]: # cxp_mimic_report, cxp_mimic_accuracy = calculate_stats(mimic_cxp_pre
# combined_accuracy, combined_report = combine_with_existing(combined_
```

## All four datasets together: MIMIC, CXP, NIH and PAD

```
In [71]: print("Loading a model trained on all four datasets: MIMIC, CXP, NIH and PAD")

model_CXP_MIMIC_NIH_PAD = EmbModel(emb_type="densenet121", feature_size=1024)
model_CXP_MIMIC_NIH_PAD.load_state_dict(torch.load("model/balanced/model_CXP_MIMIC_NIH_PAD.pth"))
model_CXP_MIMIC_NIH_PAD.eval()
```

Loading a model trained on all four datasets: MIMIC, CXP, NIH and PAD

Emb Dim:

1024

Manually setting output dim to 1024

1024

```
Out[71]: EmbModel(
  (encoder): Sequential(
    (0): Sequential(
      (conv0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
      (norm0): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu0): ReLU(inplace=True)
      (pool0): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
      (denseblock1): _DenseBlock(
        (denselayer1): _DenseLayer(
          (norm1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True)
```



In [72]: `cxp_mimic_nih_pad_predictions, cxp_mimic_nih_pad_probabilities, cxp_mi`

```
print('predictions', cxp_mimic_nih_pad_predictions)
print('labels', cxp_mimic_nih_pad_labels)
print("Done")
```

Started prediction validation

Number of batches: 17

Batch number: 1 of 17

Batch number: 2 of 17

Batch number: 3 of 17

Batch number: 4 of 17

Batch number: 5 of 17

Batch number: 6 of 17

Batch number: 7 of 17

Batch number: 8 of 17

Batch number: 9 of 17

Batch number: 10 of 17

Batch number: 11 of 17

Batch number: 12 of 17

Batch number: 13 of 17

Batch number: 14 of 17

Batch number: 15 of 17

Batch number: 16 of 17

Batch number: 17 of 17

Done.

Prediction took: 0:00:37.584432

predictions [[False False False ... False False False]

[ True False False ... False False False]

[False False False ... False False False]

...

[False False False ... False False False]

[False False False ... False False False]

[False False False ... False False False]]

labels [[0. 0. 0. ... 1. 0. 1.]

[1. 0. 0. ... 0. 0. 0.]

[1. 0. 0. ... 0. 0. 0.]

...

[0. 0. 0. ... 0. 1. 1.]

[0. 0. 0. ... 0. 0. 1.]

[0. 0. 1. ... 0. 0. 1.]]

Done

In [73]: `cxp_mimic_nih_pad_report, cxp_mimic_nih_pad_accuracy = calculate_stats`

Overall Accuracy: 33.11%

Classification Report:

	precision	recall	f1-score	support	source
No Finding	0.639155	0.523585	0.575627	636	cxp+mimic+nih+

No Finding pad	0.297297	0.035948	0.064140	306.0	cxp+mimic+nih+
Atelectasis pad	0.285714	0.008511	0.016529	235.0	cxp+mimic+nih+
Cardiomegaly pad	0.577558	0.366109	0.448143	478.0	cxp+mimic+nih+
Effusion pad	0.142857	0.040000	0.062500	75.0	cxp+mimic+nih+
Pneumonia pad	0.000000	0.000000	0.000000	90.0	cxp+mimic+nih+
Pneumothorax pad	0.000000	0.000000	0.000000	105.0	cxp+mimic+nih+
Consolidation pad	0.398551	0.243363	0.302198	226.0	cxp+mimic+nih+
Edema pad	0.798969	0.312185	0.448950	993.0	cxp+mimic+nih+
Any pad	0.628269	0.282761	0.389998	3144.0	cxp+mimic+nih+
micro avg pad	0.348900	0.169967	0.213121	3144.0	cxp+mimic+nih+
macro avg pad	0.551798	0.282761	0.357065	3144.0	cxp+mimic+nih+
weighted avg pad	0.276903	0.262579	0.266328	3144.0	cxp+mimic+nih+
samples avg pad					

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy
0	No Finding	0.761997	0.523585	0
1	Atelectasis	0.844401	0.035948	0
2	Cardiomegaly	0.884634	0.008511	0
3	Effusion	0.791081	0.366109	0
4	Pneumonia	0.956374	0.040000	0
5	Pneumothorax	0.956374	0.000000	1
6	Consolidation	0.949103	0.000000	1
7	Edema	0.876878	0.243363	0
8	Any	0.631120	0.312185	0

	source	Min Metric Value	Min Metric Name
0	cxp+mimic+nih+pad	0.523585	Positive Accuracy (TPR)
1	cxp+mimic+nih+pad	0.035948	Positive Accuracy (TPR)

```

2   cxp+mimic+nih+pad      0.008511  Positive Accuracy (TPR)
3   cxp+mimic+nih+pad      0.366109  Positive Accuracy (TPR)
4   cxp+mimic+nih+pad      0.040000  Positive Accuracy (TPR)
5   cxp+mimic+nih+pad      0.000000  Positive Accuracy (TPR)
6   cxp+mimic+nih+pad      0.000000  Positive Accuracy (TPR)
7   cxp+mimic+nih+pad      0.243363  Positive Accuracy (TPR)
8   cxp+mimic+nih+pad      0.312185  Positive Accuracy (TPR)
Worst Performing Label: Pneumothorax with a minimum metric value of 0
.00% (Metric: Positive Accuracy (TPR))

```

In [74]: combined\_accuracy, combined\_report = combine\_with\_existing(combined\_ac

```

      precision    recall  f1-score   support           source
0      0.661376    0.326371    0.437063     383.0         mimic
1      0.269231    0.080000    0.123348     175.0         mimic
2      0.104377    0.203947    0.138085     152.0         mimic
3      0.466667    0.242775    0.319392     173.0         mimic
4      0.076923    0.016393    0.027027      61.0         mimic
..      ...      ...      ...      ...      ...
138    0.798969    0.312185    0.448950     993.0    cxp+mimic+nih+pad
139    0.628269    0.282761    0.389998    3144.0    cxp+mimic+nih+pad
140    0.348900    0.169967    0.213121    3144.0    cxp+mimic+nih+pad
141    0.551798    0.282761    0.357065    3144.0    cxp+mimic+nih+pad
142    0.276903    0.262579    0.266328    3144.0    cxp+mimic+nih+pad

```

[143 rows x 5 columns]

```

      Label  Accuracy  Positive Accuracy (TPR)  Negative Accura
cy (SPC) \
0      No Finding    0.698502                0.326371
0.906569
1      Atelectasis    0.813670                0.080000
0.957447
2      Cardiomegaly    0.637640                0.203947
0.709607
3      Effusion    0.832397                0.242775
0.946369
4      Pneumonia    0.932584                0.016393
0.988083
..      ...      ...      ...
...
94      Pneumonia    0.956374                0.040000
0.990946
95      Pneumothorax    0.956374                0.000000
1.000000
96      Consolidation    0.949103                0.000000
1.000000
97      Edema    0.876878                0.243363
0.954818
98      Any    0.631120                0.312185
0.927103

```

	source	Min Metric Value	Min Metric Name
0	mimic	0.326371	Positive Accuracy (TPR)
1	mimic	0.080000	Positive Accuracy (TPR)
2	mimic	0.203947	Positive Accuracy (TPR)
3	mimic	0.242775	Positive Accuracy (TPR)
4	mimic	0.016393	Positive Accuracy (TPR)
...	...	...	...
94	cxp+mimic+nih+pad	0.040000	Positive Accuracy (TPR)
95	cxp+mimic+nih+pad	0.000000	Positive Accuracy (TPR)
96	cxp+mimic+nih+pad	0.000000	Positive Accuracy (TPR)
97	cxp+mimic+nih+pad	0.243363	Positive Accuracy (TPR)
98	cxp+mimic+nih+pad	0.312185	Positive Accuracy (TPR)

[99 rows x 7 columns]

In [ ]:

## Saving all data to data frame and to file

```
In [75]: print("Saving the evaluation stats to CSV")
combined_df_now = datetime.now()
combined_df_dt_string = combined_df_now.strftime("%d-%m-%Y-%H-%M-%S")
note = "validate_on_two"

breakdown_file_name = f"stats/{combined_df_dt_string}_combined_accuracy_breakdown_validate_on_two.csv"
totals_file_name = f"stats/{combined_df_dt_string}_combined_accuracy_totals_validate_on_two.csv"

combined_accuracy.to_csv(breakdown_file_name, index=False)
combined_report.to_csv(totals_file_name, index=False)
print("Done.")

print("File names:")
print(breakdown_file_name)
print(totals_file_name)
```

Saving the evaluation stats to CSV

Done.

File names:

stats/06-05-2024-23-18-10\_combined\_accuracy\_breakdown\_validate\_on\_two.csv  
stats/06-05-2024-23-18-10\_combined\_accuracy\_totals\_validate\_on\_two.csv

## Visualizing data

```

In [78]: timestamp_to_load = "04-05-2024-23-34-18"
two_dataset_eval_timestamp = "04-05-2024-23-34-18"
one_dataset_eval_timestamp = "05-05-2024-16-23-37"
balanced_dataset_timestamp = "05-05-2024-22-03-40"
balanced_single_dataset_timestamp = "05-05-2024-23-11-28"

balanced_single_dataset_min_max_timestamp = "06-05-2024-01-43-10" #"06-05-2024-01-43-10"

balanced_single_dataset_validate_on_two_timestamp = "06-05-2024-23-18-10"

two_dataset_accuracy_breakdown_path = f"stats/{two_dataset_eval_timestamp}_accuracy_breakdown.csv"
two_dataset_accuracy_totals_path = f"stats/{two_dataset_eval_timestamp}_accuracy_totals.csv"

one_dataset_accuracy_breakdown_path = f"stats/{one_dataset_eval_timestamp}_accuracy_breakdown.csv"
one_dataset_accuracy_totals_path = f"stats/{one_dataset_eval_timestamp}_accuracy_totals.csv"

balanced_dataset_accuracy_breakdown_path = f"stats/{balanced_dataset_timestamp}_accuracy_breakdown.csv"
balanced_dataset_accuracy_totals_path = f"stats/{balanced_dataset_timestamp}_accuracy_totals.csv"

balanced_single_dataset_accuracy_breakdown_path = f"stats/{balanced_single_dataset_timestamp}_accuracy_breakdown.csv"
balanced_single_dataset_accuracy_totals_path = f"stats/{balanced_single_dataset_timestamp}_accuracy_totals.csv"

balanced_validate_on_two_dataset_accuracy_breakdown_path = f"stats/{balanced_validate_on_two_dataset_timestamp}_accuracy_breakdown.csv"
balanced_validate_on_two_dataset_accuracy_totals_path = f"stats/{balanced_validate_on_two_dataset_timestamp}_accuracy_totals.csv"

#loading the dataframes
two_dataset_combined_accuracy = pd.read_csv(two_dataset_accuracy_breakdown_path)
two_dataset_combined_report = pd.read_csv(two_dataset_accuracy_totals_path)

one_dataset_combined_accuracy = pd.read_csv(one_dataset_accuracy_breakdown_path)
one_dataset_combined_report = pd.read_csv(one_dataset_accuracy_totals_path)

balanced_dataset_combined_accuracy = pd.read_csv(balanced_dataset_accuracy_breakdown_path)
balanced_dataset_combined_report = pd.read_csv(balanced_dataset_accuracy_totals_path)

balanced_single_dataset_combined_accuracy = pd.read_csv(balanced_single_dataset_accuracy_breakdown_path)
balanced_single_dataset_combined_report = pd.read_csv(balanced_single_dataset_accuracy_totals_path)

balanced_validate_on_two_dataset_combined_accuracy = pd.read_csv(balanced_validate_on_two_dataset_accuracy_breakdown_path)
balanced_validate_on_two_dataset_combined_report = pd.read_csv(balanced_validate_on_two_dataset_accuracy_totals_path)

print(two_dataset_combined_accuracy.head())
print(two_dataset_combined_report.head())

print(one_dataset_combined_accuracy.head())
print(one_dataset_combined_report.head())

print(balanced dataset combined accuracy.head())

```

```

print(balanced_dataset_combined_report.head())

print(balanced_single_dataset_combined_accuracy.head())
print(balanced_single_dataset_combined_report.head())

print(balanced_validate_on_two_dataset_combined_accuracy.head())
print(balanced_validate_on_two_dataset_combined_report.head())

```

	precision	recall	f1-score	support	source
0	0.654206	0.540541	0.591966	259.0	mimic
1	0.142857	0.006289	0.012048	159.0	mimic
2	0.200000	0.024590	0.043796	122.0	mimic
3	0.432692	0.300000	0.354331	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic

	Label	Accuracy	source
0	No Finding	73.92%	mimic
1	Atelectasis	77.84%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	77.84%	mimic
4	Pneumonia	93.92%	mimic

	precision	recall	f1-score	support	source
0	0.654206	0.540541	0.591966	259.0	mimic
1	0.142857	0.006289	0.012048	159.0	mimic
2	0.200000	0.024590	0.043796	122.0	mimic
3	0.432692	0.300000	0.354331	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic

	Label	Accuracy	source
0	No Finding	73.92%	mimic
1	Atelectasis	77.84%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	77.84%	mimic
4	Pneumonia	93.92%	mimic

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic

	precision	recall	f1-score	support	source
0	0.715190	0.436293	0.541966	259.0	mimic
1	0.269231	0.088050	0.132701	159.0	mimic
2	0.153846	0.016393	0.029630	122.0	mimic
3	0.471910	0.280000	0.351464	150.0	mimic
4	0.000000	0.000000	0.000000	42.0	mimic

	Label	Accuracy	source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic

	Label	Accuracy	Source
0	No Finding	74.19%	mimic
1	Atelectasis	75.27%	mimic
2	Cardiomegaly	82.30%	mimic
3	Effusion	79.05%	mimic
4	Pneumonia	93.65%	mimic

	precision	recall	f1-score	support	source
0	0.661376	0.326371	0.437063	383.0	mimic
1	0.269231	0.080000	0.123348	175.0	mimic
2	0.104377	0.203947	0.138085	152.0	mimic
3	0.466667	0.242775	0.319392	173.0	mimic
4	0.076923	0.016393	0.027027	61.0	mimic

	Label	Accuracy	Positive Accuracy (TPR)	Negative Accuracy (SPC) \
0	No Finding	0.698502	0.326371	0.906569
1	Atelectasis	0.813670	0.080000	0.957447
2	Cardiomegaly	0.637640	0.203947	0.709607
3	Effusion	0.832397	0.242775	0.946369
4	Pneumonia	0.932584	0.016393	0.988083

	source	Min Metric Value	Min Metric Name
0	mimic	0.326371	Positive Accuracy (TPR)
1	mimic	0.080000	Positive Accuracy (TPR)
2	mimic	0.203947	Positive Accuracy (TPR)
3	mimic	0.242775	Positive Accuracy (TPR)
4	mimic	0.016393	Positive Accuracy (TPR)

```

In [79]: def get_accuracy_df(report):
    metric_column_names = ['Accuracy', 'Positive Accuracy (TPR)', 'Neg

    baseline_metrics = {}
    for source in report['source'].unique():
        if '+' not in source: #only single sources
            source_data = report[report['source'] == source]
            for label in source_data['Label'].unique():
                for metric in metric_column_names:
                    key = (source, label, metric)
                    filtered_data = source_data[source_data['Label'] == label]
                    if not filtered_data.empty:
                        baseline_value = filtered_data.iloc[0][metric]
                        baseline_metrics[key] = baseline_value
                    else:
                        print(f"No data found for {key}")

    accuracy_changes = []
    for source in report['source'].unique():
        if '+' in source:
            parts = source.split('+')
            source_data = report[report['source'] == source]
            for label in source_data['Label'].unique():
                for metric in metric_column_names:
                    current_data = source_data[source_data['Label'] == label]
                    current_value = current_data.iloc[0][metric]
                    for part in parts:
                        base_key = (part, label, metric)
                        base_value = baseline_metrics.get(base_key, 0)
                        change = current_value - base_value
                        names_without_base = [x for x in parts if x != part]
                        name_base_first = [part] + names_without_base
                        if len(parts) > 2 and not(source.startswith(part)):
                            pass
                        else:
                            accuracy_changes.append({
                                'source_combination': "+".join(name_base_first),
                                'part_source': part,
                                'label': label,
                                'metric': metric,
                                'change': change
                            })

    accuracy_changes_df = pd.DataFrame(accuracy_changes)
    print(accuracy_changes_df)

    return accuracy_changes_df

```



```
In [81]: #commenting these out since they don't have the positive and negative
# two_datasets_accuracy_changes_df = get_accuracy_df(two_dataset_combi
# one_datasets_accuracy_changes_df = get_accuracy_df(one_dataset_combi
# balanced_dataset_accuracy_changes_df = get_accuracy_df(balanced_data
# balanced_single_dataset_accuracy_changes_df = get_accuracy_df(balanc

# balanced_min_max_dataset_accuracy_changes_df = get_accuracy_df(balan

balanced_validate_on_two_dataset_accuracy_changes_df = get_accuracy_df
```

	source_combination	part_source	label	metr
ic \				
0	cxp+nih	cxp	No Finding	Accura
cy				
1	nih+cxp	nih	No Finding	Accura
cy				
2	cxp+nih	cxp	No Finding	Positive Accuracy (TP
R)				
3	nih+cxp	nih	No Finding	Positive Accuracy (TP
R)				
4	cxp+nih	cxp	No Finding	Negative Accuracy (SP
C)				
..	...	...	...	.
..				
346	cxp+mimic+nih+pad	cxp	Edema	Positive Accuracy (TP
R)				
347	cxp+mimic+nih+pad	cxp	Edema	Negative Accuracy (SP
C)				
348	cxp+mimic+nih+pad	cxp	Any	Accura
cy				
349	cxp+mimic+nih+pad	cxp	Any	Positive Accuracy (TP
R)				
350	cxp+mimic+nih+pad	cxp	Any	Negative Accuracy (SP
C)				

	change
0	0.000000
1	0.134411
2	0.005435
3	-0.690217
4	-0.001292
..	...
346	-0.105952
347	0.073044
348	0.150953
349	-0.171437
350	0.451021

[351 rows x 5 columns]

In [ ]:

```
In [82]: def plot_breakdown_accuracy_changes(accuracy_changes_df, metric='Accuracy')
plt.figure(figsize=(14, 8))
label_order = ['Pneumonia', 'Cardiomegaly', 'Edema', 'Effusion', 'Pleural Effusion']
source_order = [
    'mimic+cxp', 'mimic+nih', 'mimic+pad',
    'cxp+mimic', 'cxp+nih', 'cxp+pad',
    'nih+mimic', 'nih+cxp', 'nih+pad',
    'pad+mimic', 'pad+cxp', 'pad+nih'
]

filtered_df = accuracy_changes_df[accuracy_changes_df['metric'] == metric]

ax = sns.barplot(data=filtered_df, y='change', x='label', hue='source')

for y in ax.get_yticks():
    plt.axhline(y, color='gray', linewidth=0.5, linestyle='--')

plt.axhline(0, color='gray', linewidth=0.8)

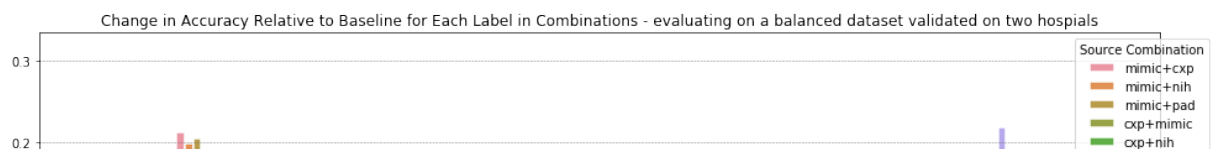
plt.title(f'Change in {metric} Relative to Baseline for Each Label')
plt.ylabel(f'Change in {metric}')
plt.xlabel('Disease Label')
plt.xticks()
plt.legend(title='Source Combination', loc='upper right', bbox_to_anchor=(1.05, 0.9))
plt.tight_layout()
plt.show()
```

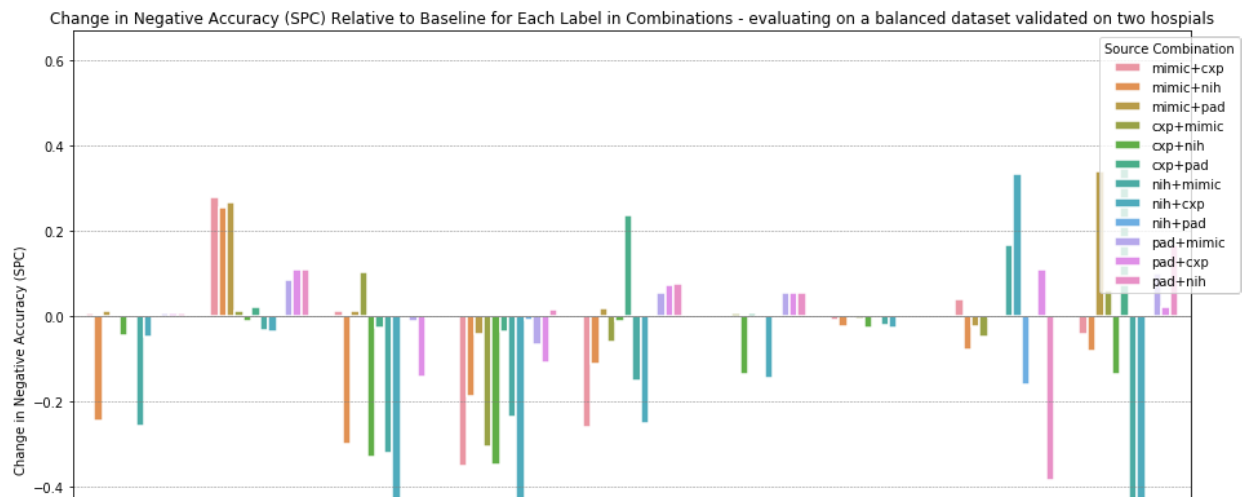


```
In [84]: #commenting these out since they don't have the positive and negative
# plot_breakdown_accuracy_changes(two_datasets_accuracy_changes_df, "Accuracy")
# plot_breakdown_accuracy_changes(one_datasets_accuracy_changes_df, "Accuracy")
# plot_breakdown_accuracy_changes(balanced_dataset_accuracy_changes_df, "Accuracy")
# plot_breakdown_accuracy_changes(balanced_single_dataset_accuracy_changes_df, "Accuracy")


# plot_breakdown_accuracy_changes(balanced_min_max_dataset_accuracy_changes_df, "Accuracy")

# plot_breakdown_accuracy_changes(balanced_min_max_dataset_accuracy_changes_df, "Accuracy")
# plot_breakdown_accuracy_changes(balanced_min_max_dataset_accuracy_changes_df, "Accuracy")
# plot_breakdown_accuracy_changes(balanced_min_max_dataset_accuracy_changes_df, "Accuracy")

plot_breakdown_accuracy_changes(balanced_validate_on_two_dataset_accuracy_changes_df, "Accuracy")
plot_breakdown_accuracy_changes(balanced_validate_on_two_dataset_accuracy_changes_df, "Accuracy")
plot_breakdown_accuracy_changes(balanced_validate_on_two_dataset_accuracy_changes_df, "Accuracy")
```







```
In [85]: def plot_accuracy_per_dataset(report, chart_name=""):
    dataset_names = ['nih', 'mimic', 'pad', 'cxp']

    for dataset_name in dataset_names:
        dataset_specific = report[report['source'].str.contains(dataset_name)]

        sorted_sources = sorted(dataset_specific['source'].unique(), key=lambda x: x)
        dataset_specific = dataset_specific[dataset_specific['source'].isin(sorted_sources)]

        plt.figure(figsize=(12, 6))

        ax = sns.barplot(data=dataset_specific, x='Label', y='Accuracy')

        for y in ax.get_yticks():
            plt.axhline(y, color='gray', linewidth=0.5, linestyle='--')

        plt.title(f'Comparison of Model Accuracies by Disease Label for {dataset_name}')
        plt.ylabel('Accuracy')
        plt.xlabel('Disease Label')
        plt.xticks()
        plt.legend(title='Source', loc='lower right')


        plt.tight_layout()


        # Show the plot
        plt.show()
```

```
In [86]: #commenting these out since they don't have the positive and negative
# plot_accuracy_per_dataset(two_dataset_combined_report, "evaluating on a dataset from two hospitals")
# plot_accuracy_per_dataset(one_dataset_combined_report, "evaluating on a dataset from one hospital")
# plot_accuracy_per_dataset(balanced_dataset_combined_report, "evaluating on a balanced dataset")
# plot_accuracy_per_dataset(balanced_single_dataset_combined_report, "evaluating on a balanced single dataset")

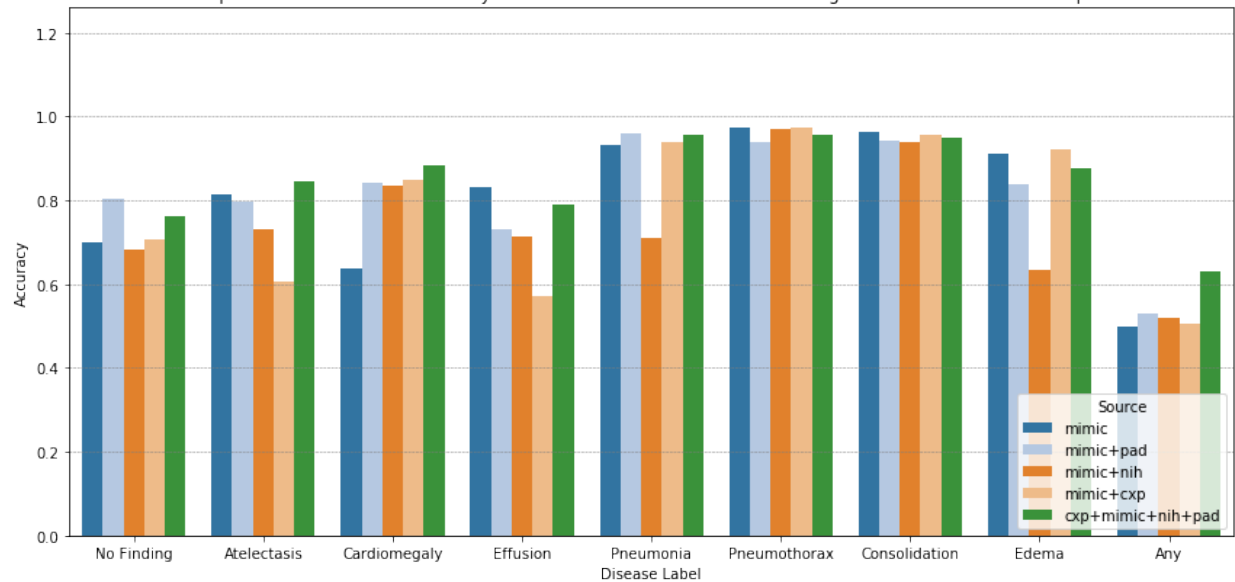
# plot_accuracy_per_dataset(balanced_min_max_dataset_combined_report, "evaluating on a balanced min_max dataset")

plot_accuracy_per_dataset(balanced_validate_on_two_dataset_combined_report, "evaluating on a balanced validate on two dataset")
```

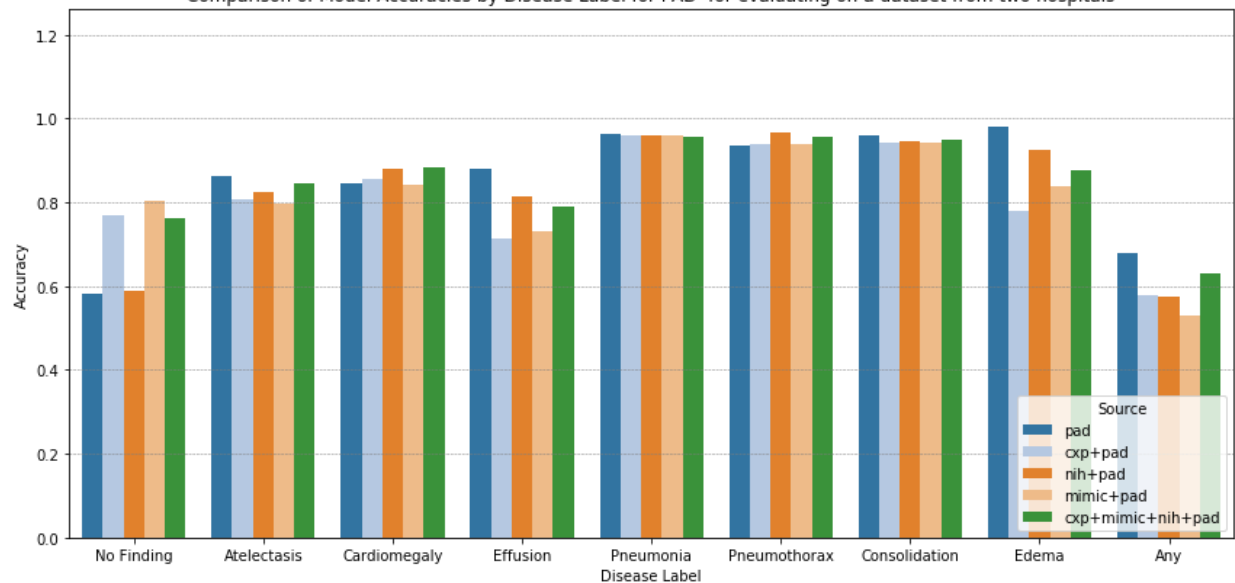





Comparison of Model Accuracies by Disease Label for MIMIC for evaluating on a dataset from two hospitals



Comparison of Model Accuracies by Disease Label for PAD for evaluating on a dataset from two hospitals




Comparison of Model Accuracies by Disease Label for CXP for evaluating on a dataset from two hospitals



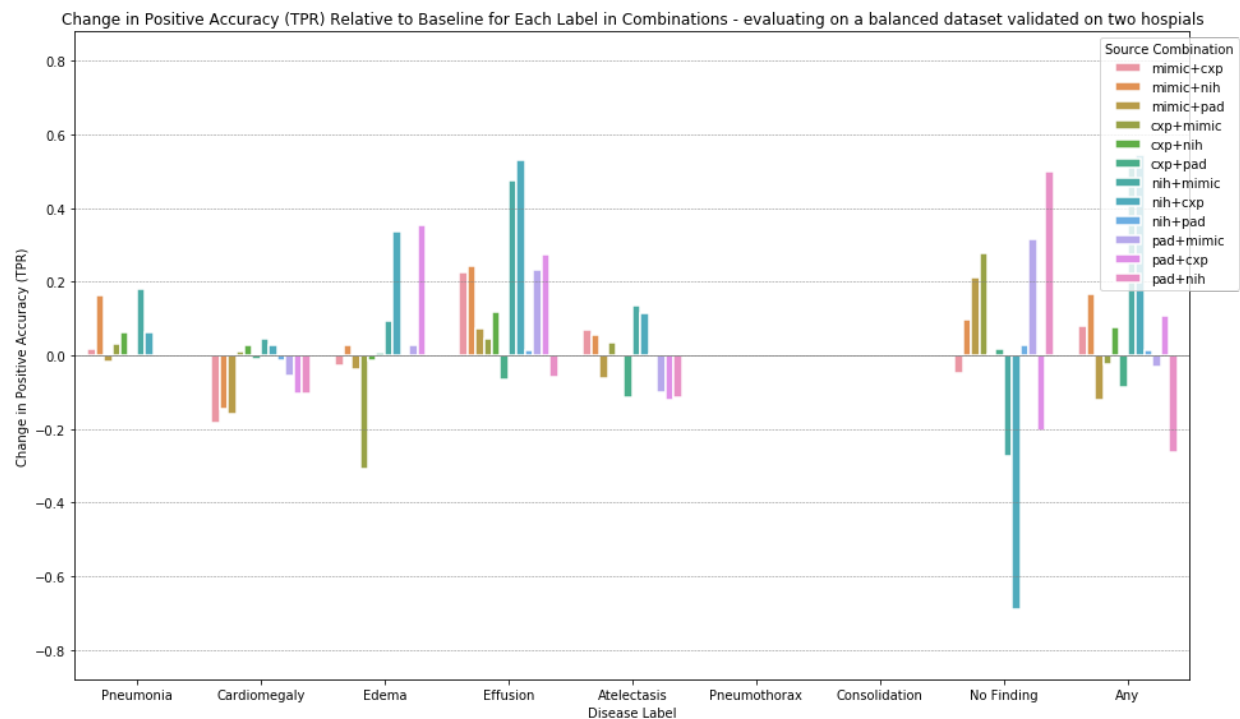
## Results

I had a mixed bag results as is seen from the graphs above. Let's dive in.

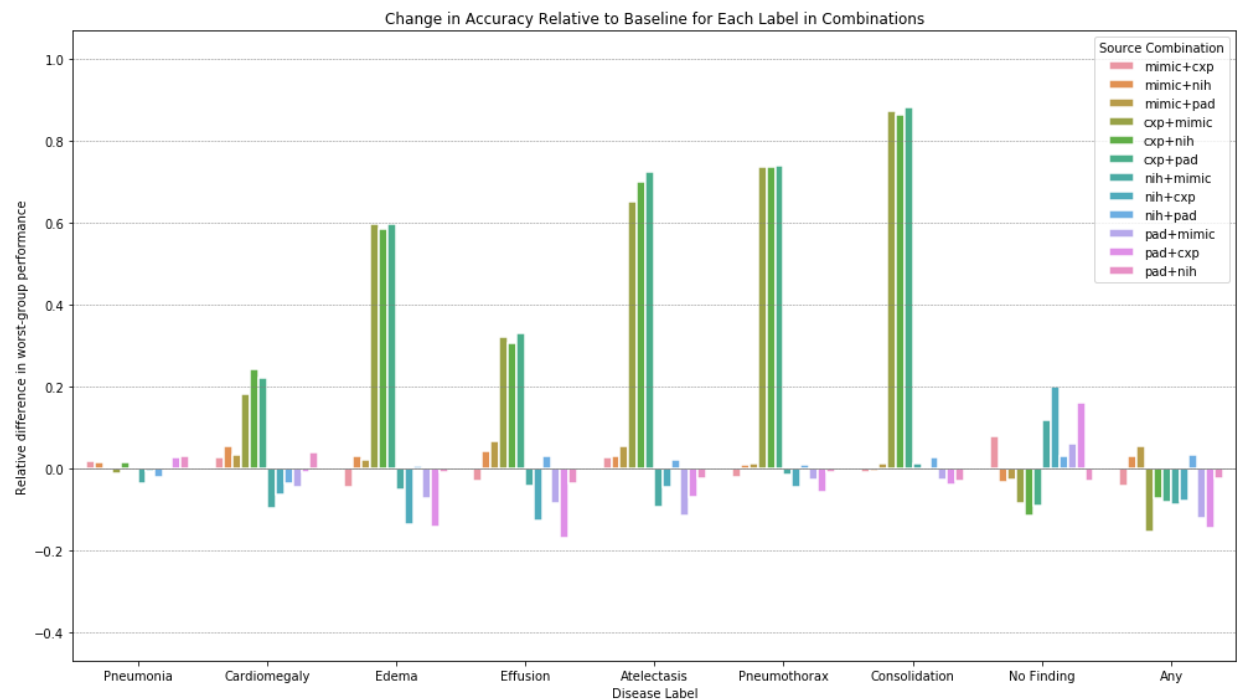
Original paper results:



My closest results:



My initial results:



As it is seen from the tables above, I got a significantly different result than the article authors. Most accuracy changes from incorporating a second dataset were positive (most impacting in the CXP case), and the negative ones were not as bad as the article stated.

## Ablation study

To make sure I am not missing anything, I have repeated the training and the study on the following dataset combinations for training, test and validation:


- Unbalanced dataset
  - train on single (for example, CXP), validate on combined dataset (for example, CXP+NIH)
  - train on combined (for example, CXP+NIH), validate on combined dataset (for example, CXP+NIH)
  - train on combined (for example, CXP+NIH), validate on single dataset (for example, just CXP)
- Balanced dataset
  - train on single (for example, CXP), validate on combined dataset (for example, CXP+NIH)
  - train on combined (for example, CXP+NIH), validate on combined dataset (for example, CXP+NIH)
  - train on combined (for example, CXP+NIH), validate on single dataset (for example, just CXP)

I also re-run the training on different sized results (skip every 20th entry vs every 30th), and in one case ran the study on the full amount of data for one combination.

Balancing affected the results slightly, but I still didn't get the picture that the authors of the article got.


When comparing the balanced vs unbalanced dataset validation, the pattern stays the same.

Unbalanced dataset:






## Balanced dataset:



When both training and validating on two datasets, the margin of error on accuracy was very small.

## Reproducing the results (the closest I got)

The closest results to the paper study I got with a balanced dataset that was evaluating on the two hospital data even if trained with a single hospital.




## Model comparison

## Discussion

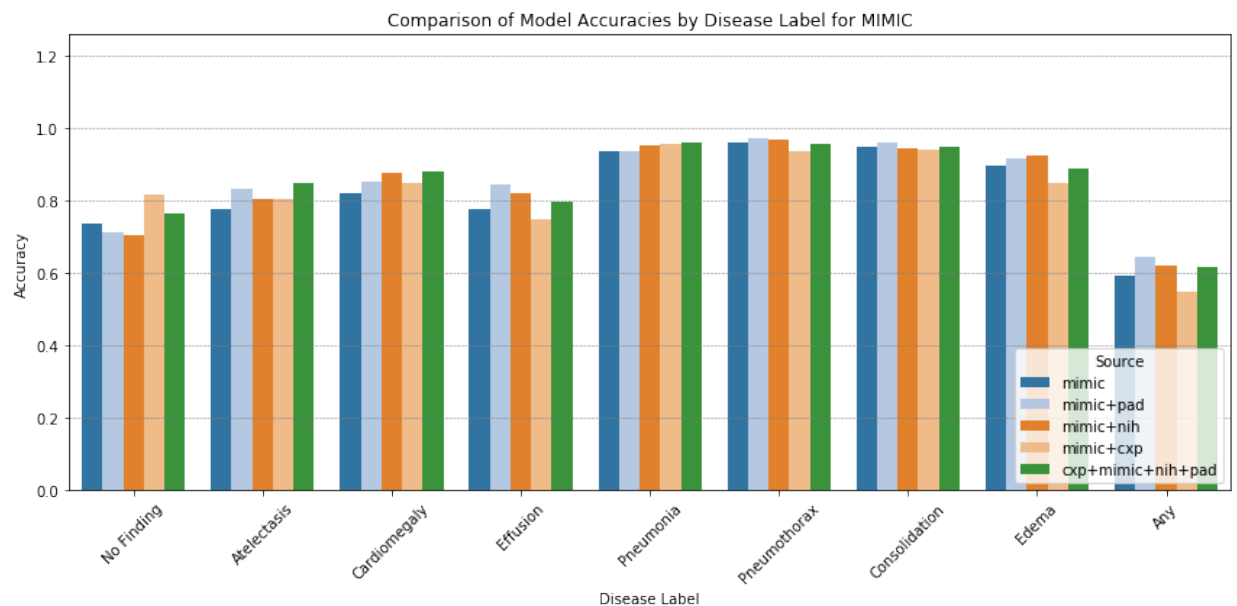
### Is the paper reproducible?

Partially. My resulting visualization ended up being different from the one authors used in the article. However, looking at the results, especially when they are broken down by dataset, one can get to pretty much the same response as the authors of the article did: "It depends", and it looks like balancing the dataset, just as the article stated, doesn't always improve the outcome.


When you look at the results for individual datasets and the resulting combinations, the answer to the main question is very different depending on the dataset:




For PAD, the improvements from adding the second dataset were mainly for the "No finding" category, and the rest generally, even though very slightly, performed worse.



For MIMIC, the answer is "more-or-less", although we can see the very slight positive improvement for most labels.



For NIH, again, the performance mostly decreased.



One interesting exception, CXP almost always significantly benefitted from adding more datasets.

## If the paper is *not* reproducible, explain the results

A portion of the code provided by the article authors was runnable with minimal updates. I was able to reproduce the initial dataset statistics, but wasn't able to run the training code at all.

As a result, I followed the general guidelines given by the authors of the article and wrote my own code - so this may have affected the resulting data I got.

Another factor that might have influenced the resulting data was in data pre-processing. While trying to adapt the authors' code, I discovered an inconsistency in the dataset processing: in some cases the values for certain labels were `True/False`, in some cases numeric `0/1`, and in some cases other numeric value or even a `NaN`.

When adapting the code, I corrected the processing functions so they yield similar-looking results for ease of the combination in a single dataset.

## What was easy

The authors did a great job documenting some parts of the project, for example, access to data. Following the instructions was very easy, and while MIMIC-CXR-JPG dataset access took some time to get, overall the process was a breeze.

## What was difficult

Downloading the datasets is a hassle though, I ran out of space on my laptop, had to buy an external drive and restart the download process for MIMIC-CXR-JPG a few times. In the end, I got a message from my provider that my namely "unlimited" internet for the months was used 100% and I will be charged for each extra GB I use.

There are a few notebooks and standalone scripts provided to process the data. While it is possible to figure out what steps need to be done in what order, many of the parts of the process are not documented. 'pyproject.toml' did not run successfully for me, and I got stuck for a while, trying to figure out why and how to run it (I have a suspicion my processor architecture is not supported).

As a result, I opted to re-implement the training and model validation myself. There is code for training and validation in the project, which has a lot of comments (great!), but the process itself is not well documented, so the reproducer is left figuring out which steps in the code are needed and which are not, and how to adapt it to use for their experiment. The code is very general and there is a lot of it. There are some pointers in the readme, but they were not sufficient for reproducing things successfully without additional modification. wandb wasn't working for me either.

The data is not processed evenly / equally for each dataset, there are different values for the same labels (`NaN`, `True/False`, `1/0`, `1.1/0.0`). I had to write some processing code to make sure to mitigate those differences.

Additional complication was due to the fact that the amount of data is very large. Any training or processing takes a long time, the notebook kernel dies frequently and the overall process is frustrating. Downloading MIMIC took a week and ate all my provider's internet allowance for a month. Running the training on a full dataset proved it to be difficult due to

the whole different set of circumstances: I encountered out of memory issues, kernel panic, random automatic Windows updates, power outage, kids getting to the computer and switching the power supply off. In the end, the estimation to run the training on the full dataset was circa 16+ days, so I opted for running the code on the subset of the data instead.

I tried to avoid multiple separate files and scripts, and pulled many of the data preprocessing into my notebook. However, this increased the runtime of the notebook significantly. Additional factor affecting the runtime is the size of the input data, even when working on one dataset. I was never able to achieve the 8-minute runtime, in fact, my record was around 10 hours.

## Suggestions for the author

Trim the codebase leaving only relevant parts. Add documentation for the training and validation process. Add some background on why wandb is used and how to use it for this project correctly. Provide a suggested order of execution for the notebooks. Provide the instructions and code to plot the results.

## Future plans

I am still very interested in answering the question why my results, even though providing a similar answer, looked so different from the article authors'. I plan to do a few more things to try and figure this out:

- Try different models (for example, `LinearSVC` and other models tried by the authors of the article before arriving on their final model), and hyperparameters
- Try and incorporate more of the training and validation code provided by the article authors, and/or triple check mine
- Run the training on the full dataset (ensuring the computer is on backup power supply and is unreachable to anyone trying to switch it off)

# References

1. Rhys Compton; Lily Zhang; Aahlad Puli; Rajesh Ranganath, When More is Less: Incorporating Additional Datasets Can Hurt Performance By Introducing Spurious Correlations, arXiv preprint, 2023-08-09, Accepted at MLHC 2023, doi: [10.48550/arXiv.2308.04431](https://doi.org/10.48550/arXiv.2308.04431) (<https://doi.org/10.48550/arXiv.2308.04431>)
2. Haoran Zhang, Natalie Dullerud, Laleh Seyyed-Kalantari, Quaid Morris, Shalmali Joshi, and Marzyeh Ghassemi. An empirical framework for domain generalization in clinical settings. In Proceedings of the Conference on Health, Inference, and Learning, pages 279–290, 2021, doi: [10.48550/arXiv.2103.11163](https://doi.org/10.48550/arXiv.2103.11163) (<https://doi.org/10.48550/arXiv.2103.11163>)
3. Gao Huang, Zhuang Liu, Laurens Van Der Maaten, and Kilian Q Weinberger. Densely connected convolutional networks. In Proceedings of the IEEE conference on computer vision and pattern recognition, pages 4700–4708, 2017, doi: [10.48550/arXiv.1608.06993](https://doi.org/10.48550/arXiv.1608.06993) (<https://doi.org/10.48550/arXiv.1608.06993>)
4. Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In 2009 IEEE Conference on Computer Vision and Pattern Recognition, pages 248–255, 2009, doi: [10.1109/CVPR.2009.5206848](https://doi.org/10.1109/CVPR.2009.5206848) (<https://doi.org/10.1109/CVPR.2009.5206848>)
5. John R Zech, Marcus A Badgeley, Manway Liu, Anthony B Costa, Joseph J Titano, and Eric Karl Oermann. Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: a cross-sectional study. PLoS medicine, 15(11): e1002683, 2018, doi: [10.1371/journal.pmed.1002683](https://doi.org/10.1371/journal.pmed.1002683) (<https://doi.org/10.1371/journal.pmed.1002683>)

In [ ]: